Behavioral Cloning and Interactive Imitation Learning

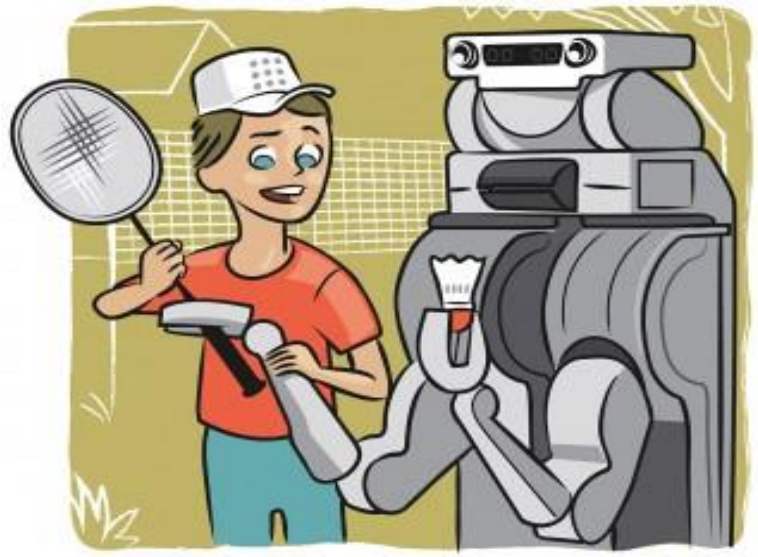

Instructor: Daniel Brown

[Some slides adapted from Sergey Levine (CS 285) and Alina Vereshchaka (CSE4/510)]

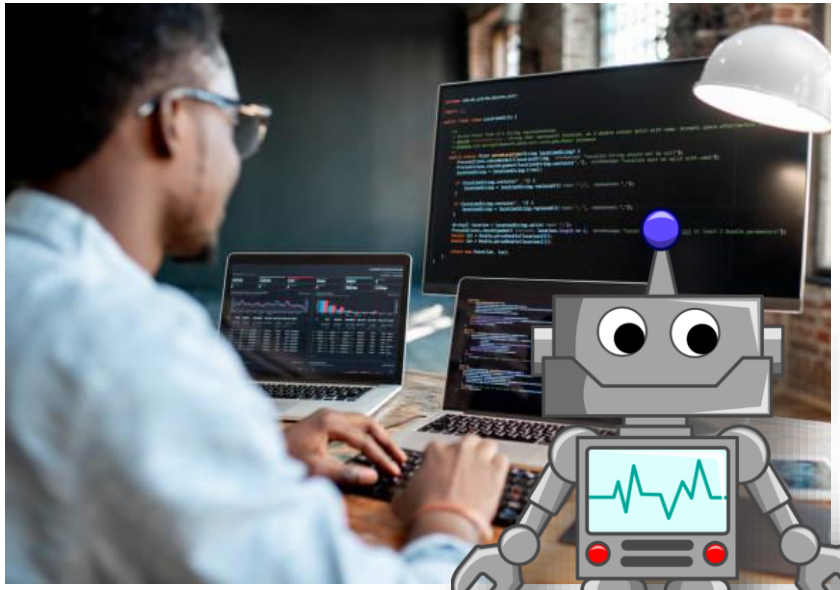

Brief Machine Learning Refresher

There are roughly 3 main branches of machine learning

- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

Supervised Learning

- **Setting/Assumptions**: In supervised learning, the model is trained on labeled data, where the input data is paired with the correct output (i.e., the "ground truth").
- **Goal**: To learn a mapping from inputs to outputs so that the model can predict the output for new, unseen inputs.
- **Common Use Cases**:
  - Classification (e.g., spam email detection, image recognition).
  - Regression (e.g., predicting house prices, stock market trends).
- **Example models**:
  - Linear regression, decision trees, support vector machines, and neural networks.

Classification

$$\text{Cross-Entropy Loss} = -\sum_{i=1}^{C} y_i \log(\hat{y}_i)$$

Regression

$$\text{MSE Loss} = \frac{1}{N}\sum_{i=1}^{N}(\hat{y}_i - y_i)^2$$

Unsupervised Learning

- **Setting/Assumptions**: In unsupervised learning, the model is trained on data without labeled outputs. It seeks to find patterns, structures, or relationships in the data. No “ground truth” labels.
- **Goal**: To explore the data and identify meaningful clusters, associations, or representations.
- **Common Use Cases**:
  - Clustering (e.g., customer segmentation).
  - Dimensionality reduction (e.g., PCA for visualization).
  - Anomaly detection (e.g., fraud detection).
- **Example models**:
  - K-means clustering, hierarchical clustering, and autoencoders.

Auto-Encoders

Reinforcement Learning

- **Setting/Assumptions**: Reinforcement learning (RL) involves training an agent to make decisions by interacting with an environment. The agent learns through trial and error (receiving rewards and penalties), optimizing its behavior to maximize cumulative rewards.
- **Goal**: To learn a policy that maps states of the environment to actions that achieve the highest reward.
- **Common Use Cases**:
  - Game-playing AI (e.g., AlphaGo, chess-playing bots).
  - Robotics (e.g., autonomous navigation).
  - Dynamic resource allocation (e.g., in networking or traffic management).
- **Examples**:
  - Q-learning, Deep Q-Networks (DQN), and Proximal Policy Optimization (PPO).

Reinforcement Learning

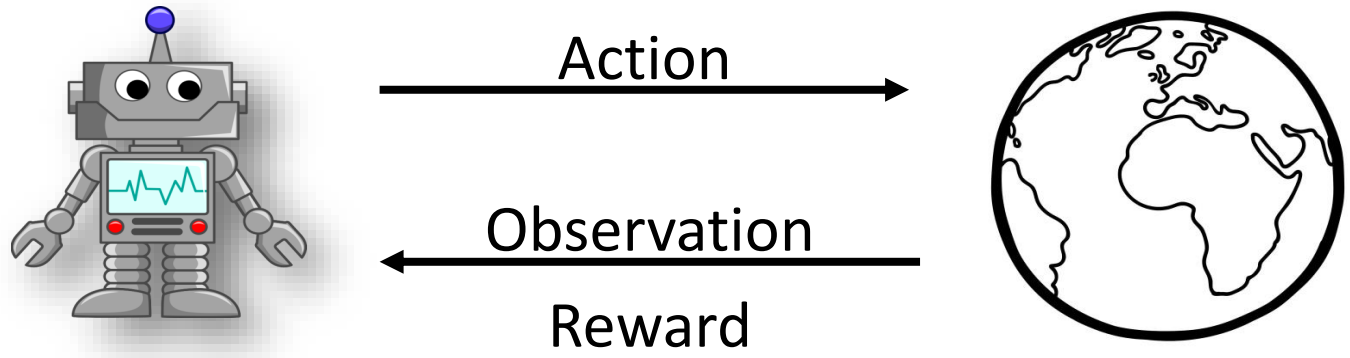

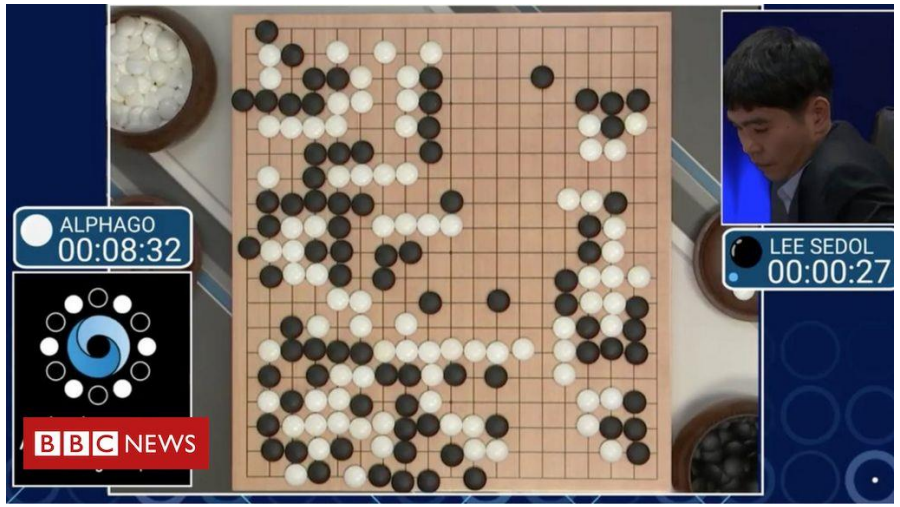

Reinforcement Learning

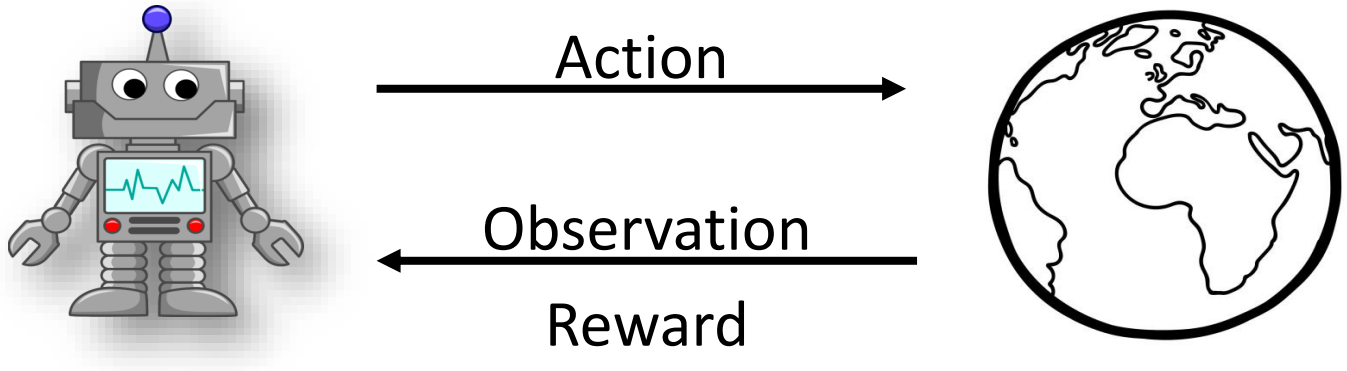

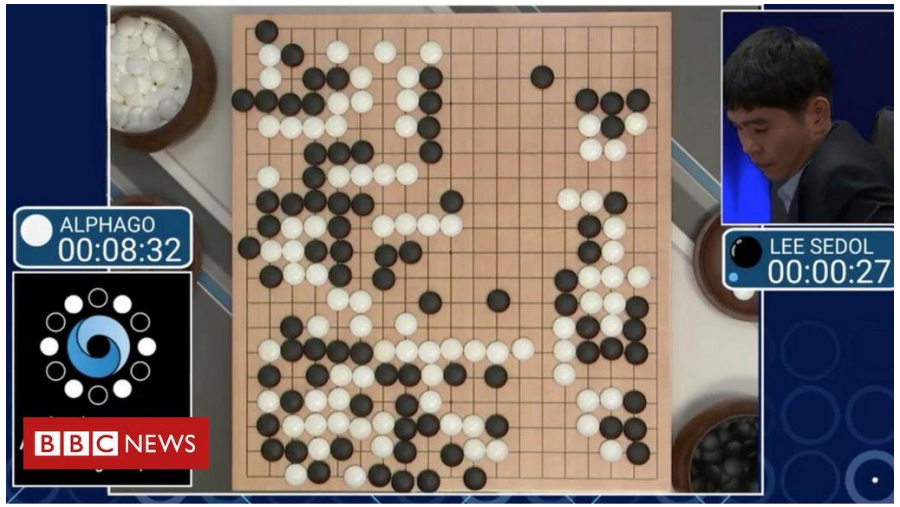

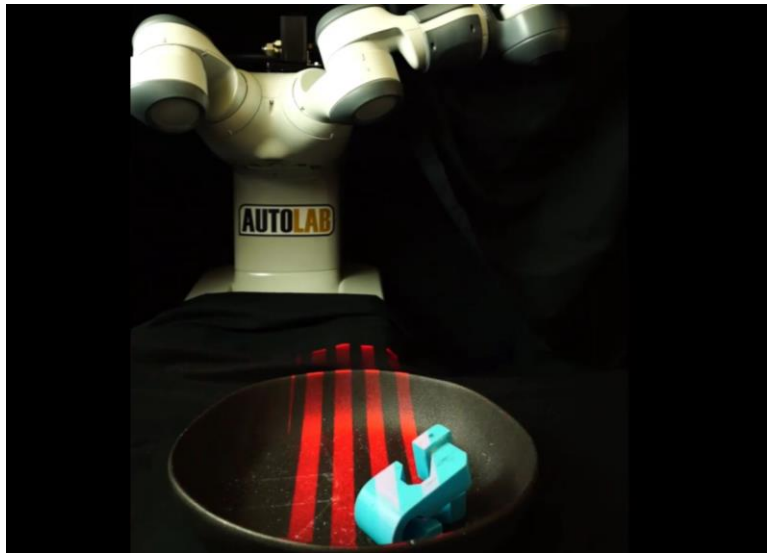

Reward engineering is hard!

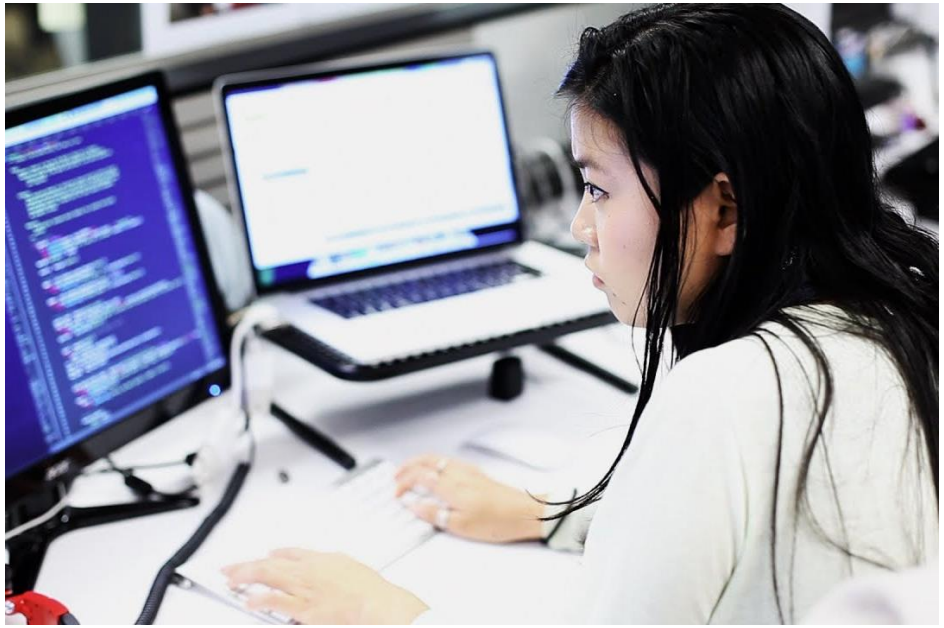

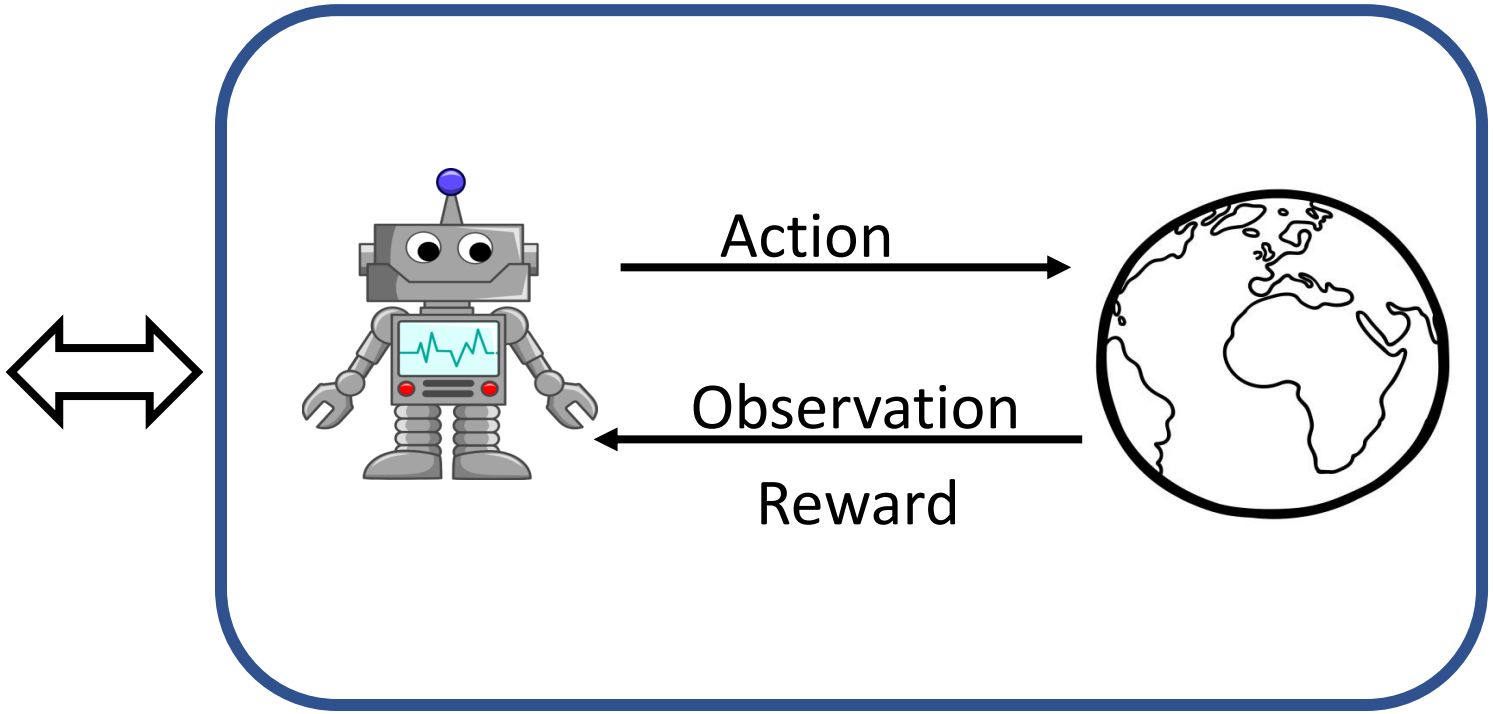

Reward engineering is hard!

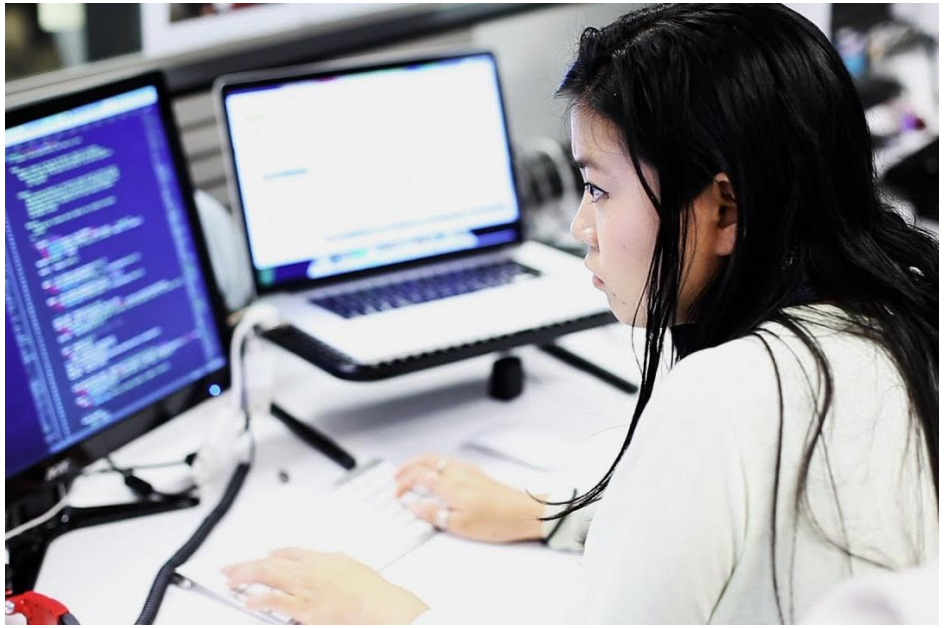

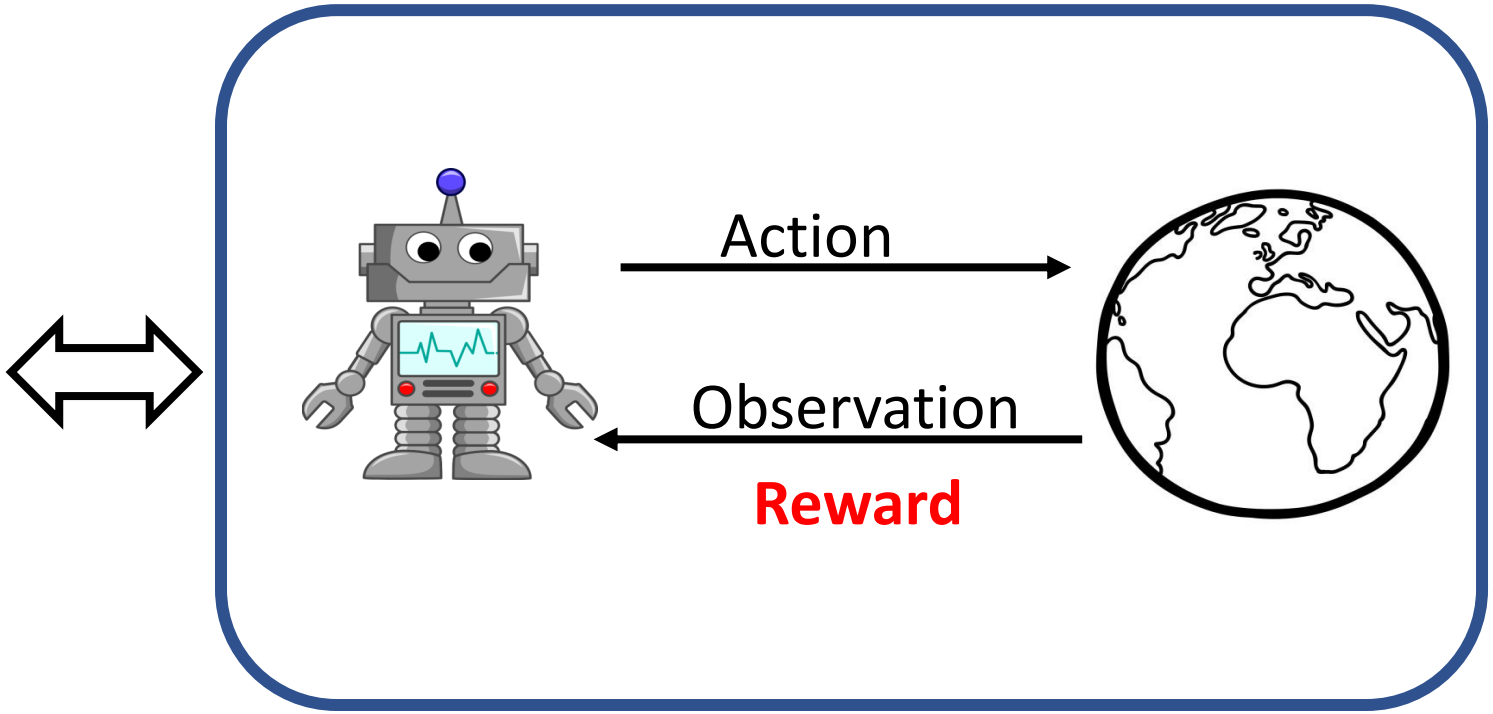

Reward engineering is hard!

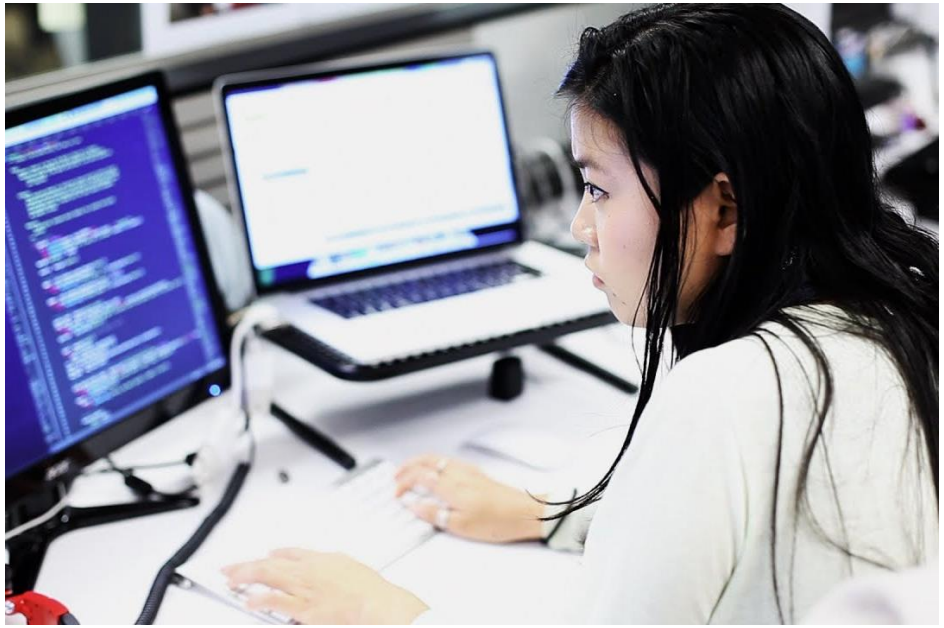

Reinforcement learning is hard…even with a reward function!

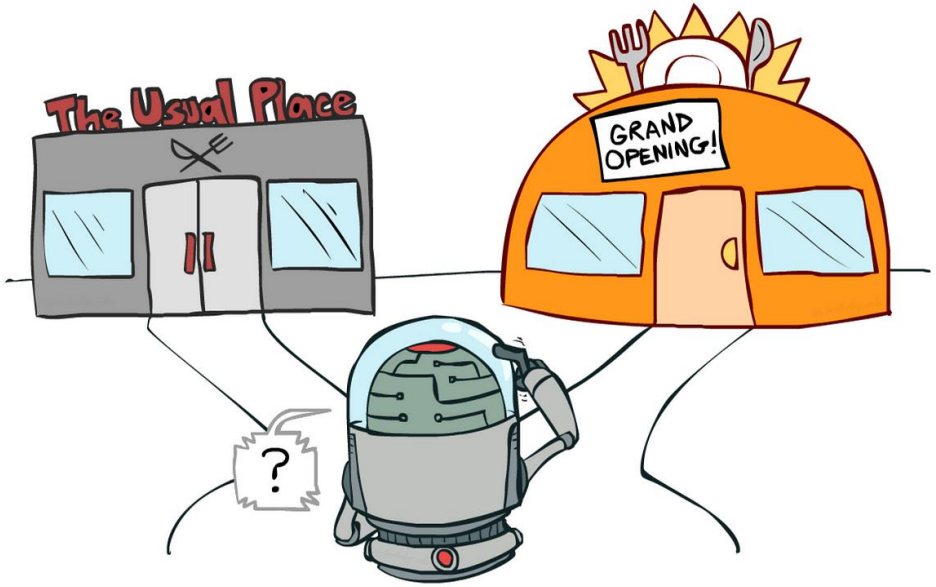

Imitation Learning:

Learn a policy from examples of good behavior.

- Often showing is easier than telling.
- Alleviates problem of exploration.

Behavioral Cloning

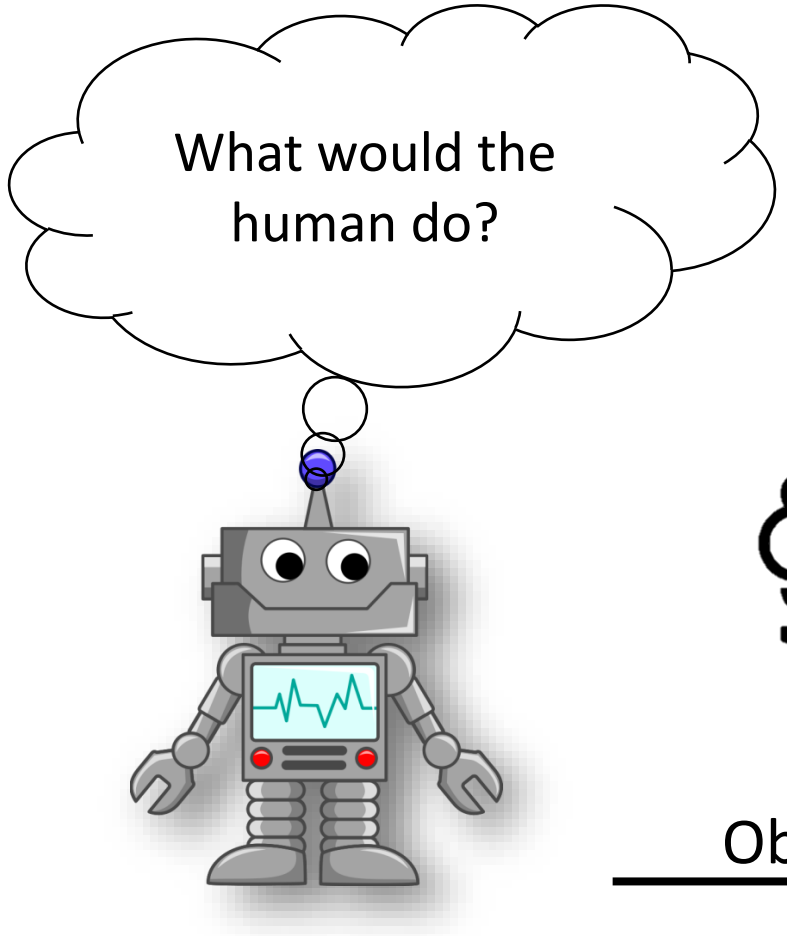

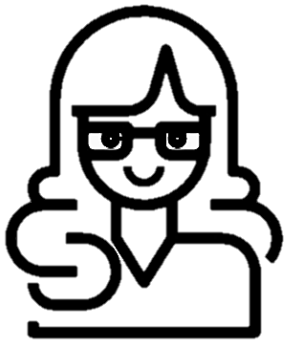

Action

Observation

Observation

Policy π

Action

Inverse Reinforcement Learning

We'll talk about this next time!

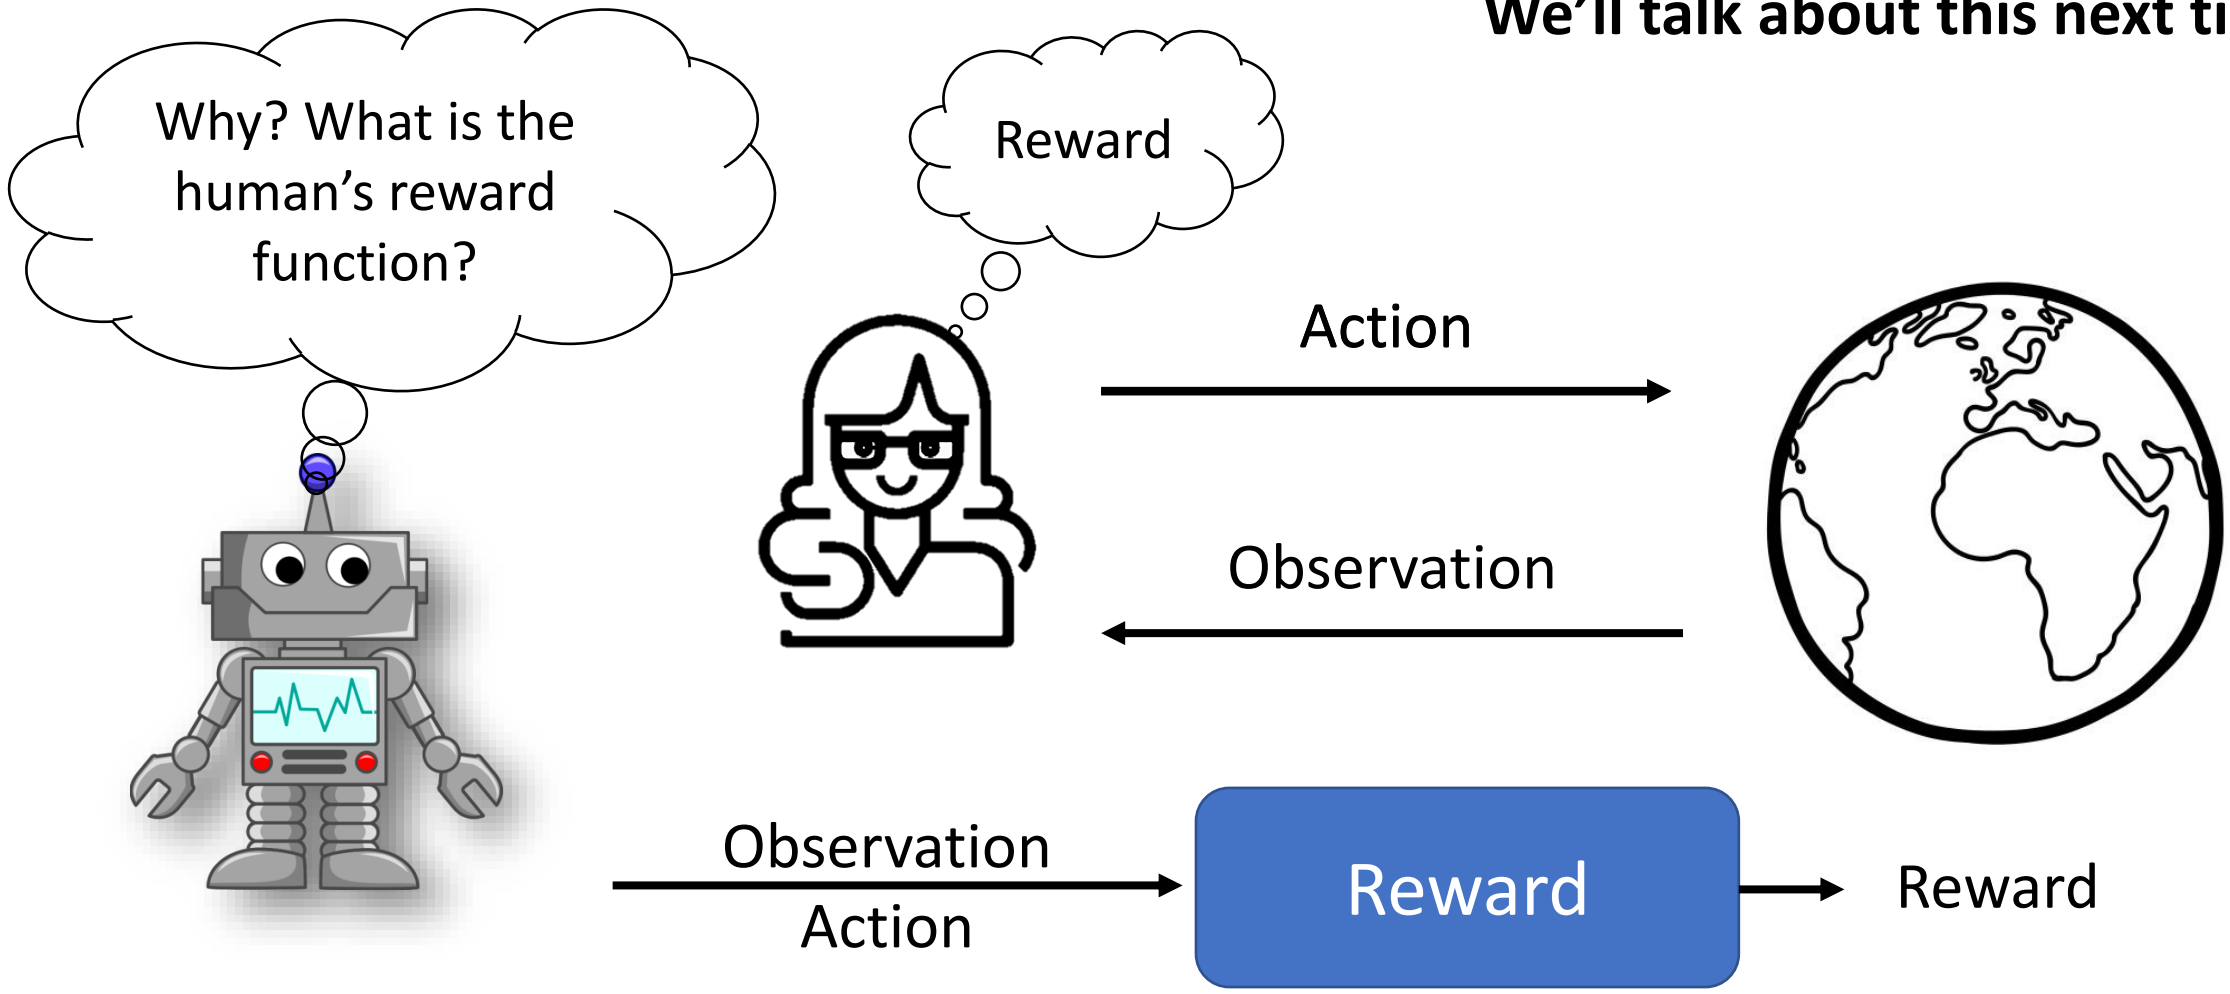

Imitation Learning via Behavioral Cloning

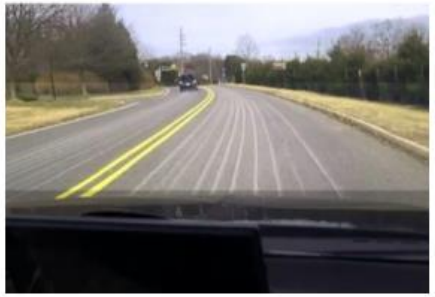

$\mathbf{o}_t$

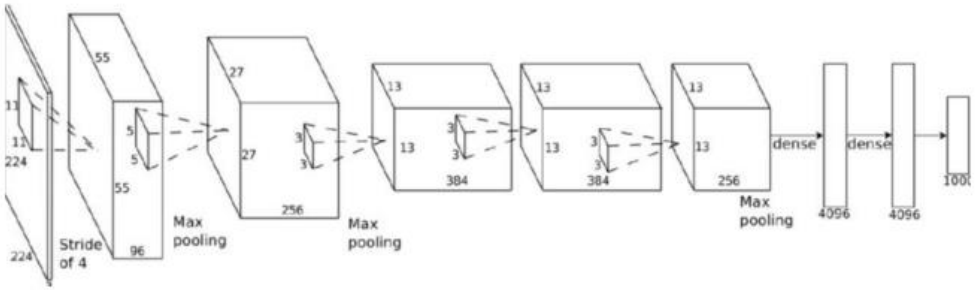

$\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

$\mathbf{a}_t$

$\mathbf{o}_t$

$\mathbf{a}_t$

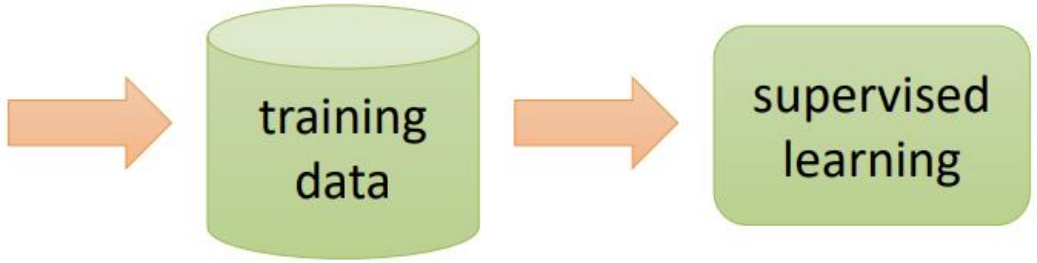

$\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

Live demo

python test_gym.py

python mountain_car_bc.py --num_demos 1

ALVINN: One of the first imitation learning systems

ALVINN: **A**utonomous **L**and **V**ehicle **I**n a **N**eural **N**etwork
1989

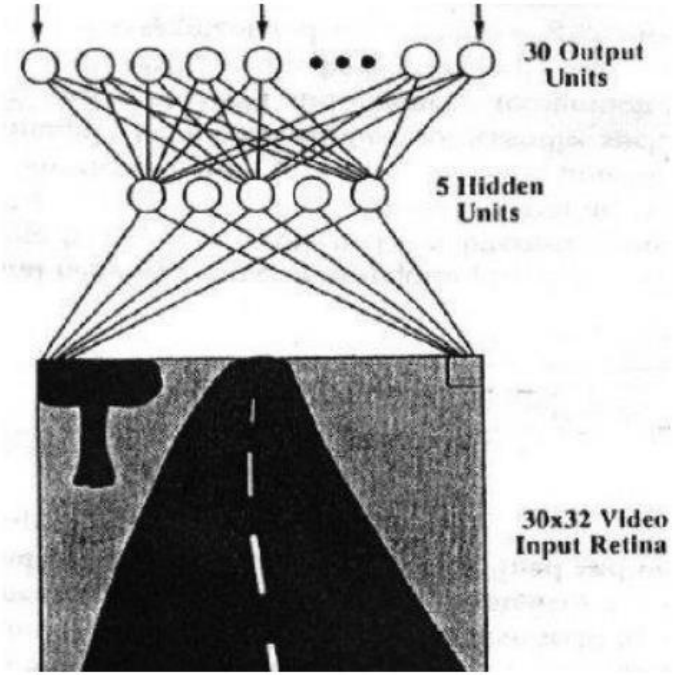

ALVINN: One of the first imitation learning systems

ALVINN: **A**utonomous **L**and **V**ehicle **I**n a **N**eural **N**etwork
1989

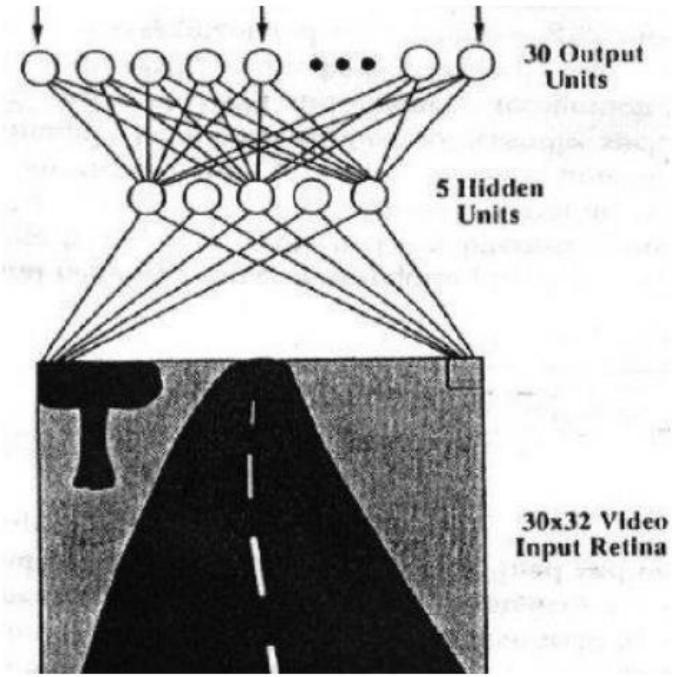

What if you don’t have actions?

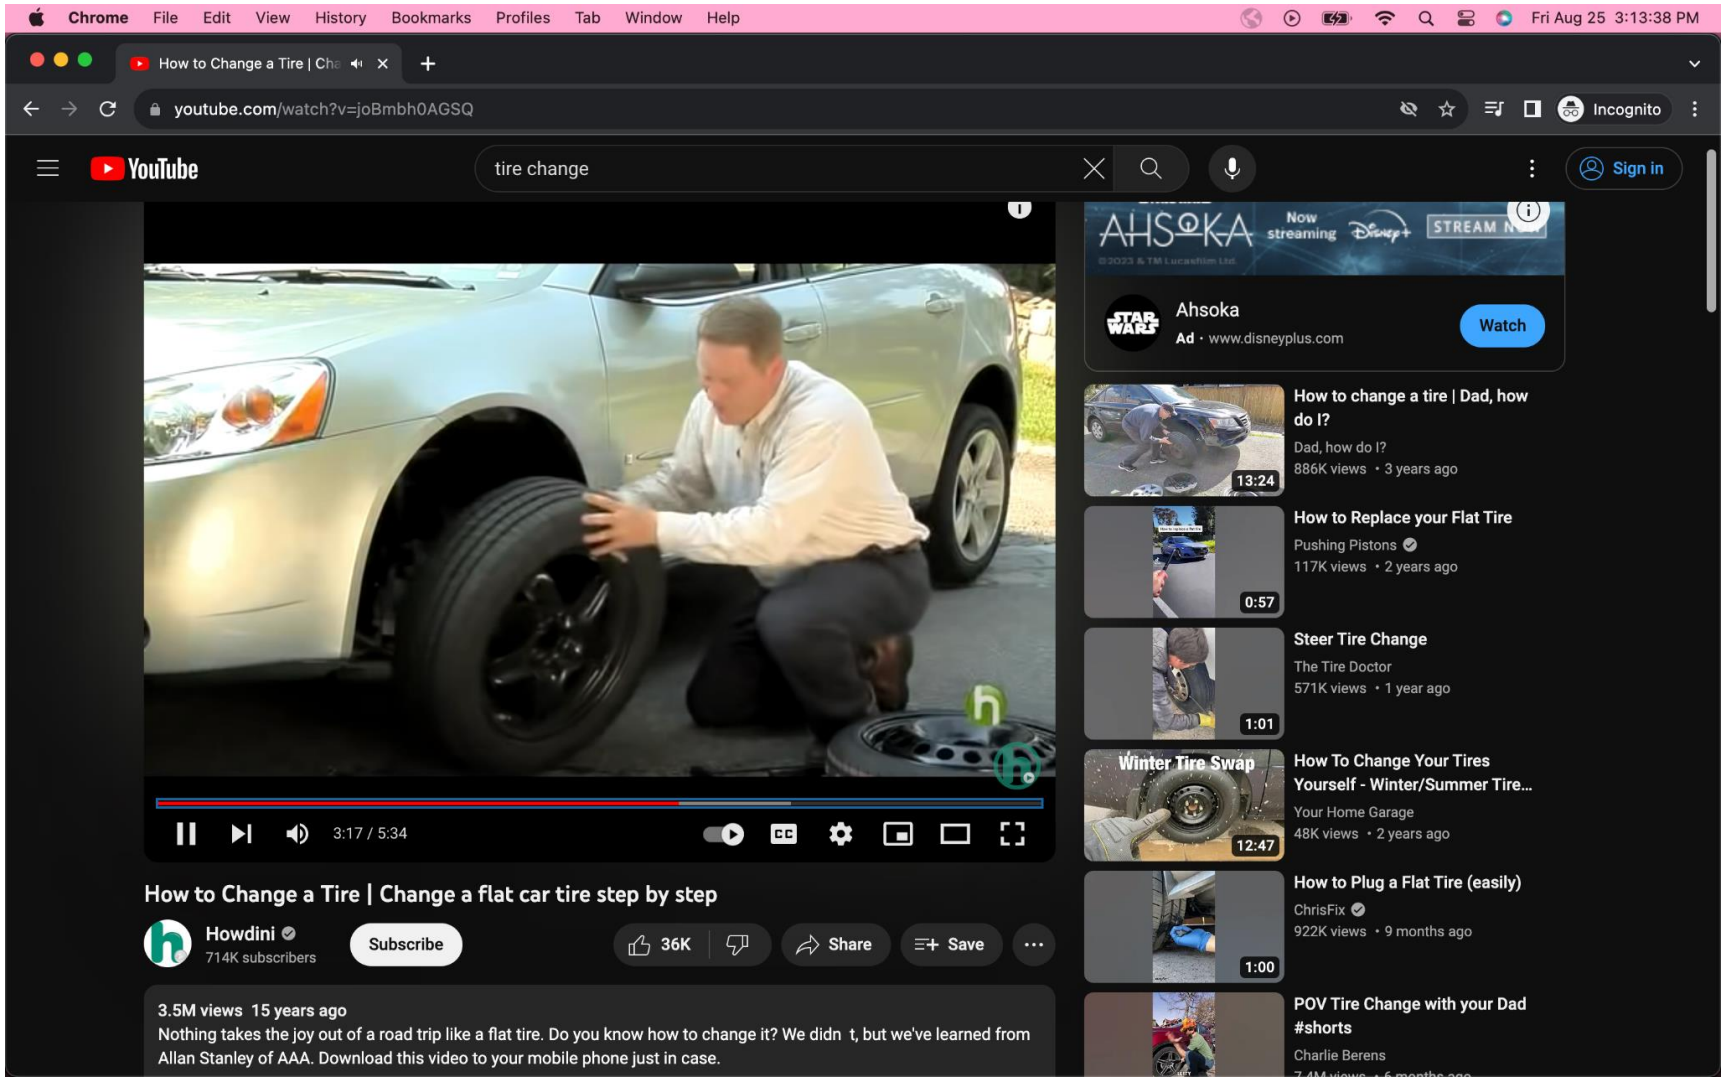

Behavioral Cloning from Observation (Torabi et al. 2018)

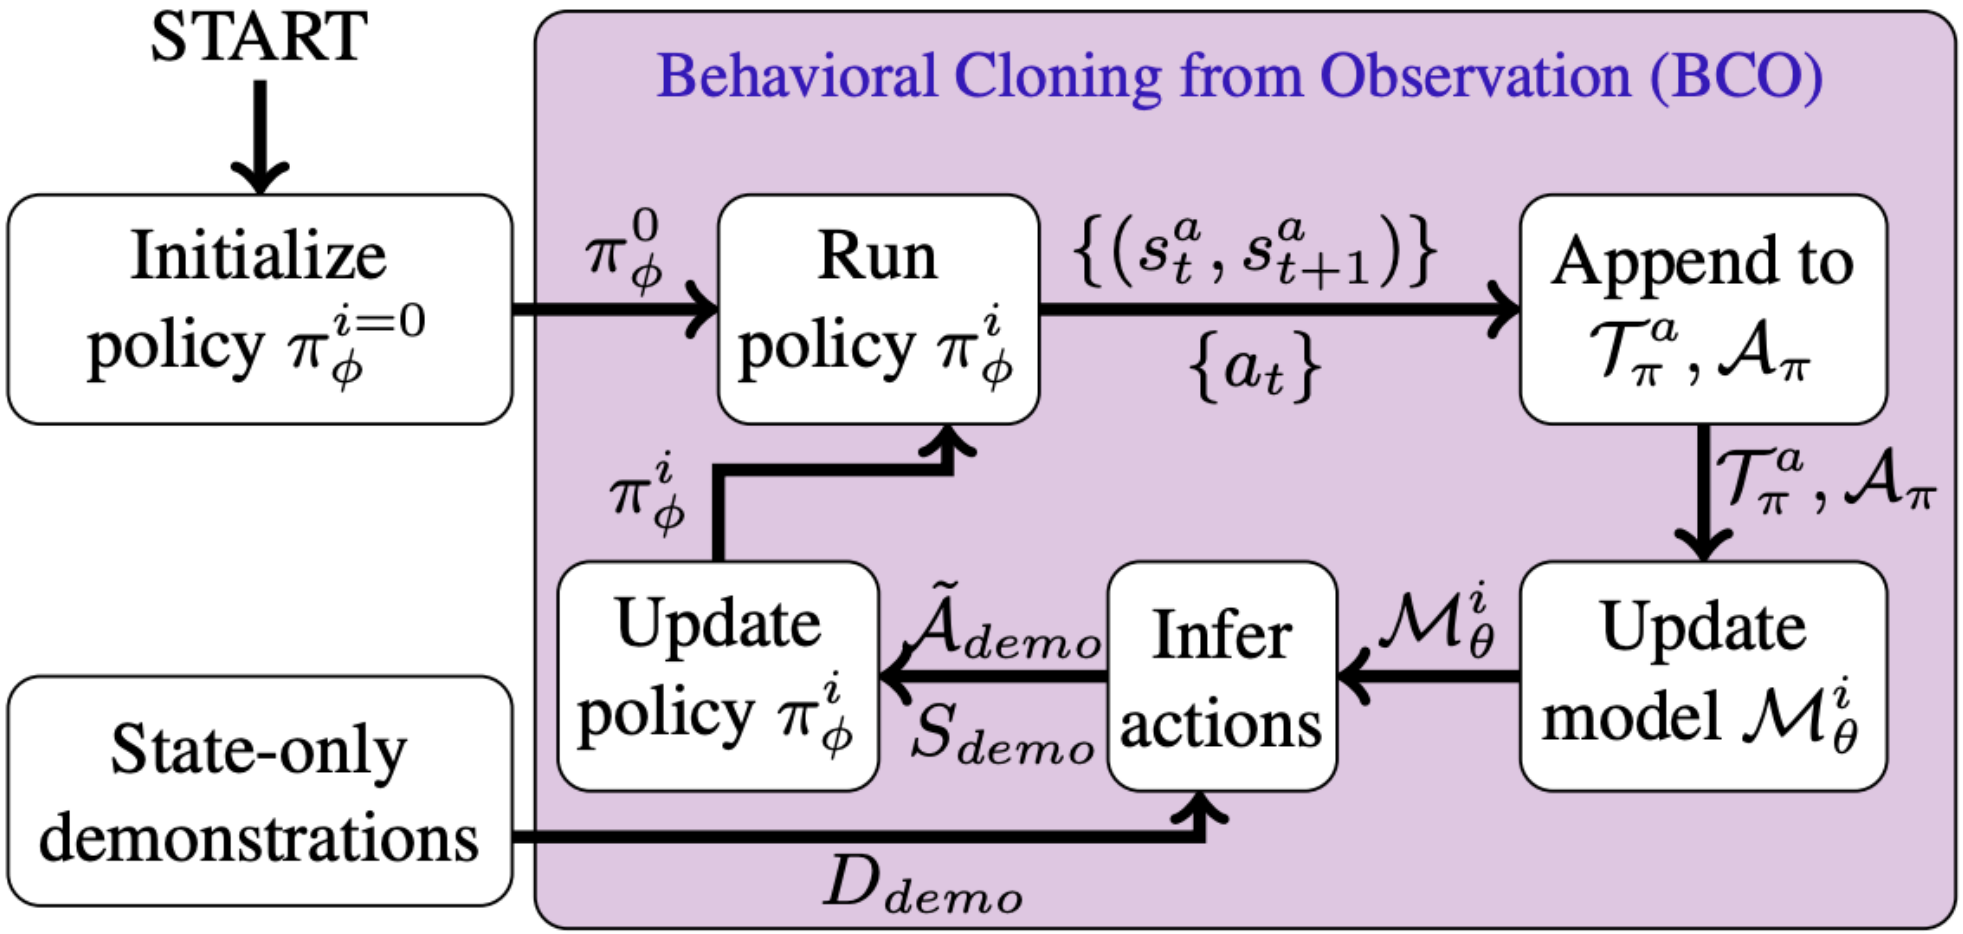

What could go wrong?

$\mathbf{o}_t$
$\mathbf{a}_t$

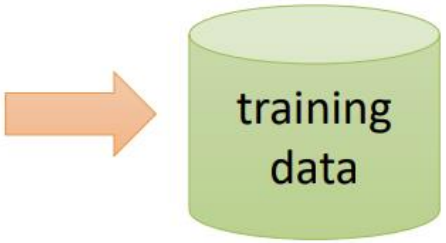

$\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

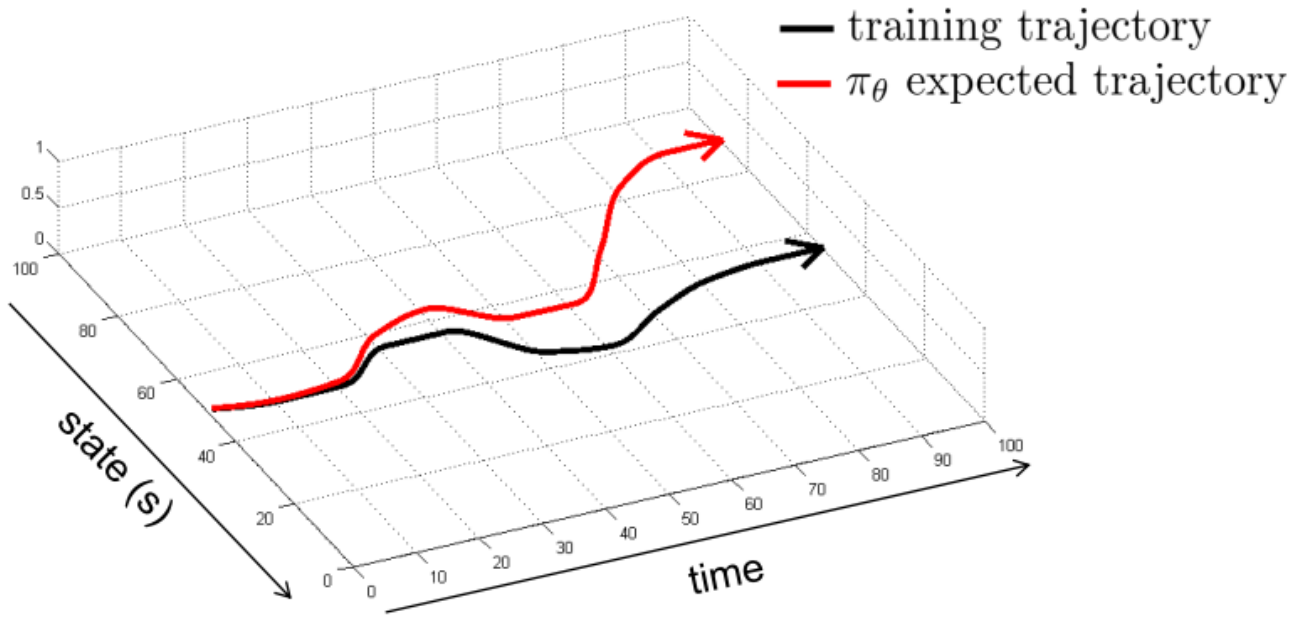

Distribution Shift

$$p_{\pi^*}(o_t) \neq p_{\pi_\theta}(o_t)$$

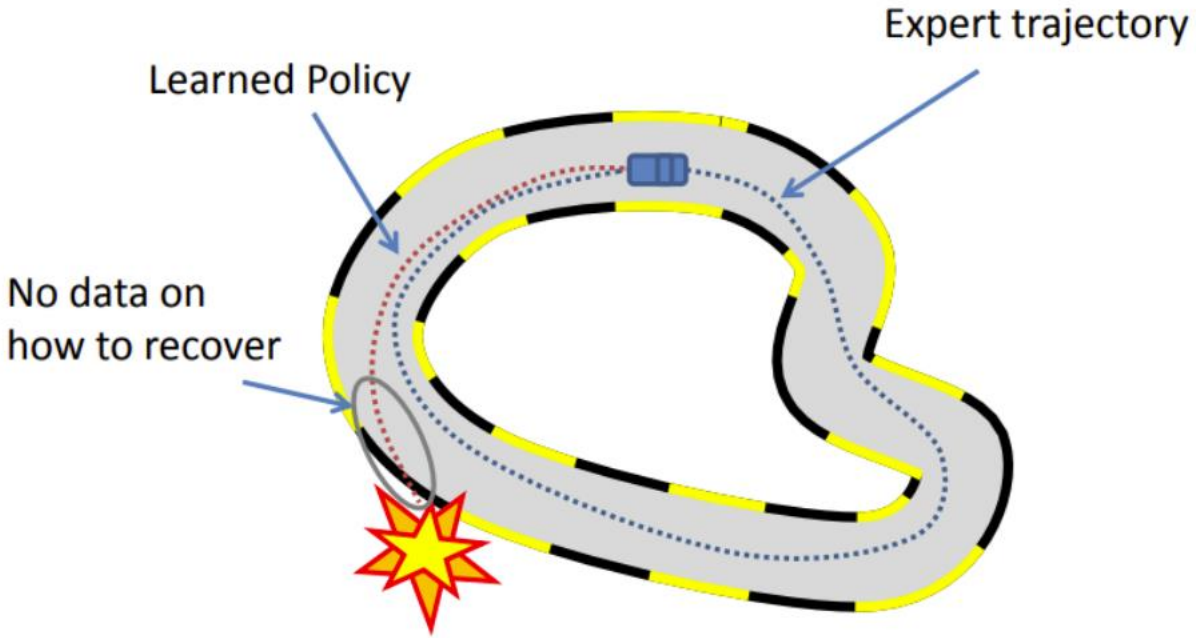

	Supervised Learning	Supervised Learning + Control
Train	$(x, y) \sim D$	$s \sim P(\cdot \mid s, \pi^*(s))$
Test	$(x, y) \sim D$	$s \sim P(\cdot \mid s, \pi(s))$

But it still can work in practice…

Bojarski et al. ‘16, NVIDIA

How?

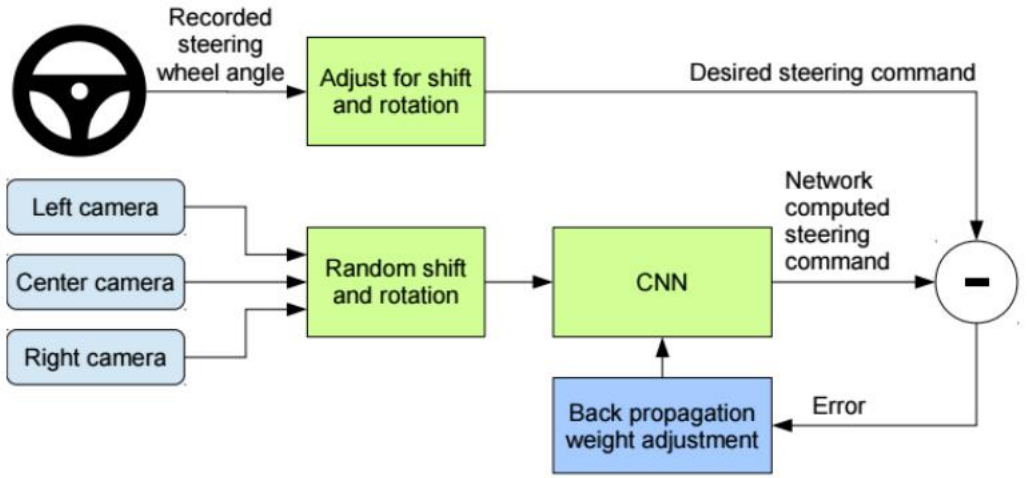

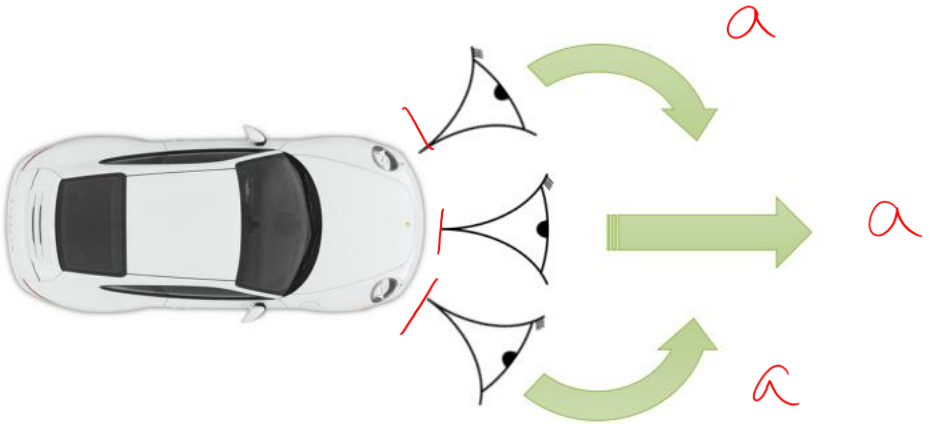

Bojarski et al. ‘16, NVIDIA

A Machine Learning Approach to Visual Perception of Forest Trails for Mobile Robots

Alessandro Giusti[1], Jérôme Guzzi[1], Dan C. Cireşan[1], Fang-Lin He[1], Juan P. Rodríguez[1]
Flavio Fontana[2], Matthias Faessler[2], Christian Forster[2]
Jürgen Schmidhuber[1], Gianni Di Caro[1], Davide Scaramuzza[2], Luca M. Gambardella[1]

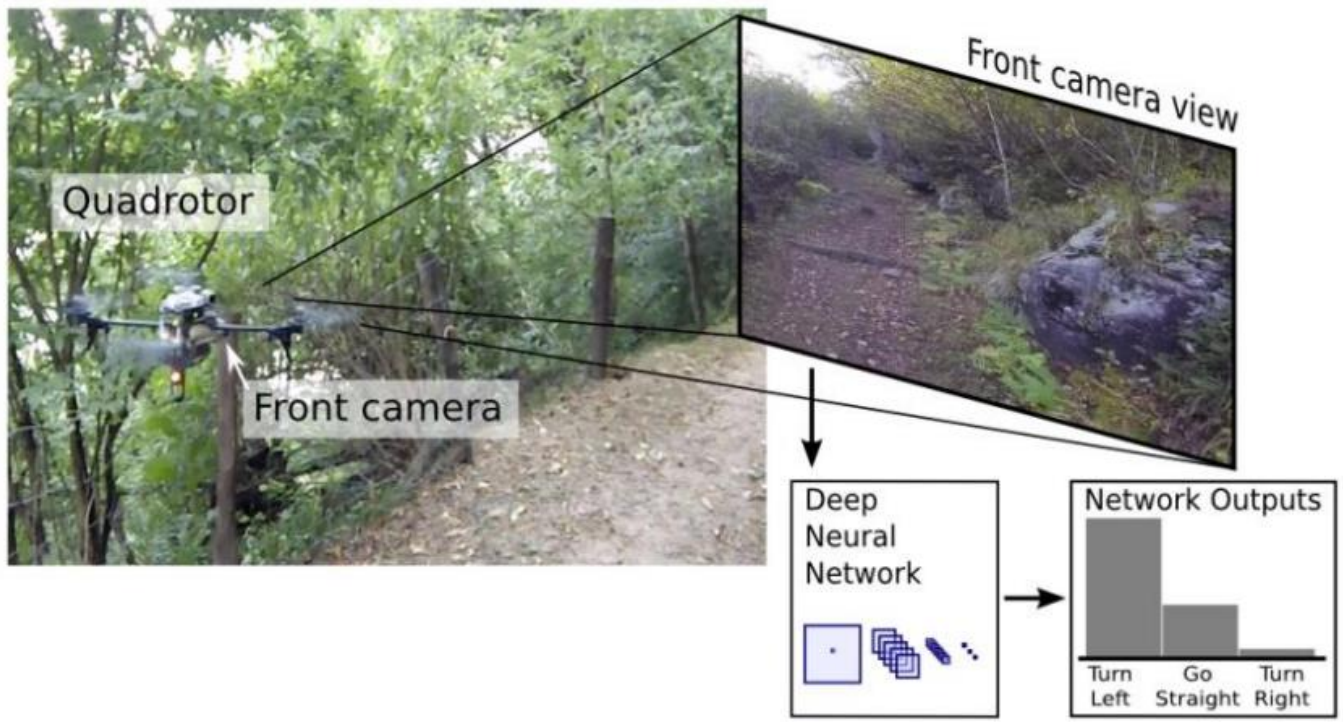


Deep Neural Network
Control Signal

Can we make it work more often?

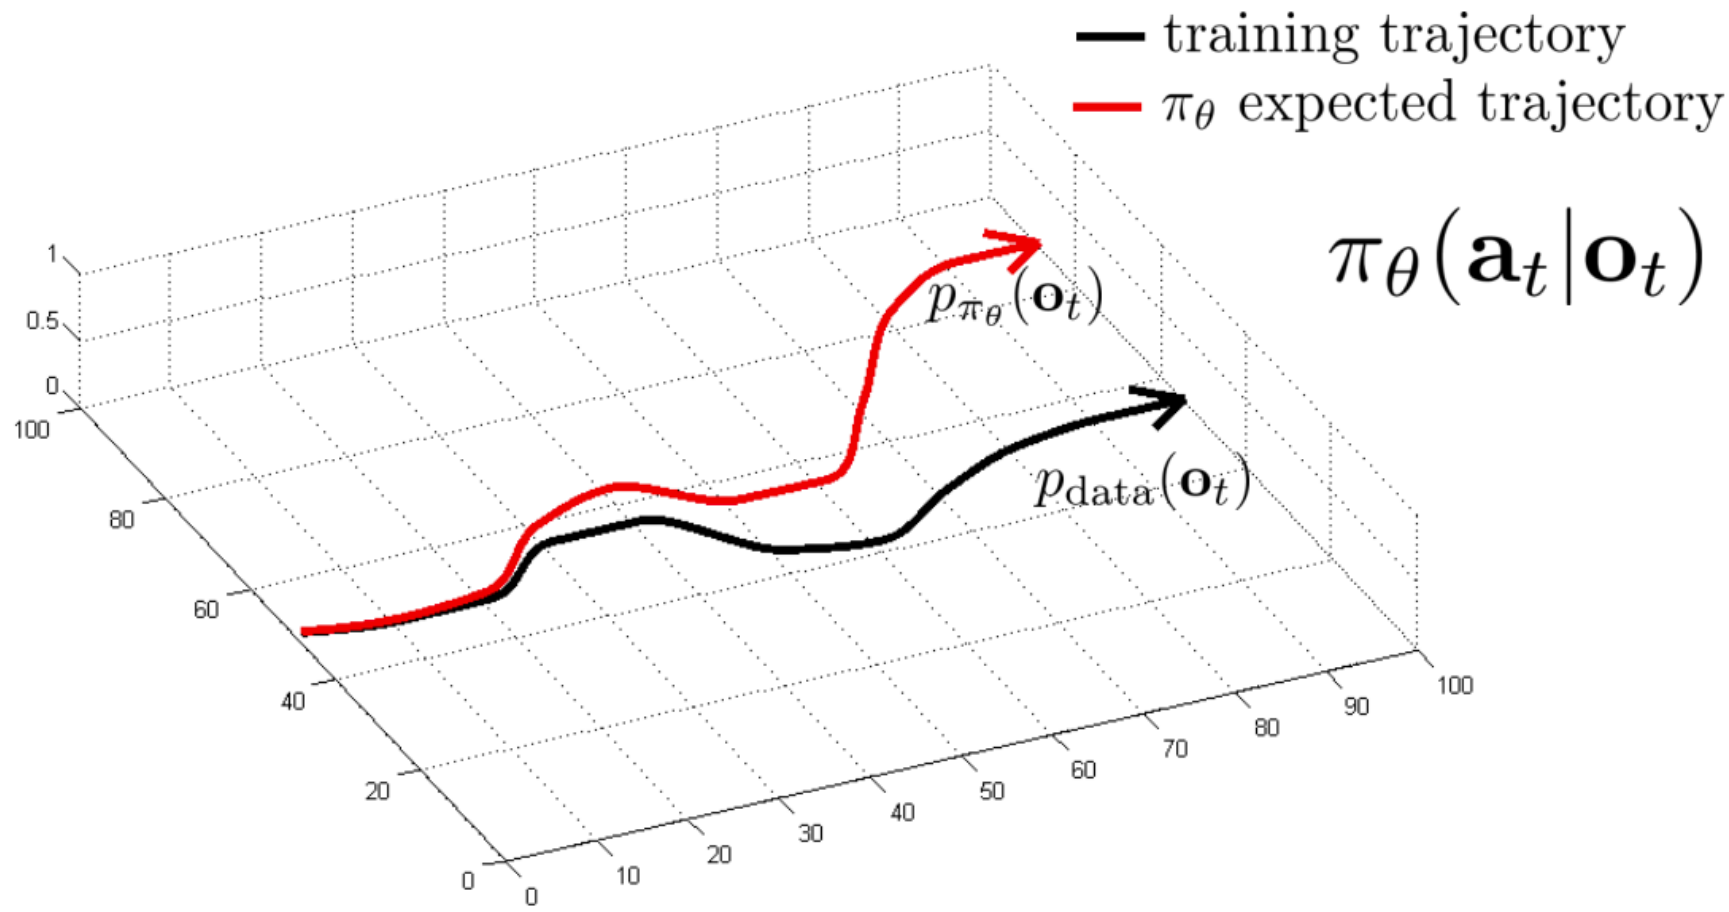

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

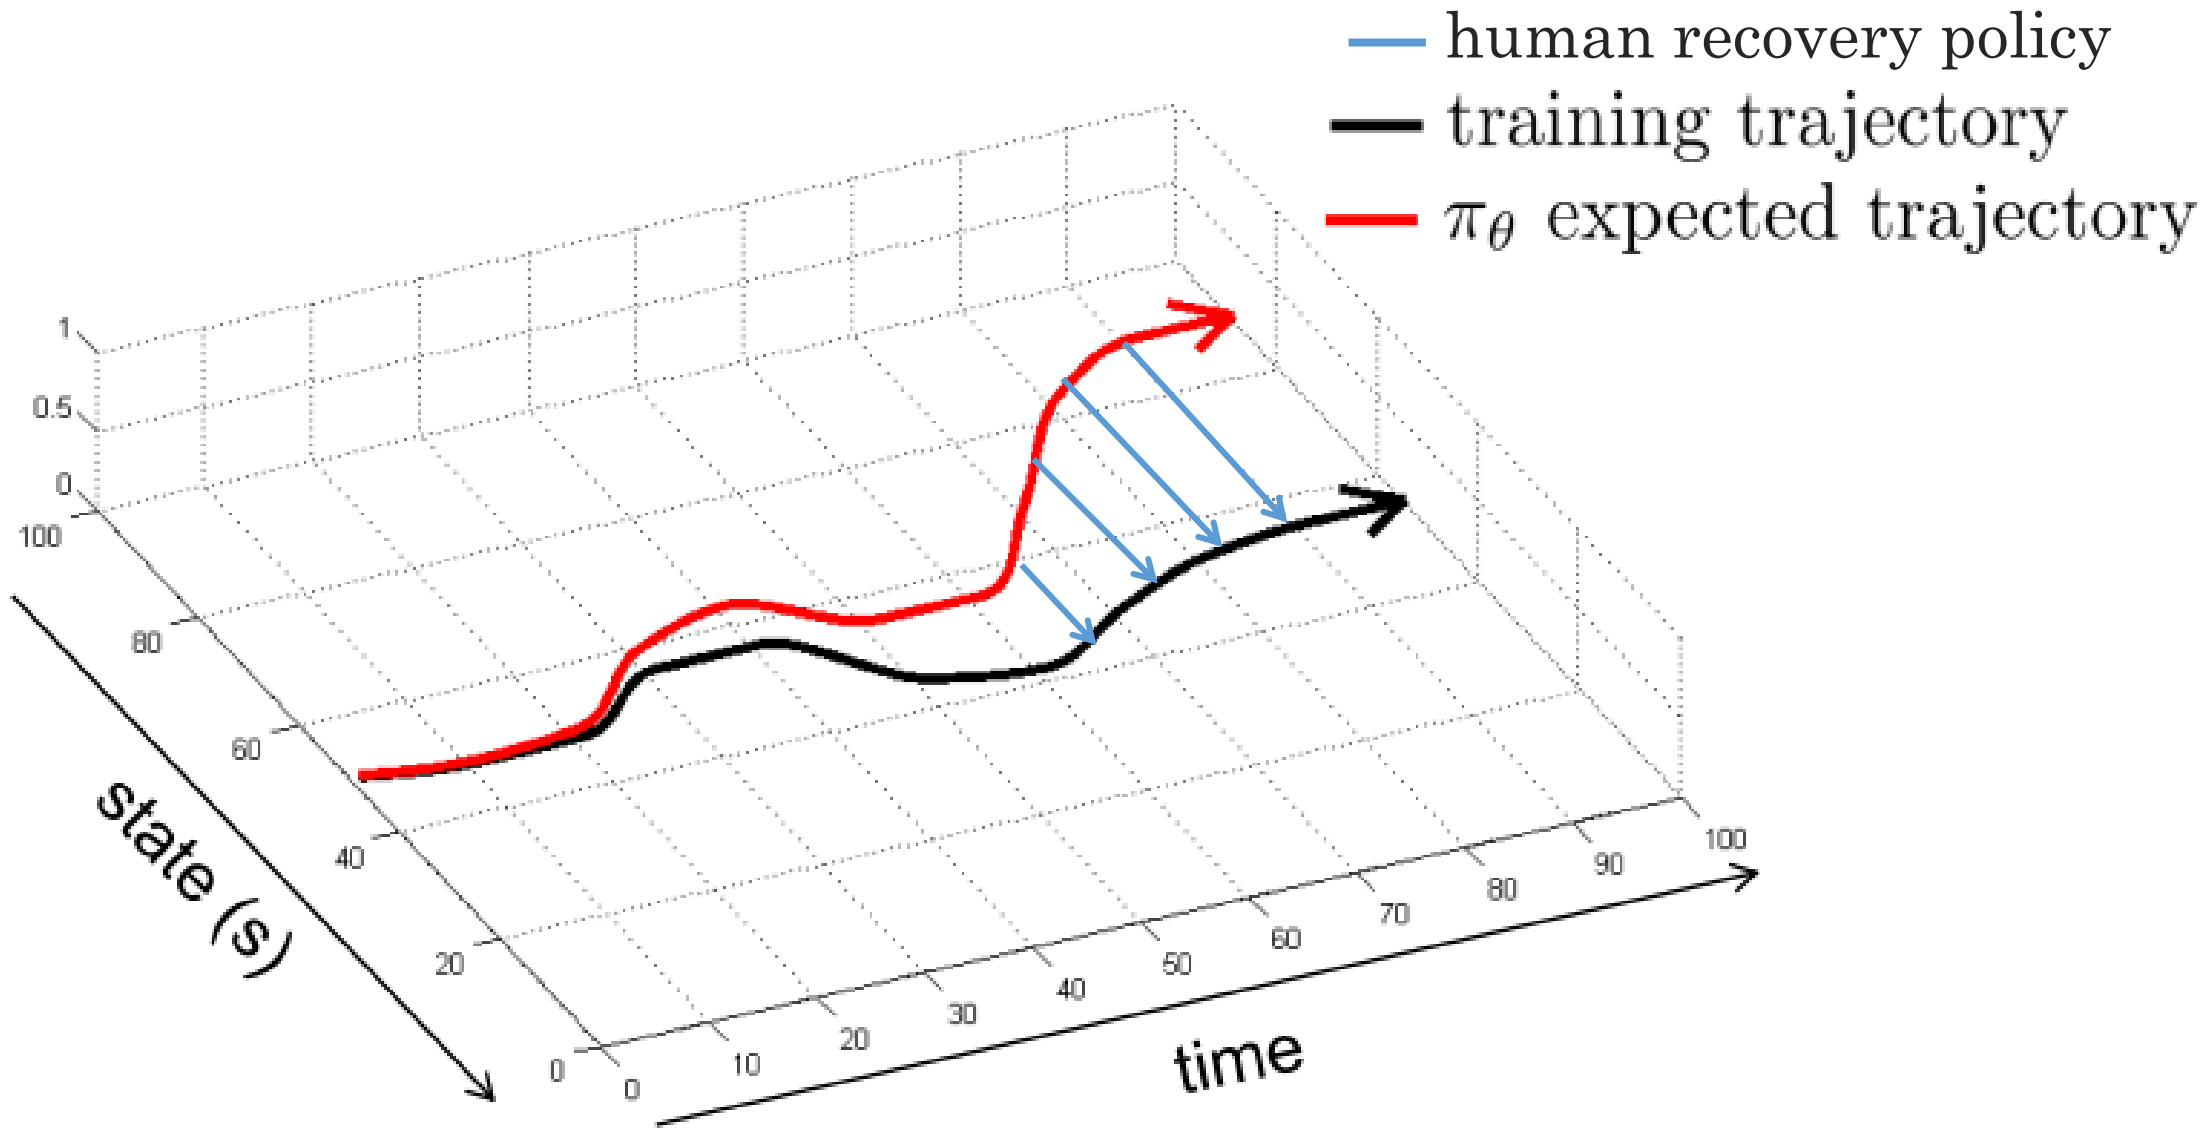
human recovery policy
training trajectory
π_θ expected trajectory
1
0.5
0
100
80
60
40
20
0
state (s)
0
10
20
30
40
50
60
70
80
90
100
time

DAgger

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

but need labels $\mathbf{a}_t$!

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \ldots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \ldots, \mathbf{o}_M\}$
3. Ask human to label $\mathcal{D}_\pi$ with actions $\mathbf{a}_t$
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Ross et al. ‘11

DAgger has very nice theoretical guarantees.

Why might it be **hard** to implement in practice?

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

but need labels $\mathbf{a}_t$!

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \ldots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \ldots, \mathbf{o}_M\}$
3. Ask human to label $\mathcal{D}_\pi$ with actions $\mathbf{a}_t$
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Ross et al. ‘11

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \ldots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \ldots, \mathbf{o}_M\}$
3. Ask human to label $\mathcal{D}_\pi$ with actions $\mathbf{a}_t$
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

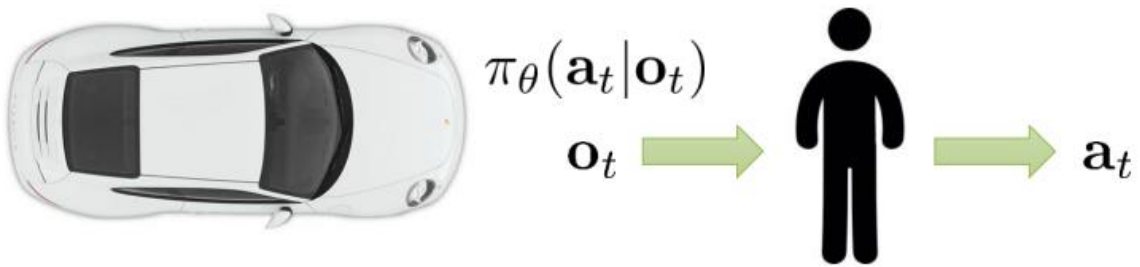

DAgger has very nice theoretical guarantees.

Why might it be **easy** to implement in practice?

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$

but need labels $\mathbf{a}_t$!

1. train $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{a}_1, \ldots, \mathbf{o}_N, \mathbf{a}_N\}$
2. run $\pi_\theta(\mathbf{a}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \ldots, \mathbf{o}_M\}$
3. Ask human to label $\mathcal{D}_\pi$ with actions $\mathbf{a}_t$
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Ross et al. ‘11

Learn from an Algorithmic Supervisor!

Seita et al. 2020. “Deep Imitation Learning of Sequential Fabric Smoothing From an Algorithmic Supervisor”

But we don’t always have access to an algorithmic supervisor…

Can we make DAgger more practical when dealing with real human labeling?

PLUS ONE ROBOTICS

WAYMO

ZOOX

nimble

Phantom Auto

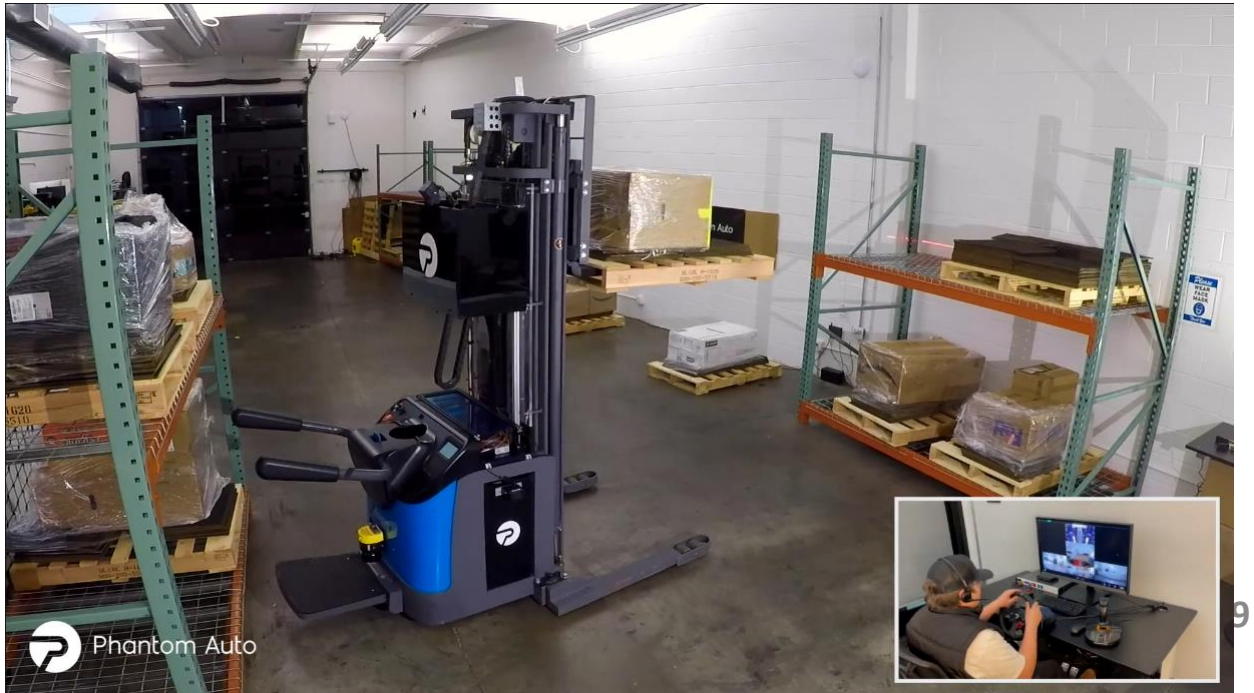

Interactive IL

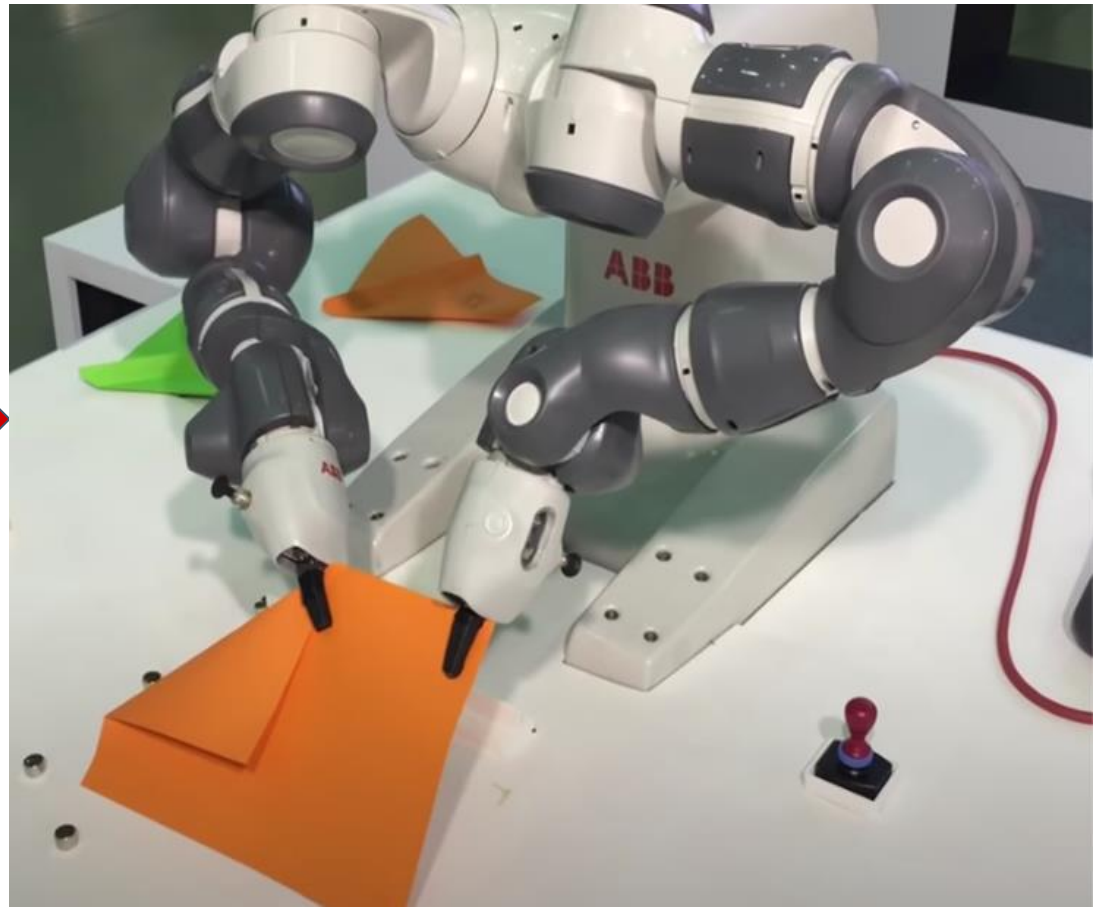

Interactive IL

$\pi_H(s)$

$\pi_{\text{meta}}(s)$

???

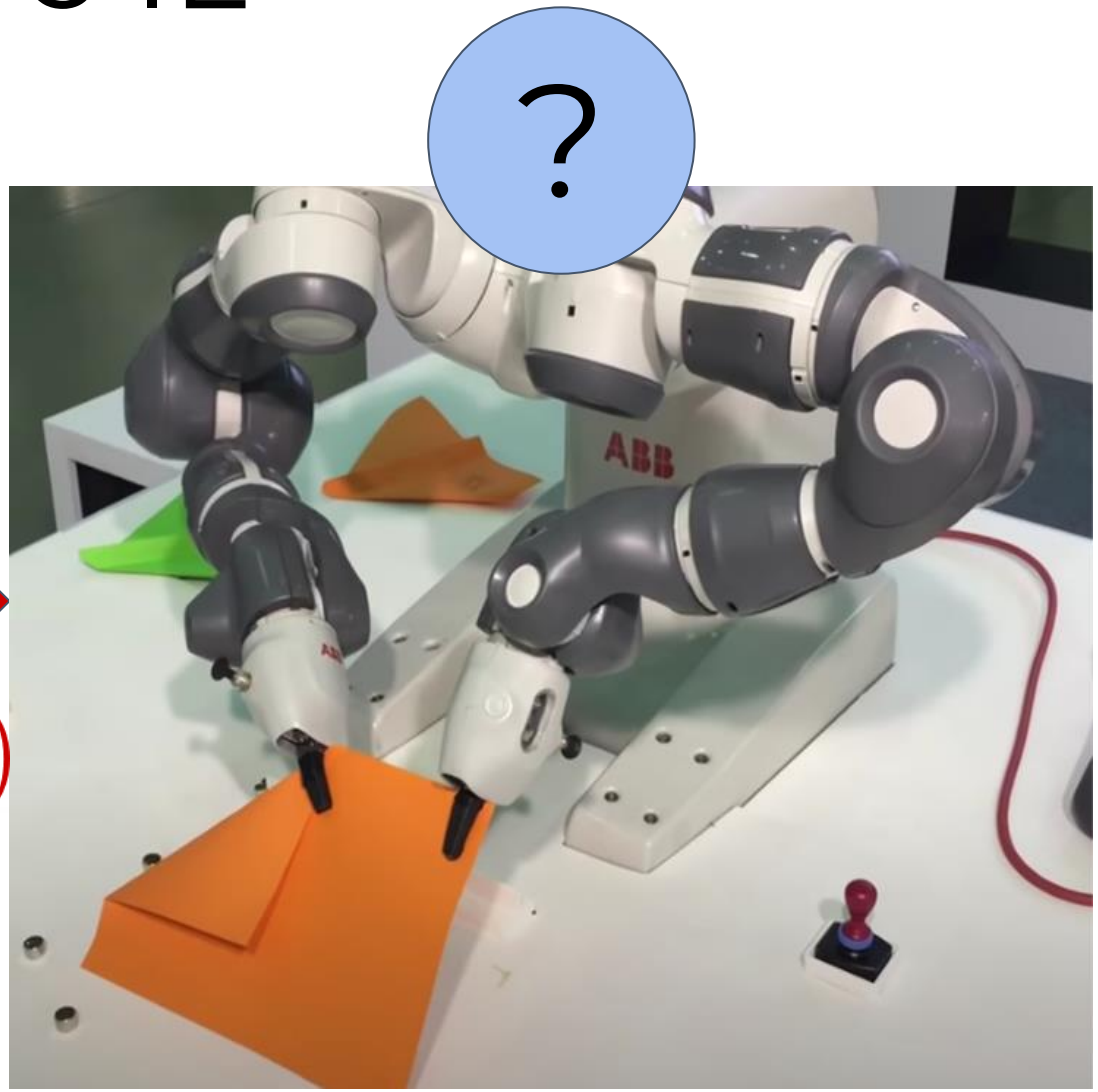

$\pi_R(s)$

Human-Gated Interactive IL

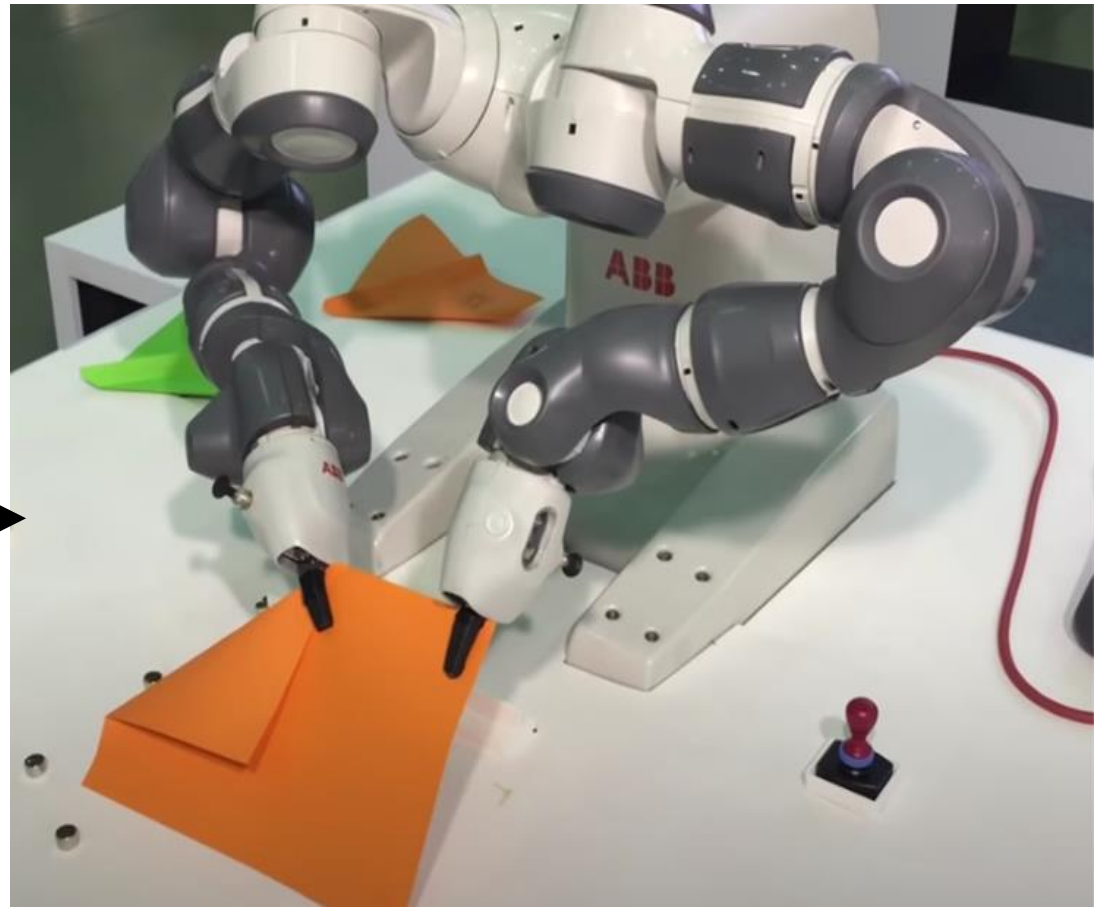

[3] M. Kelly, C. Sidrane, K. Driggs-Campbell, and M. J. Kochenderfer. HG-DAgger: Interactive Imitation Learning with Human Experts. ICRA 2019.

Human-Gated Interactive IL

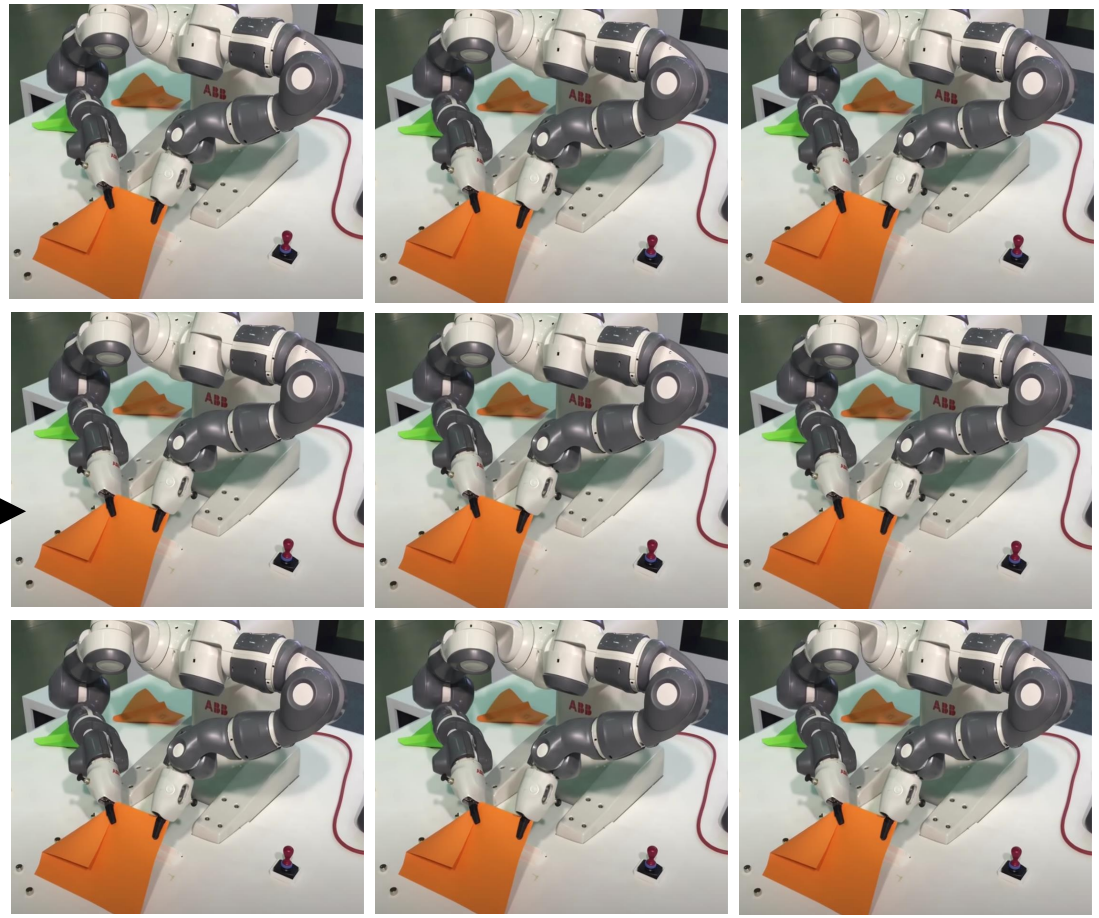

[3] M. Kelly, C. Sidrane, K. Driggs-Campbell, and M. J. Kochenderfer. HG-DAgger: Interactive Imitation Learning with Human Experts. ICRA 2019.

Robot-Gated Interactive IL

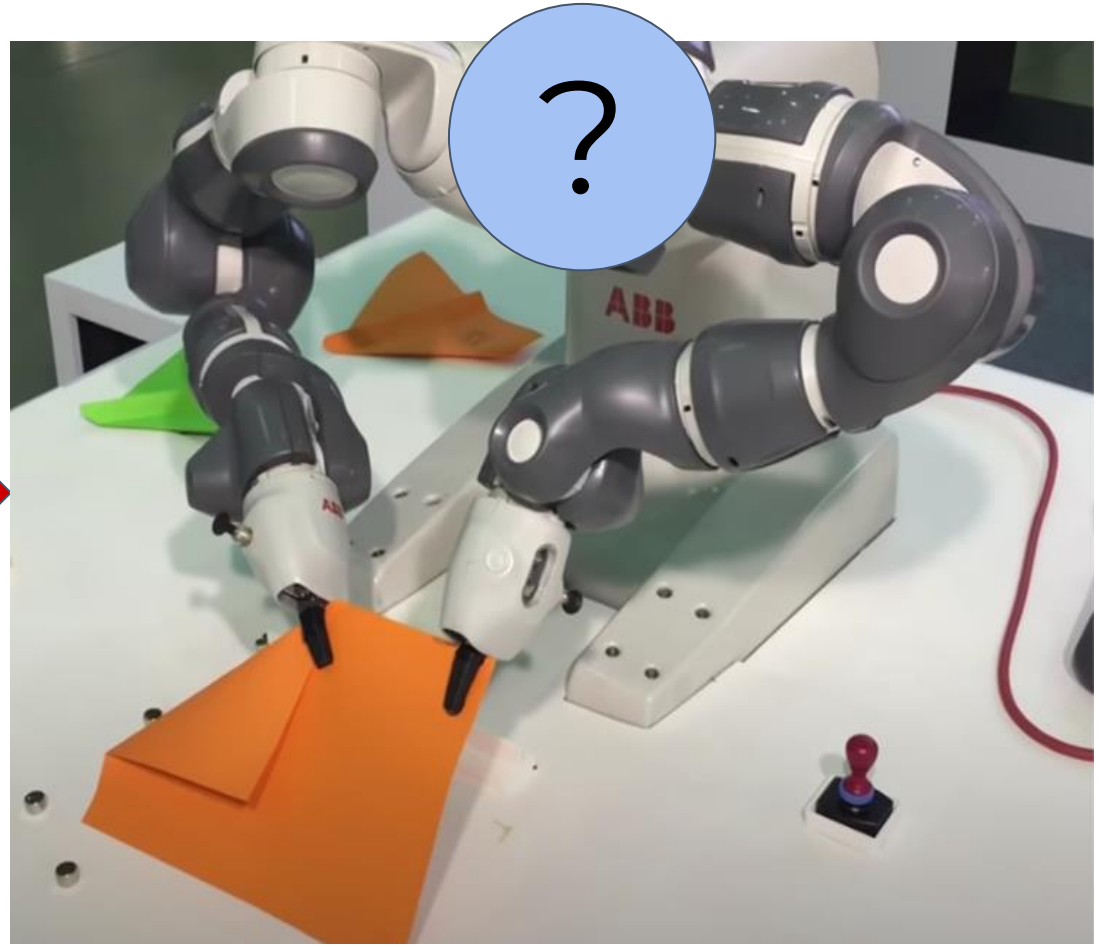

[4] J. Zhang, K. Cho. Query-Efficient Imitation Learning for End-to-End Autonomous Driving. AAAI 2017.
[5] K. Menda, K. Driggs-Campbell, M. Kochenderfer. EnsembleDAgger: A Bayesian Approach to Safe Imitation Learning. IROS 2019.

Minimizing Supervisor Burden

- C = Number of context switches
- L = Latency of context switching
- I = Expected number of supervisor actions per intervention

$$B(\pi) \triangleq C(\pi) \cdot (L + I(\pi))$$

Ideally, we want

$$\pi = \arg\min_{\pi' \in \Pi} L(\pi'_r)$$
$$s.t.\ B(\pi') \leq \Gamma_b$$

Minimizing Supervisor Burden

- C = number of context switches
- L = Latency of context switching
- I = expected number of supervisor actions per intervention

$$B(\pi) \triangleq \boldsymbol{C}(\pi) \cdot (L + I(\pi))$$

In practice, we approximate this by focusing on limiting the number of interventions (number of context switches)

$$\pi = \arg\min_{\pi' \in \Pi} L(\pi'_r)$$
$$s.t.\ B(\pi') \leq \Gamma_b$$

SafeDAgger

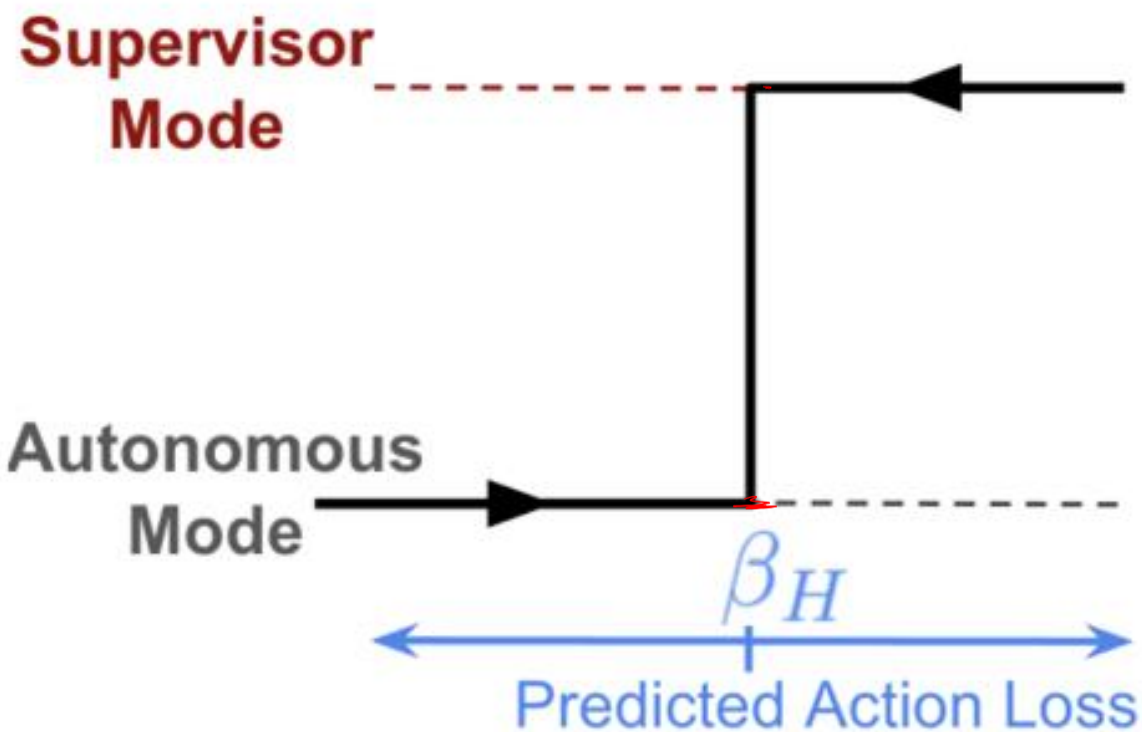

Predicted action loss = predicted difference between human and robot action.

Trained using held-out set of data from human.

J. Zhang, K. Cho. Query-Efficient Imitation Learning for End-to-End Autonomous Driving. AAAI 2017.

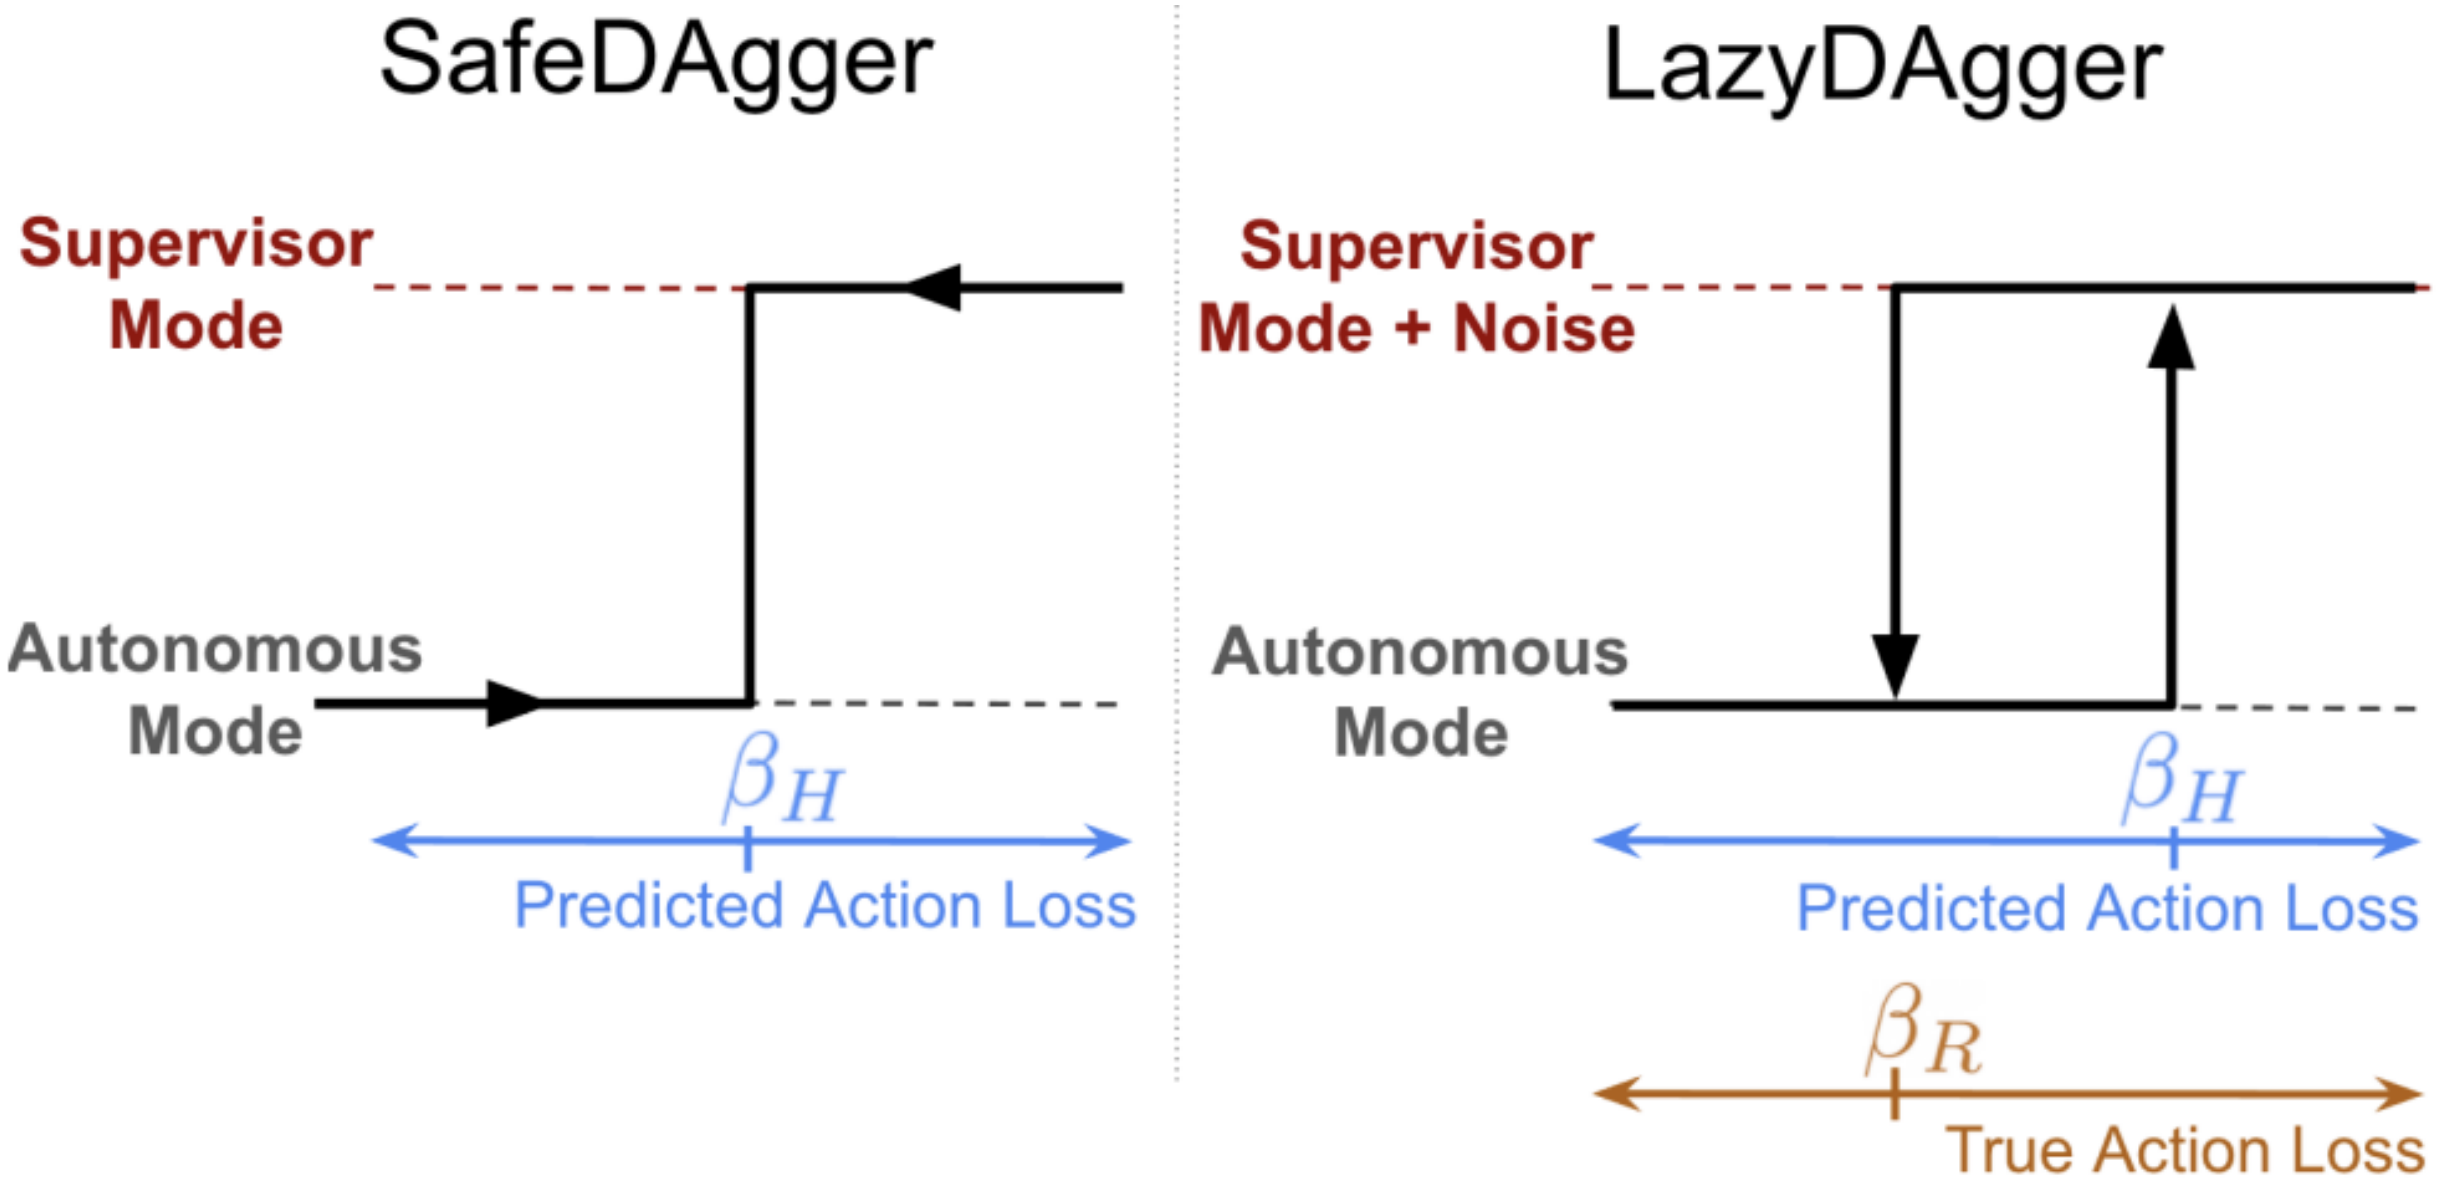

Hoque et al. 2021.

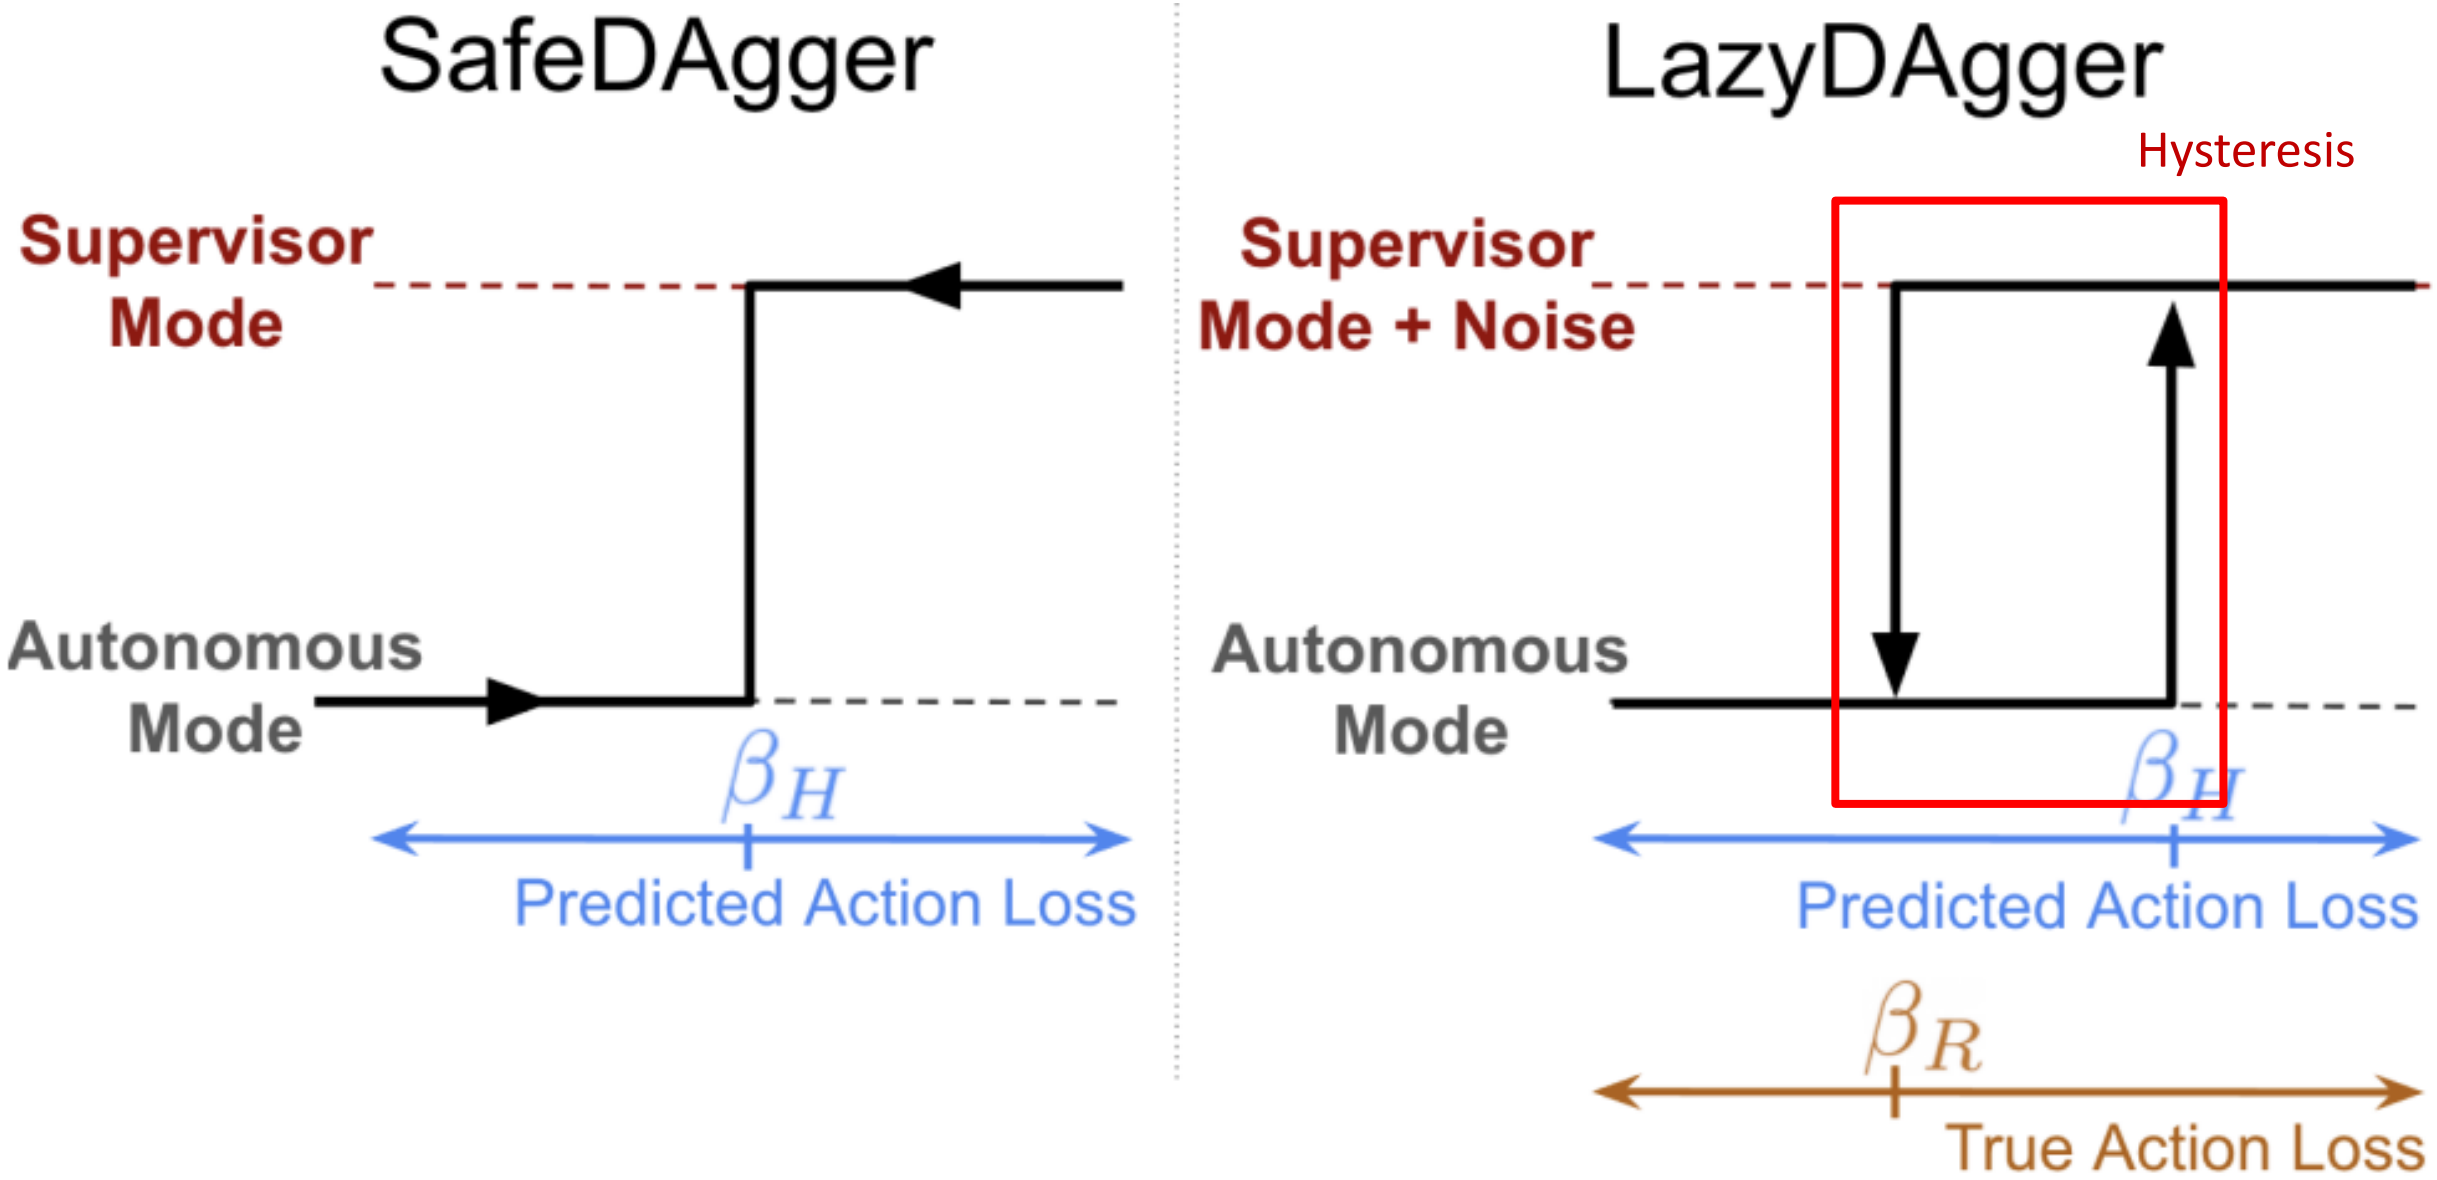

Hoque et al. 2021.

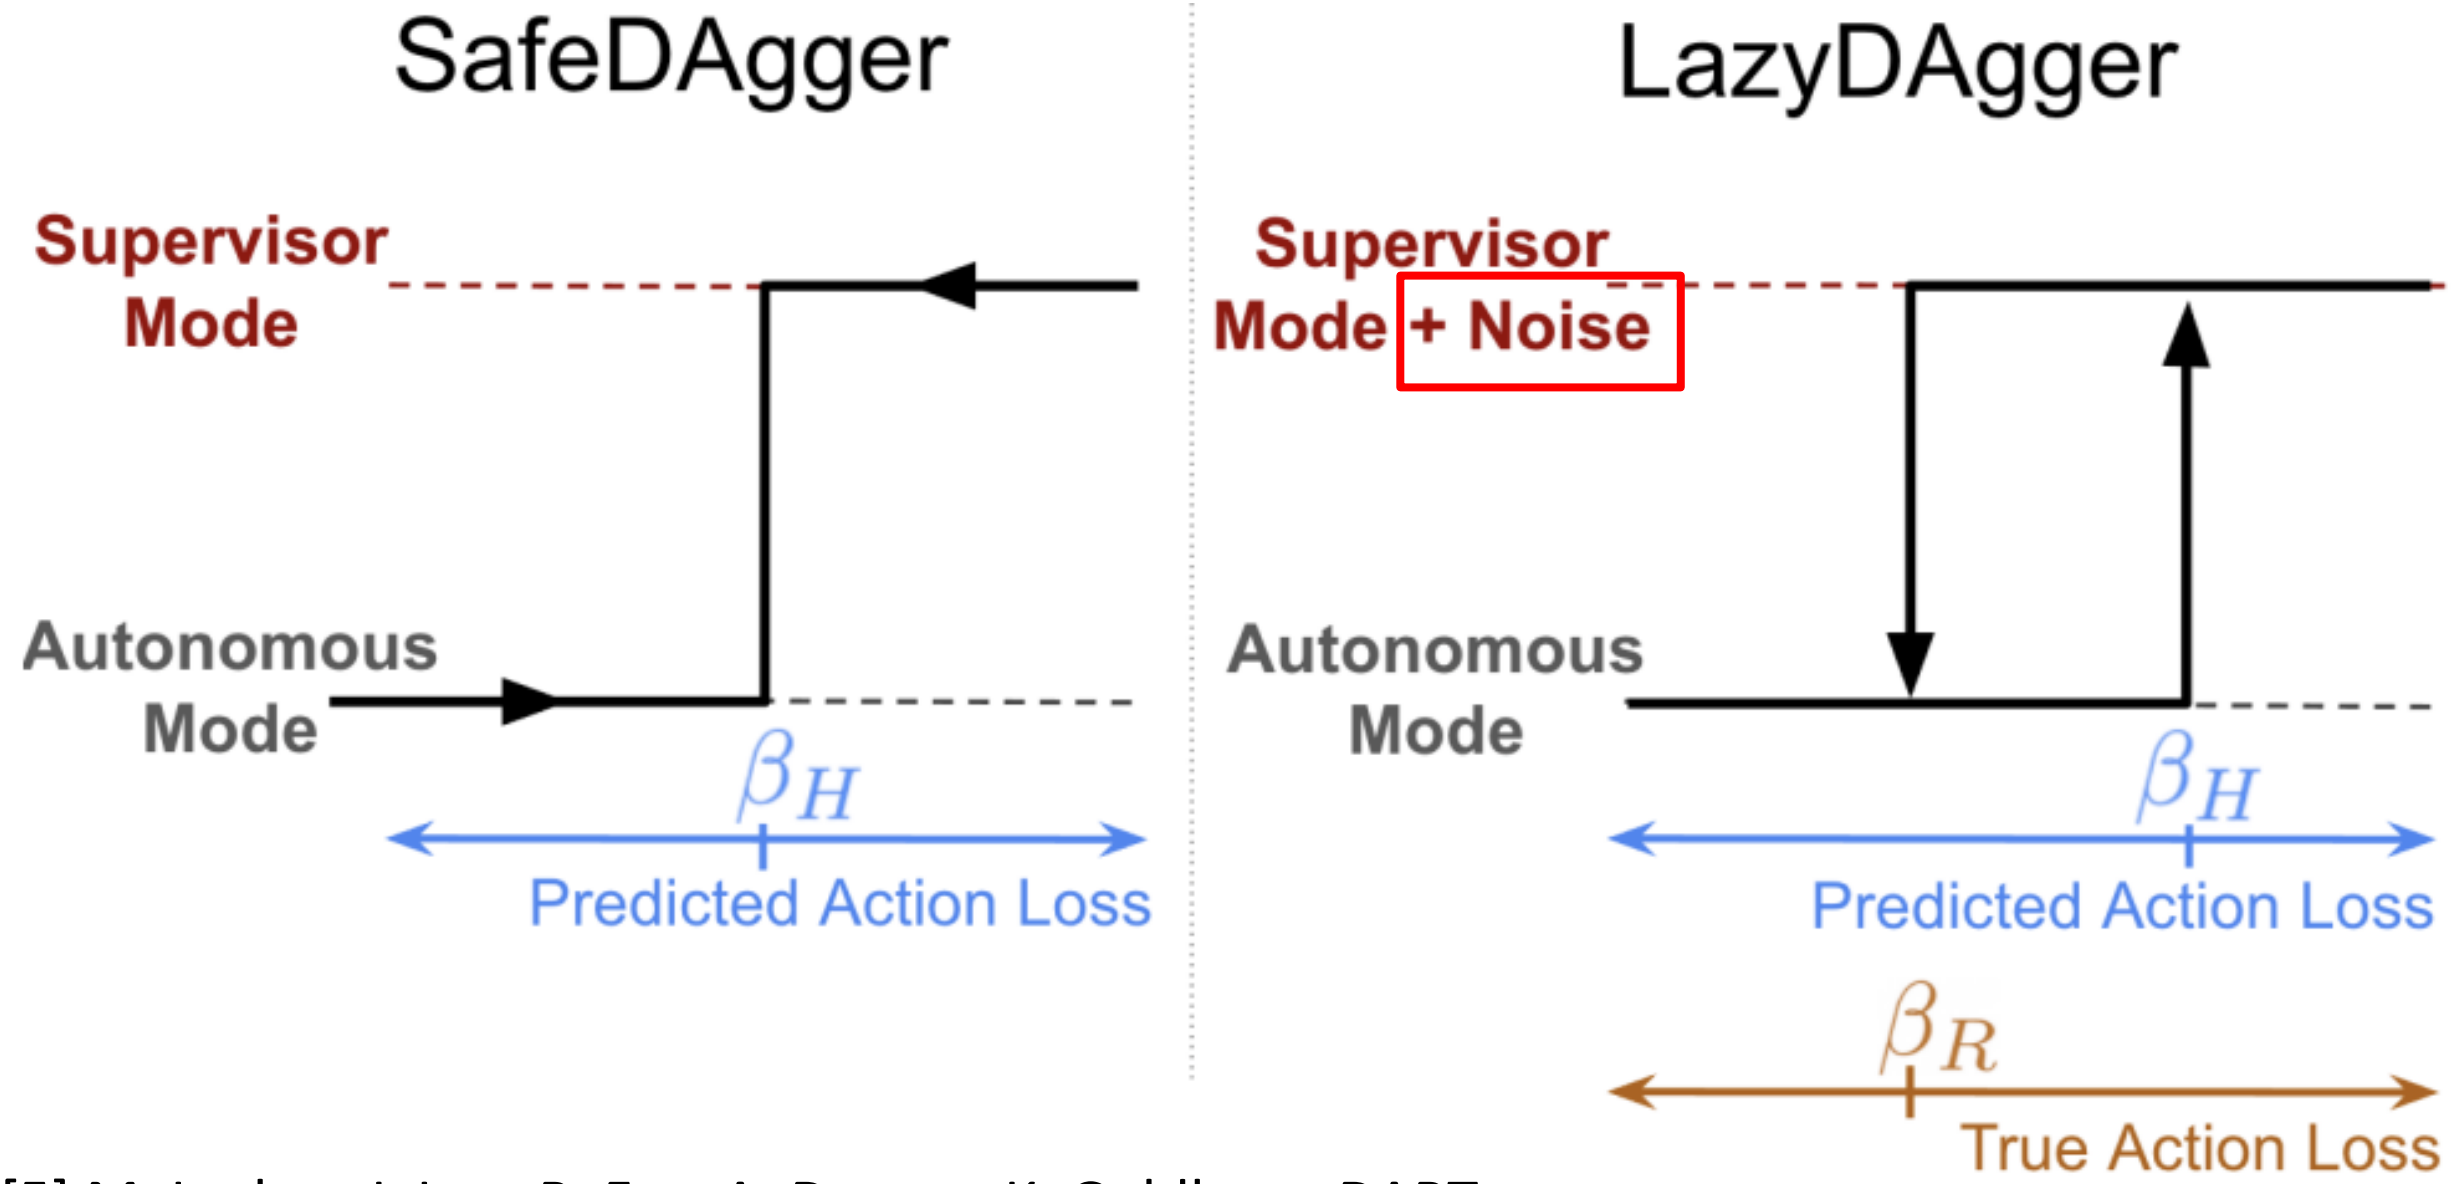

[5] M. Laskey, J. Lee, R. Fox, A. Dragan, K. Goldberg. DART: Noise Injection for Robust Imitation Learning. CoRL 2017.

LazyDAgger

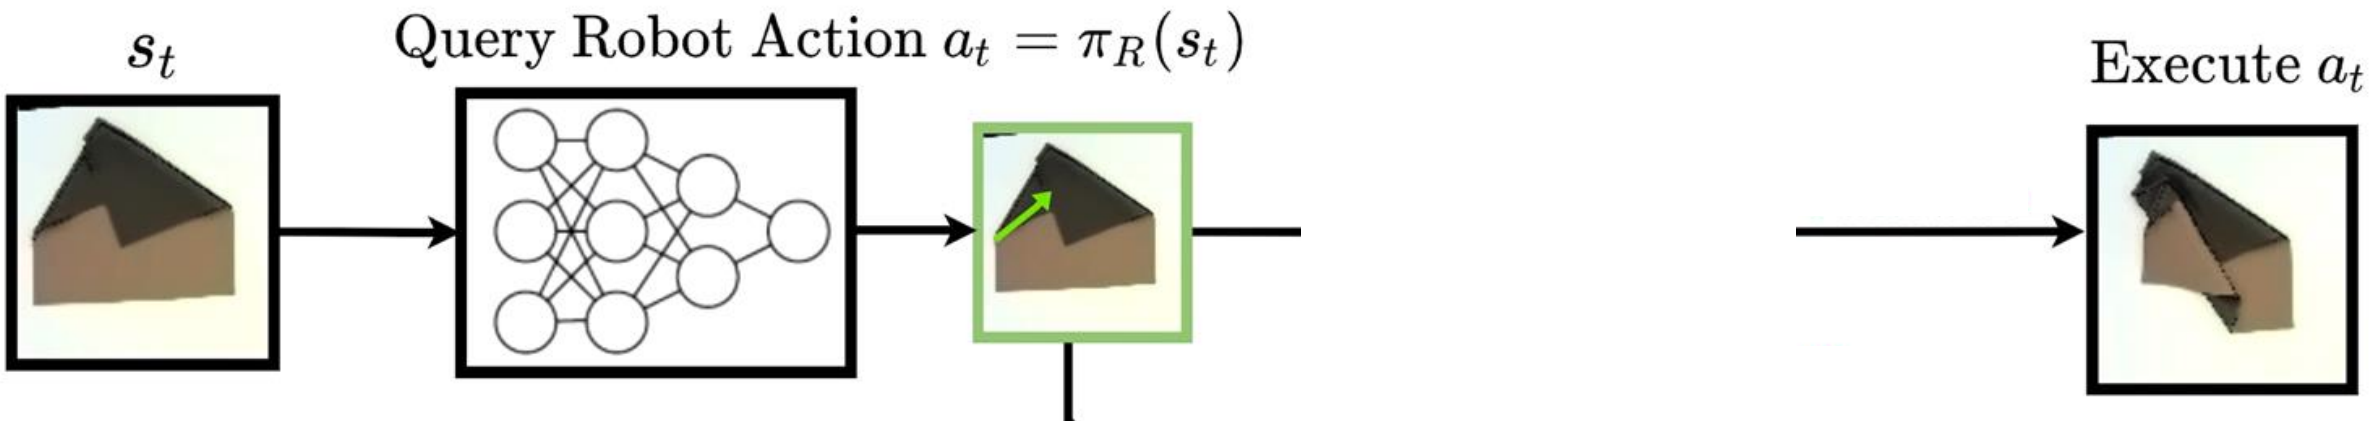

Hoque et al. 2021.

LazyDAgger

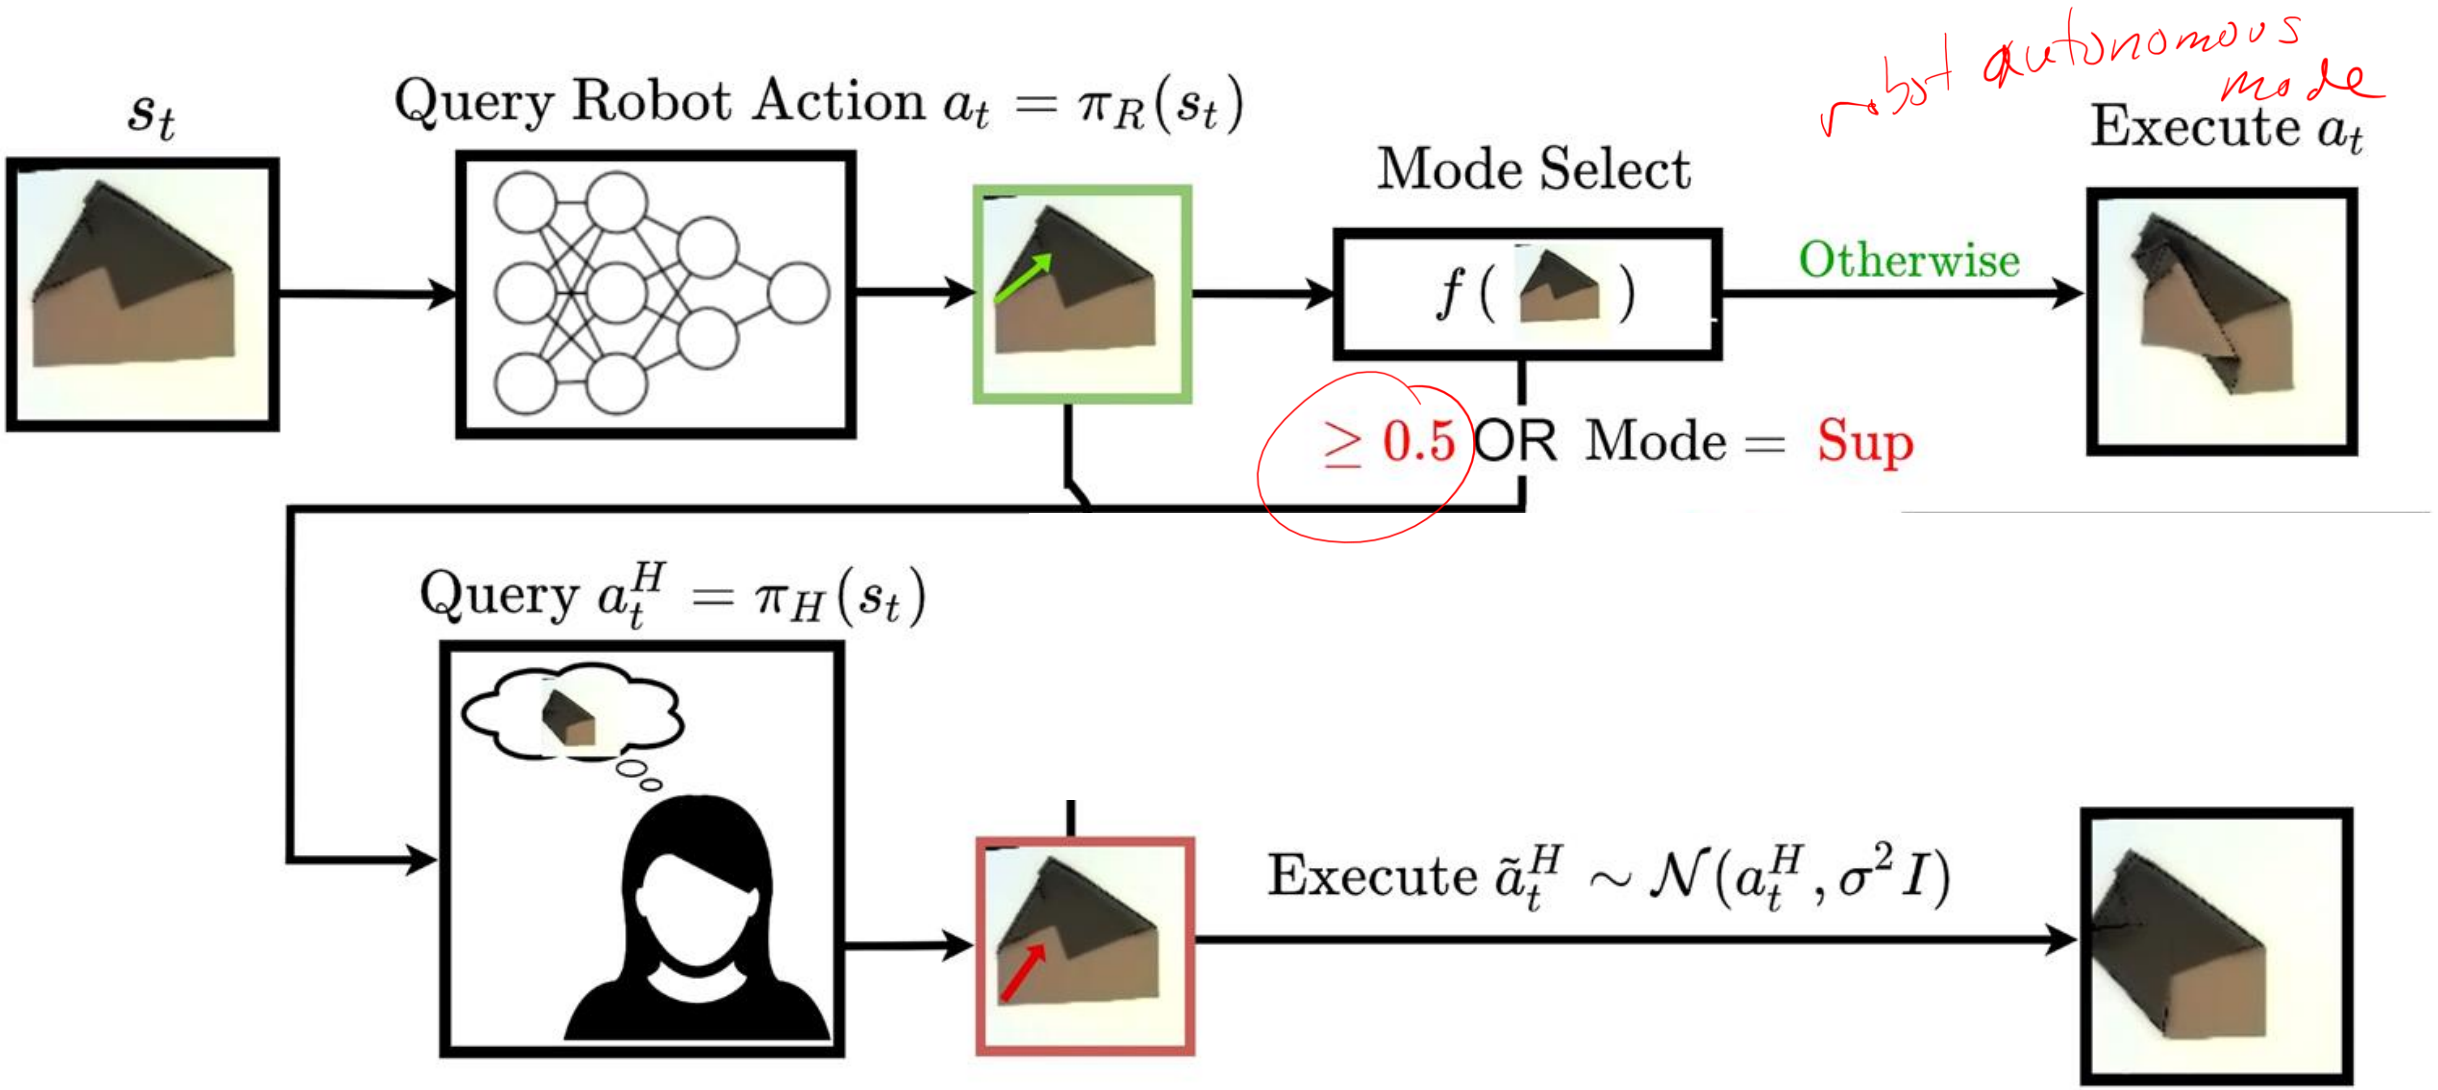
s_t
Query Robot Action $a_t = \pi_R(s_t)$
Mode Select
$f(\ )$
Otherwise
Execute a_t
robot autonomous mode
≥ 0.5 OR Mode = Sup
Query $a_t^H = \pi_H(s_t)$
Execute $\tilde{a}_t^H \sim \mathcal{N}(a_t^H, \sigma^2 I)$

LazyDAgger

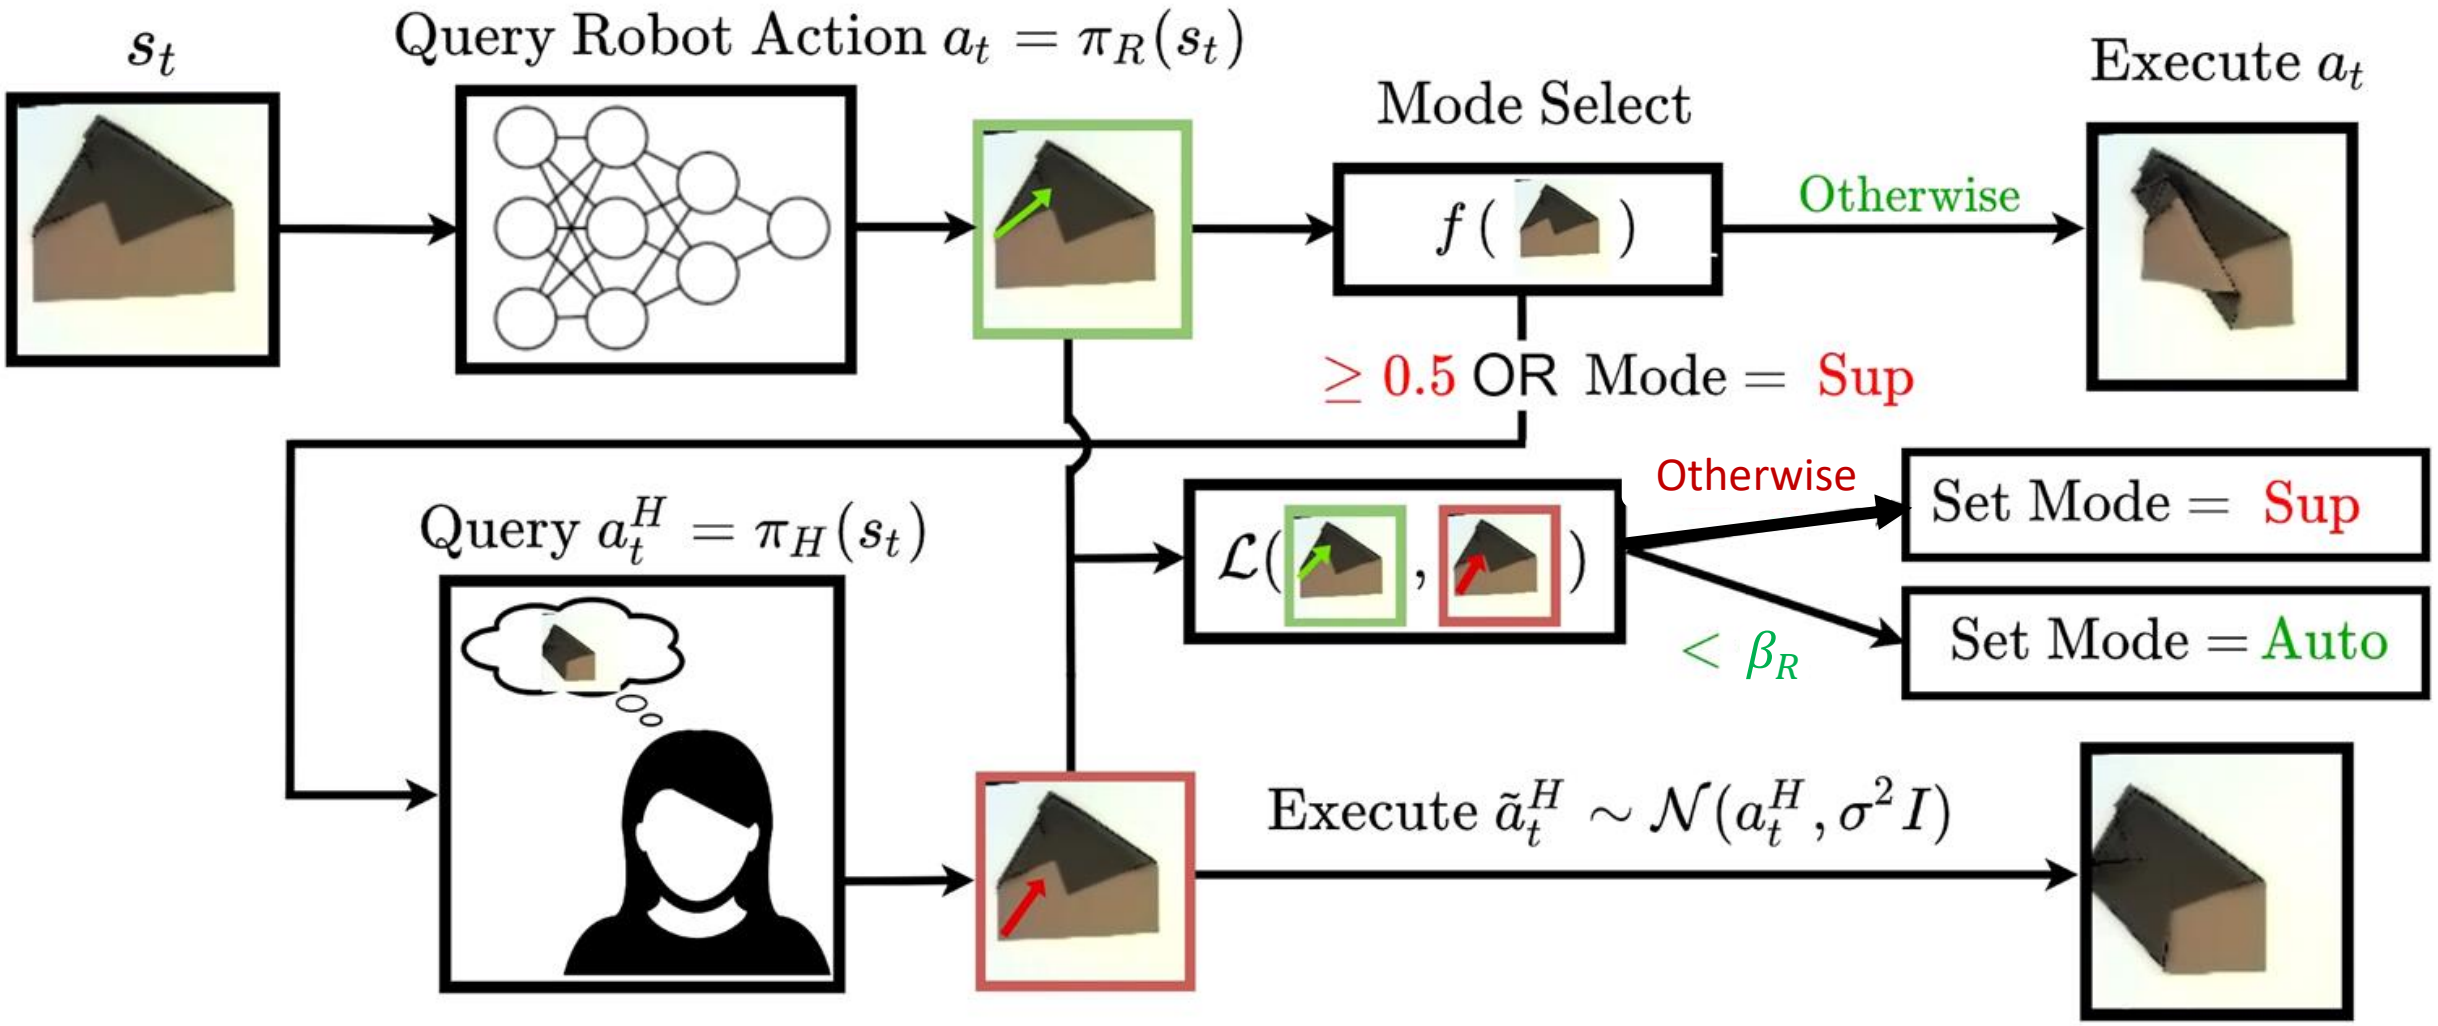
s_t
Query Robot Action $a_t = \pi_R(s_t)$
Mode Select
$f($ $)$
Otherwise
Execute a_t
≥ 0.5 OR Mode = Sup
Query $a_t^H = \pi_H(s_t)$
$\mathcal{L}($, $)$
Otherwise
Set Mode = Sup
$< \beta_R$
Set Mode = Auto
Execute $\tilde{a}_t^H \sim \mathcal{N}(a_t^H, \sigma^2 I)$

Simulation Experiments

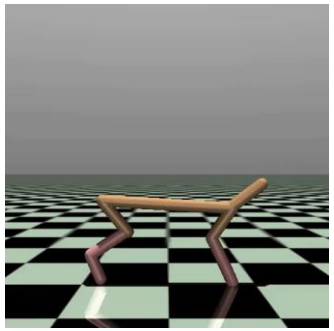

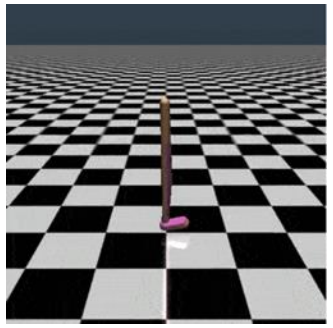

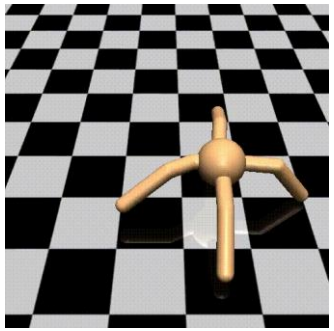

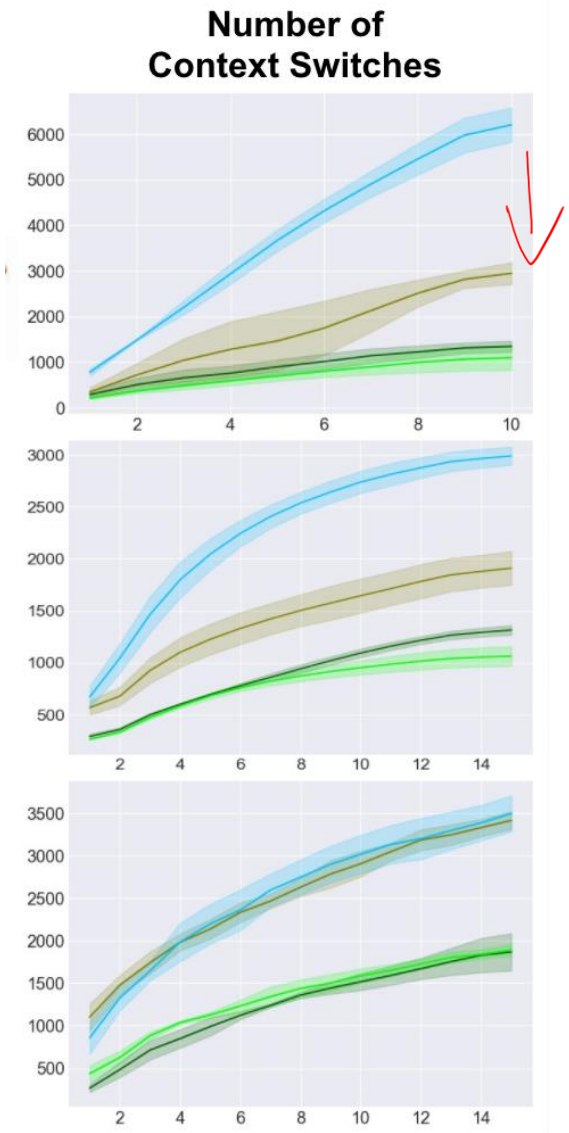

Context Switching Reduction

79%

56%

46%

Behavior Cloning DAgger SafeDAgger Ours Ours (- Noise) Ours (- Switch to Auto)

Simulation Experiments

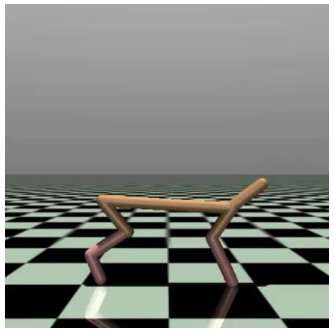

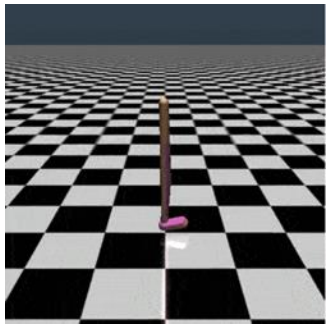

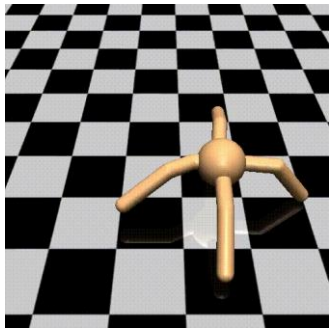

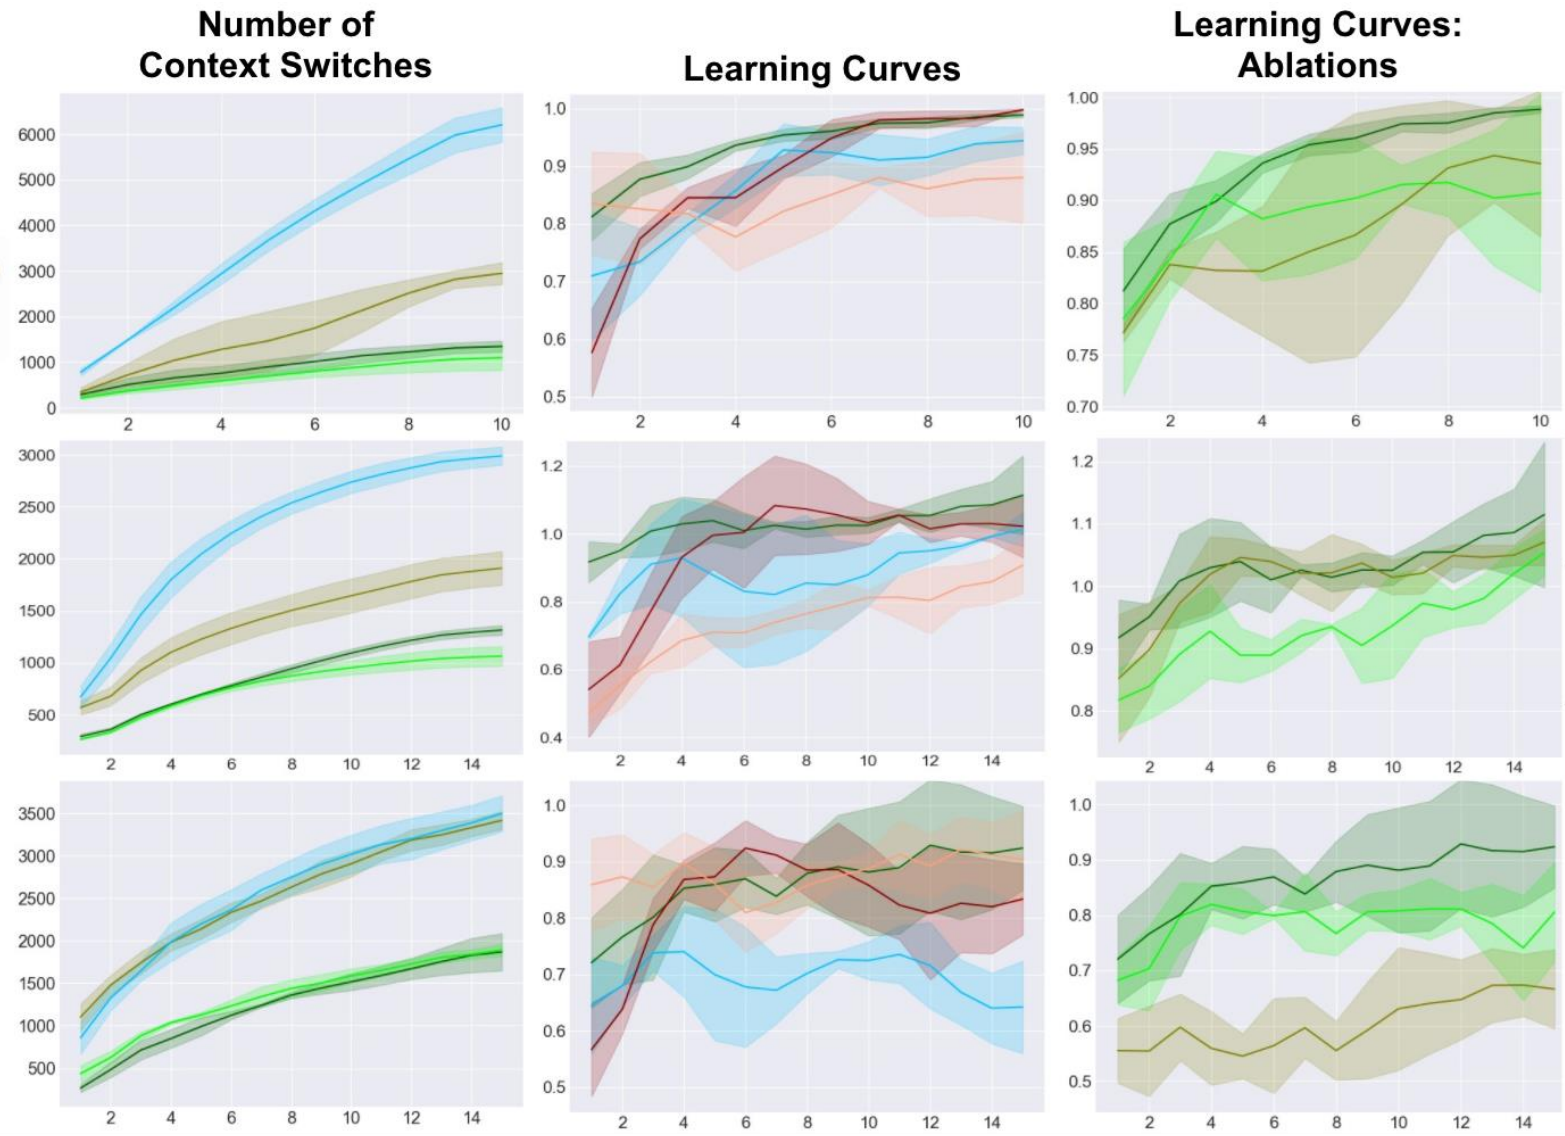

Simulation Experiments

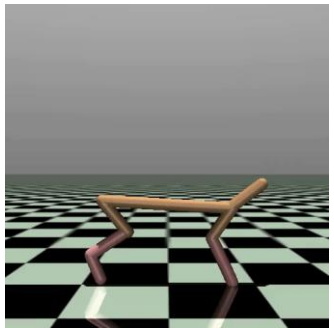

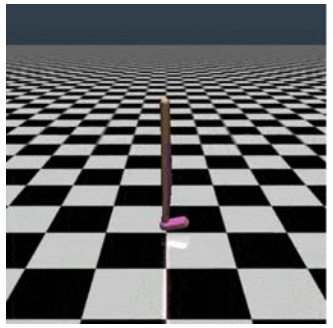

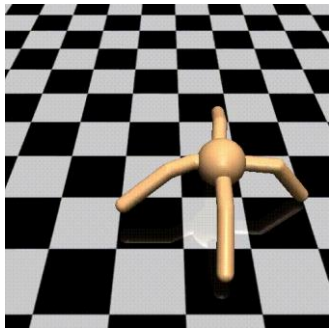

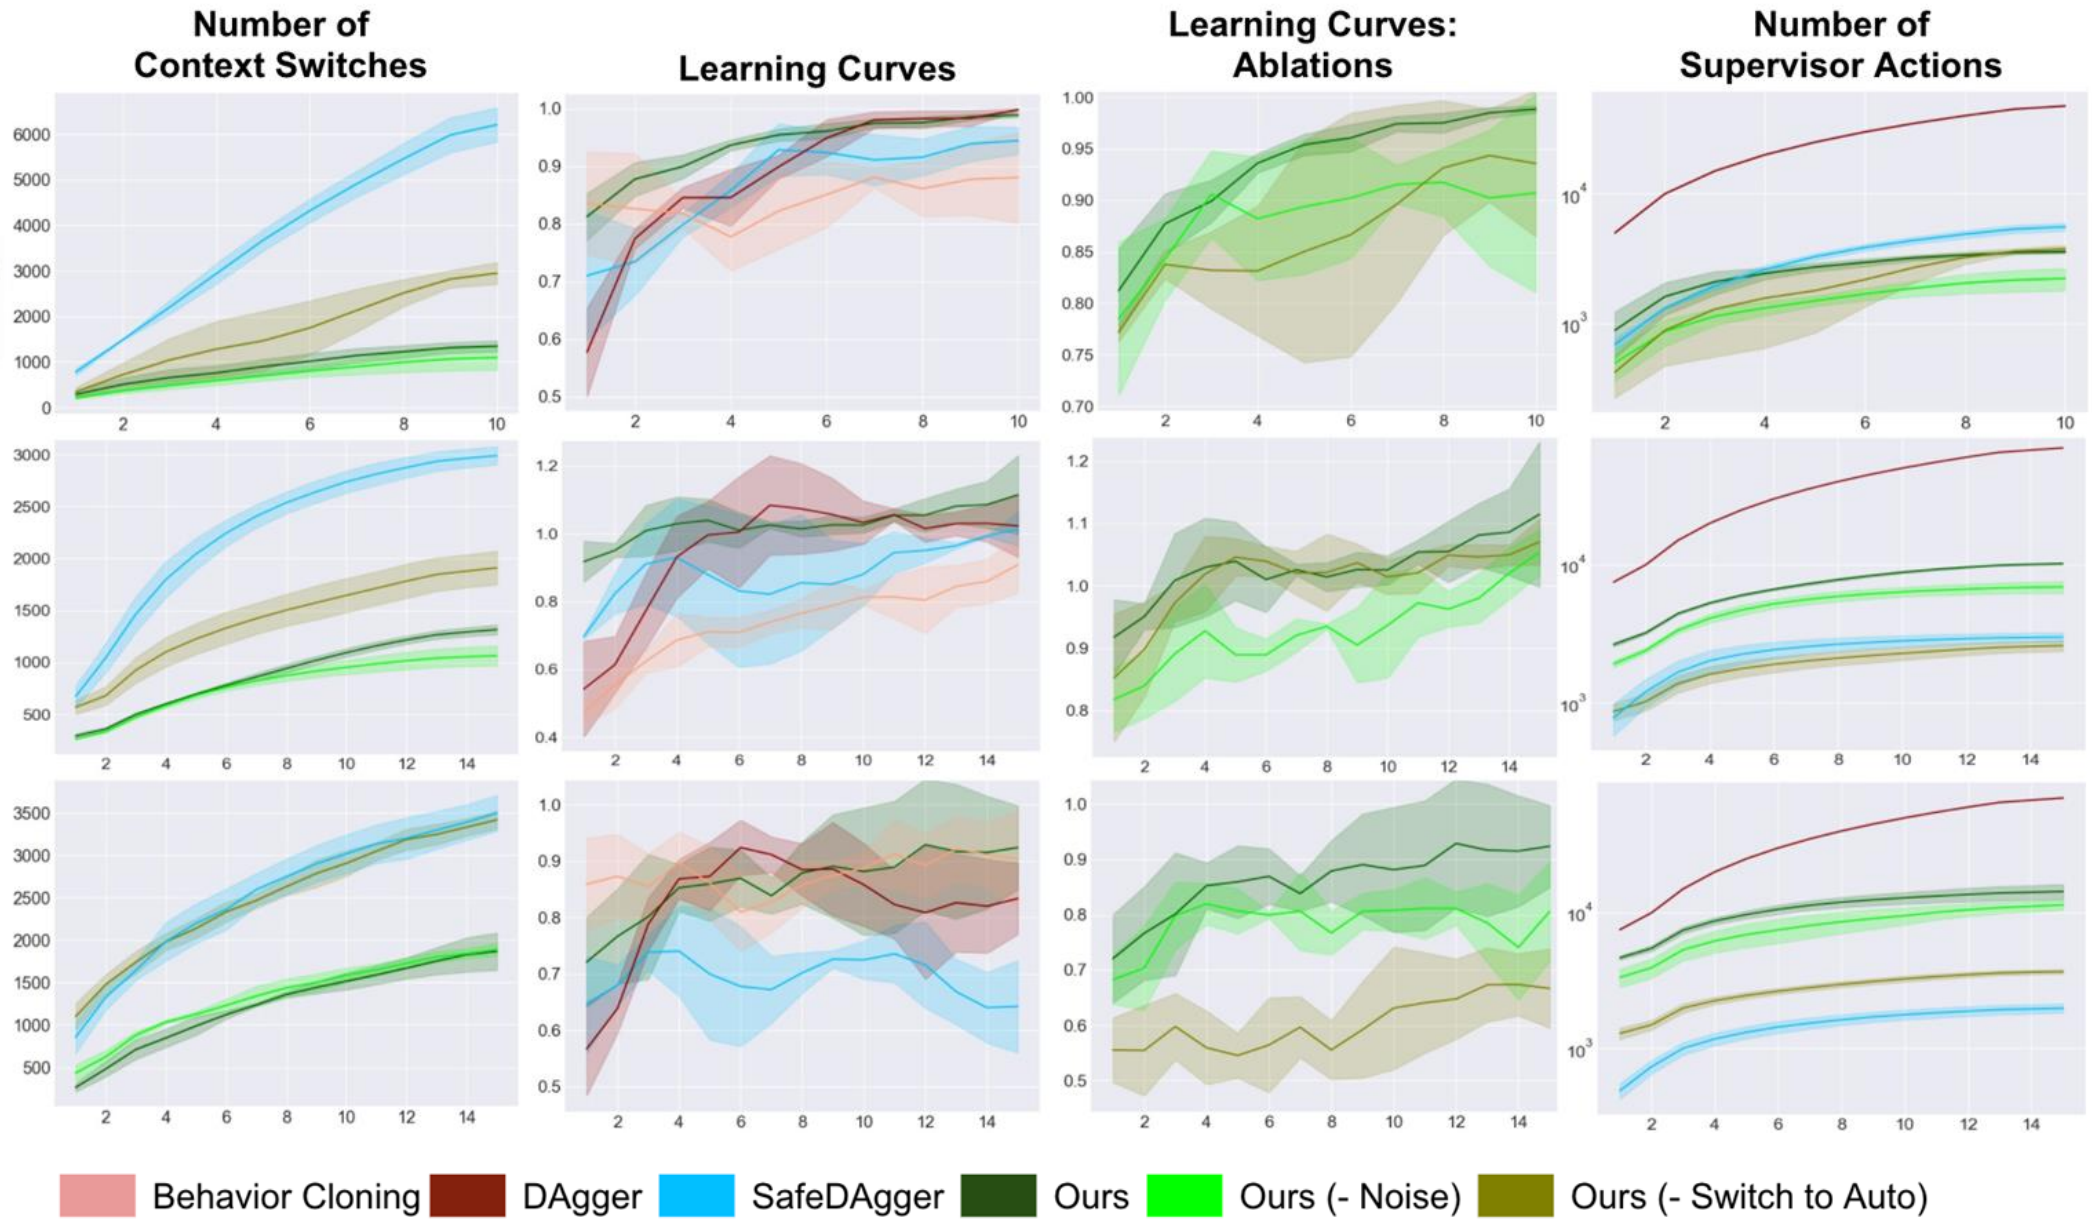

Hoque et al. 2021.

(1) SMOOTH

(2) ALIGN

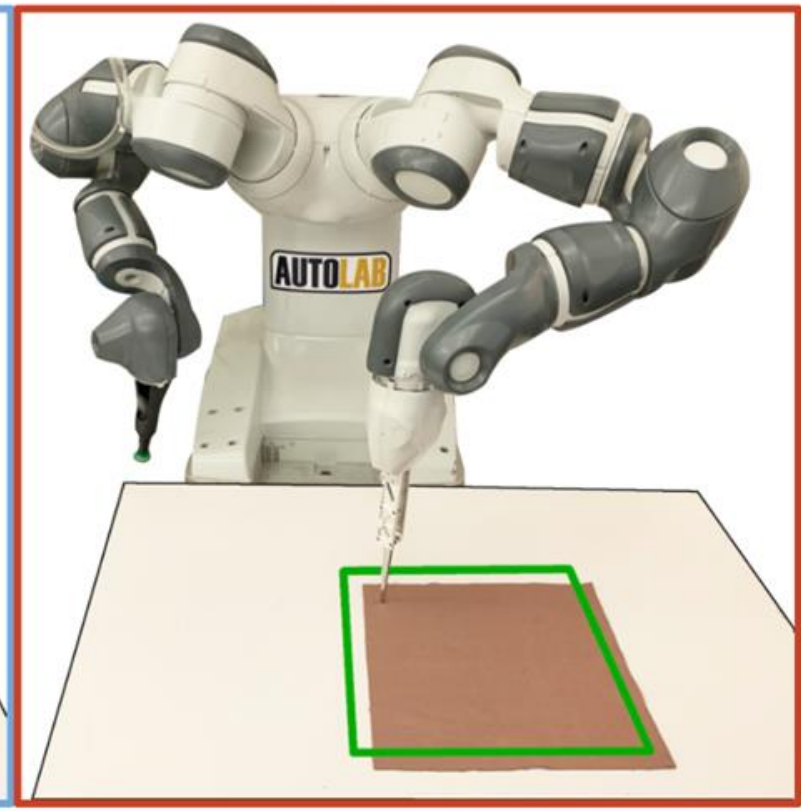

(3) FOLD

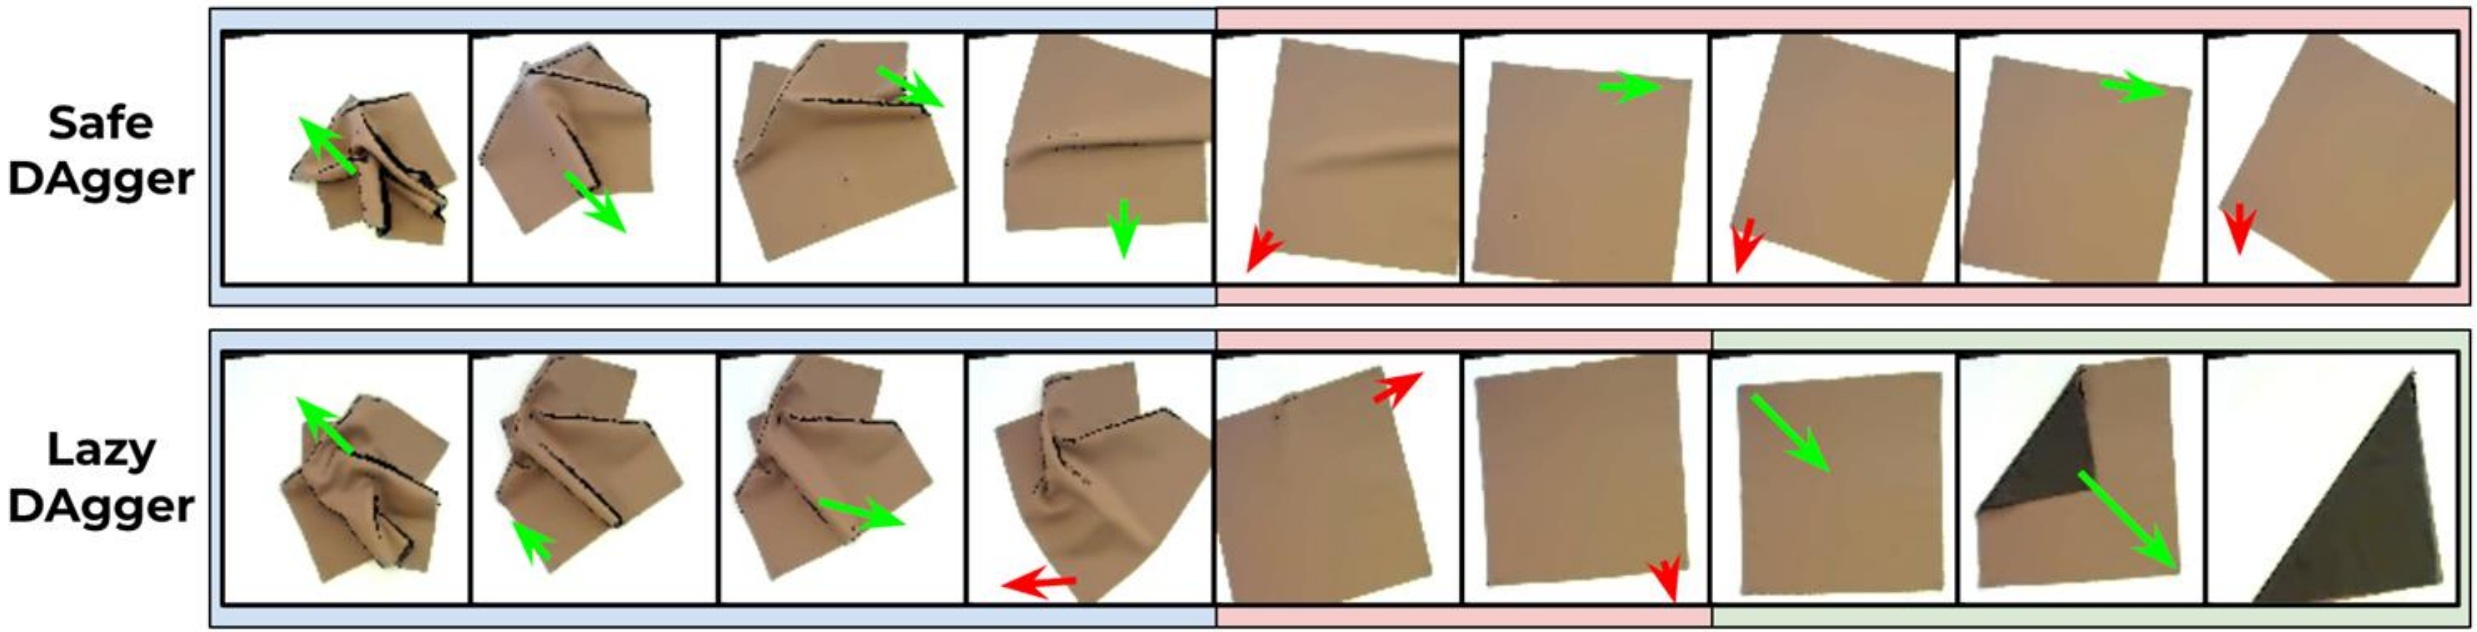
Safe
DAgger
Lazy
DAgger

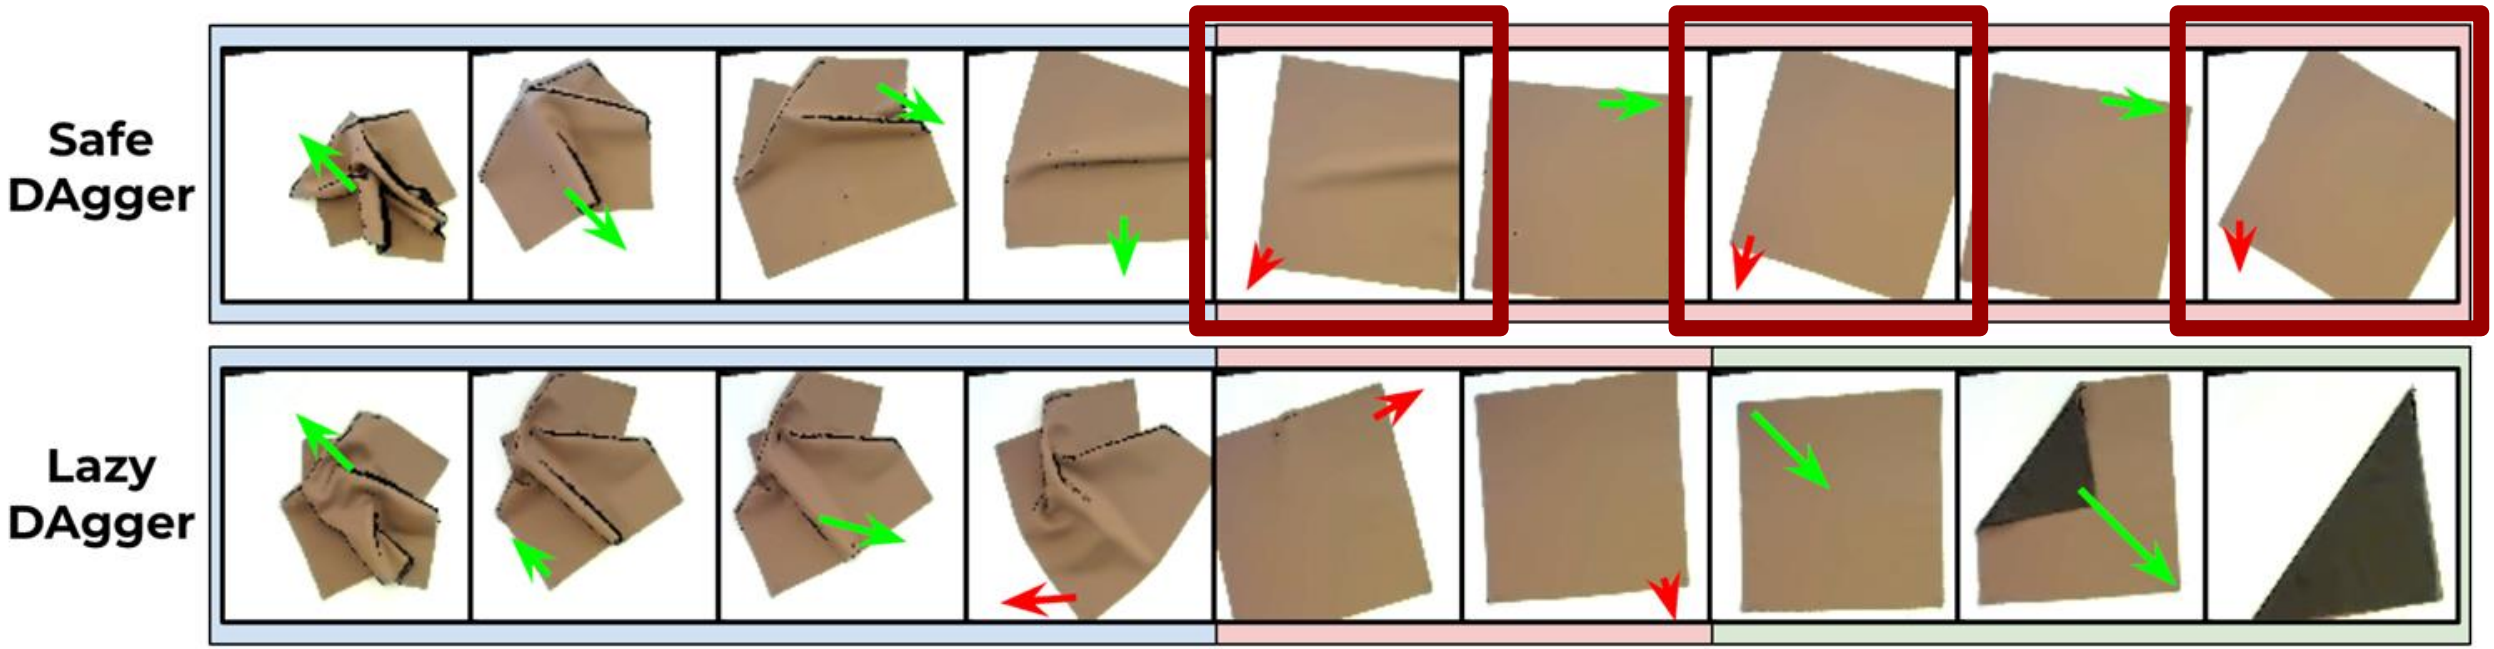
Safe
DAgger
Lazy
DAgger

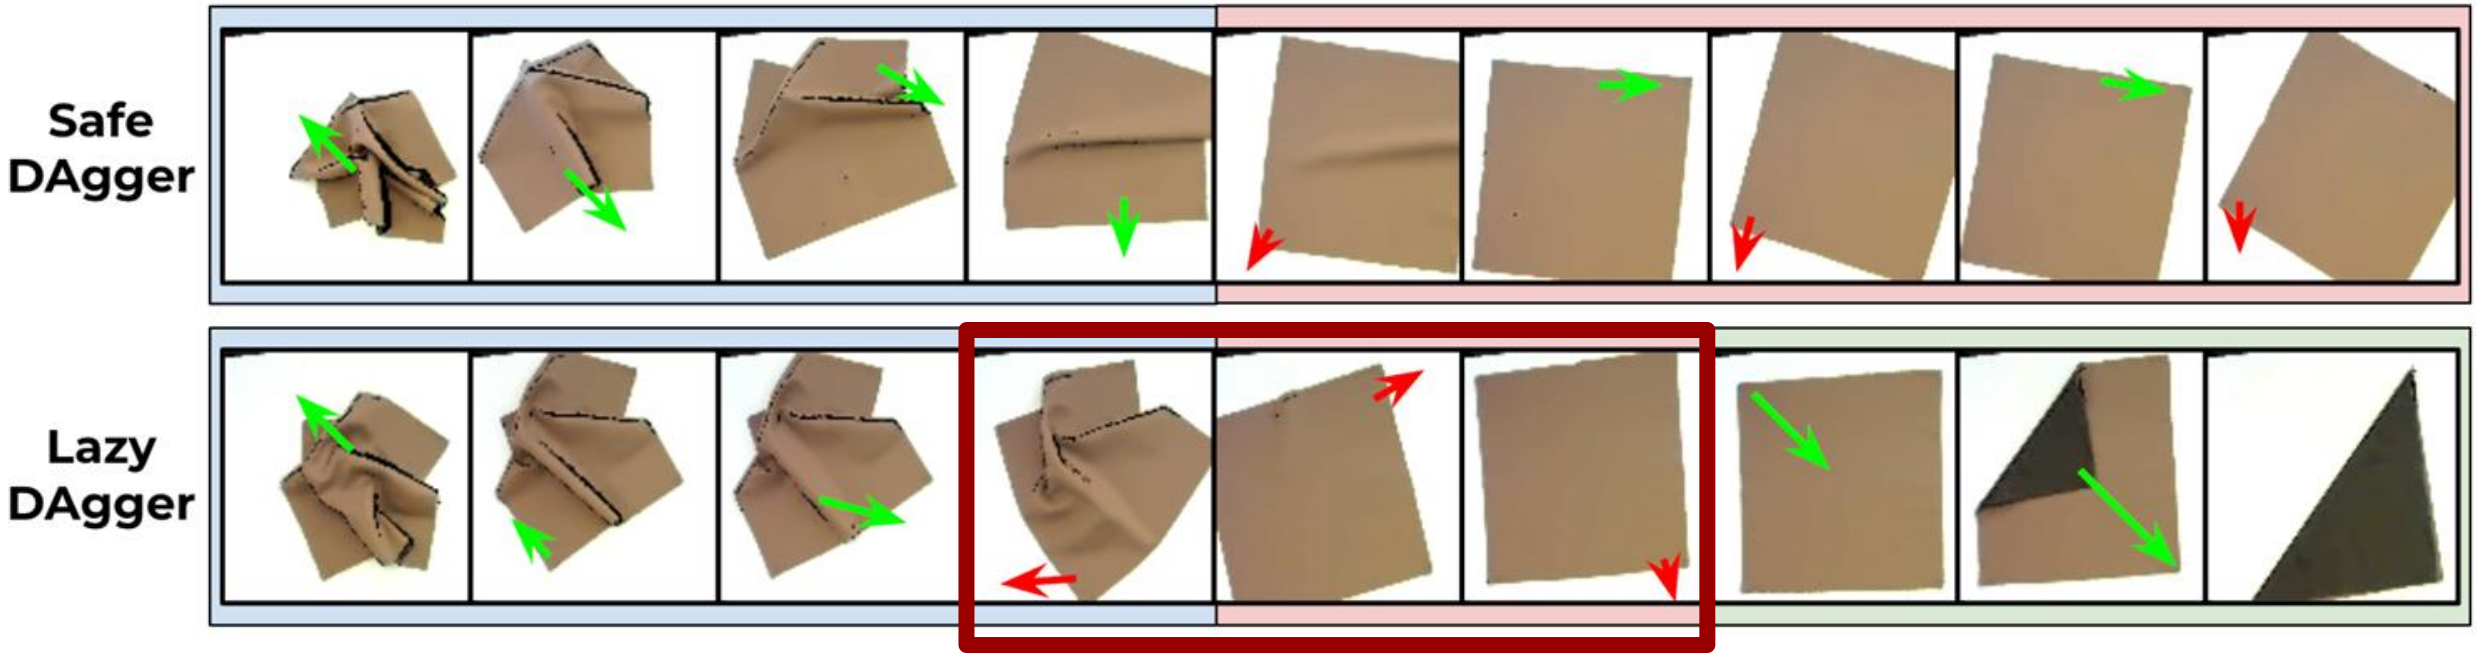
Safe
DAgger
Lazy
DAgger

Limitations

- Parameter tuning
- Hard to know how many interventions will be requested.
- One human managing one robot.

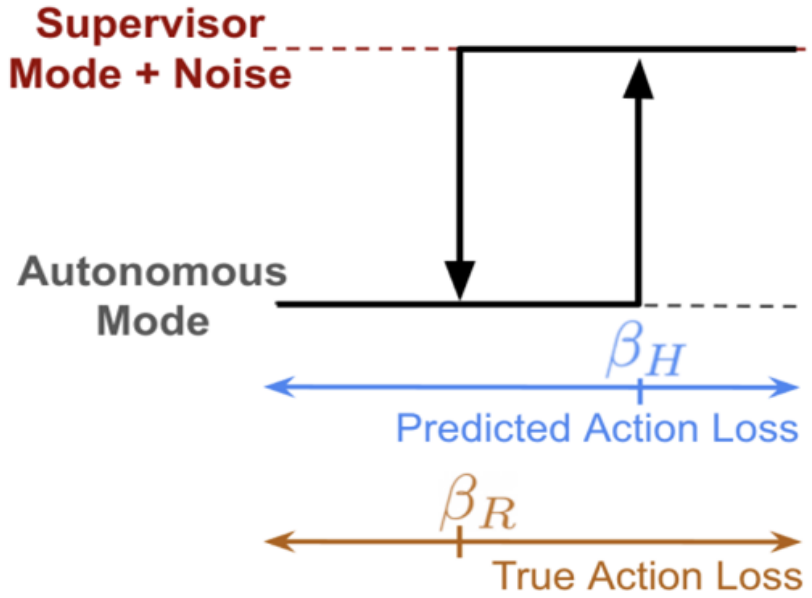

When should a robot ask for help?

Novel (and risky)

When should a robot ask for help?

Novel (and risky)

Risky (but not novel)

Novelty Estimation

Novelty Estimation: Supervisor Mode

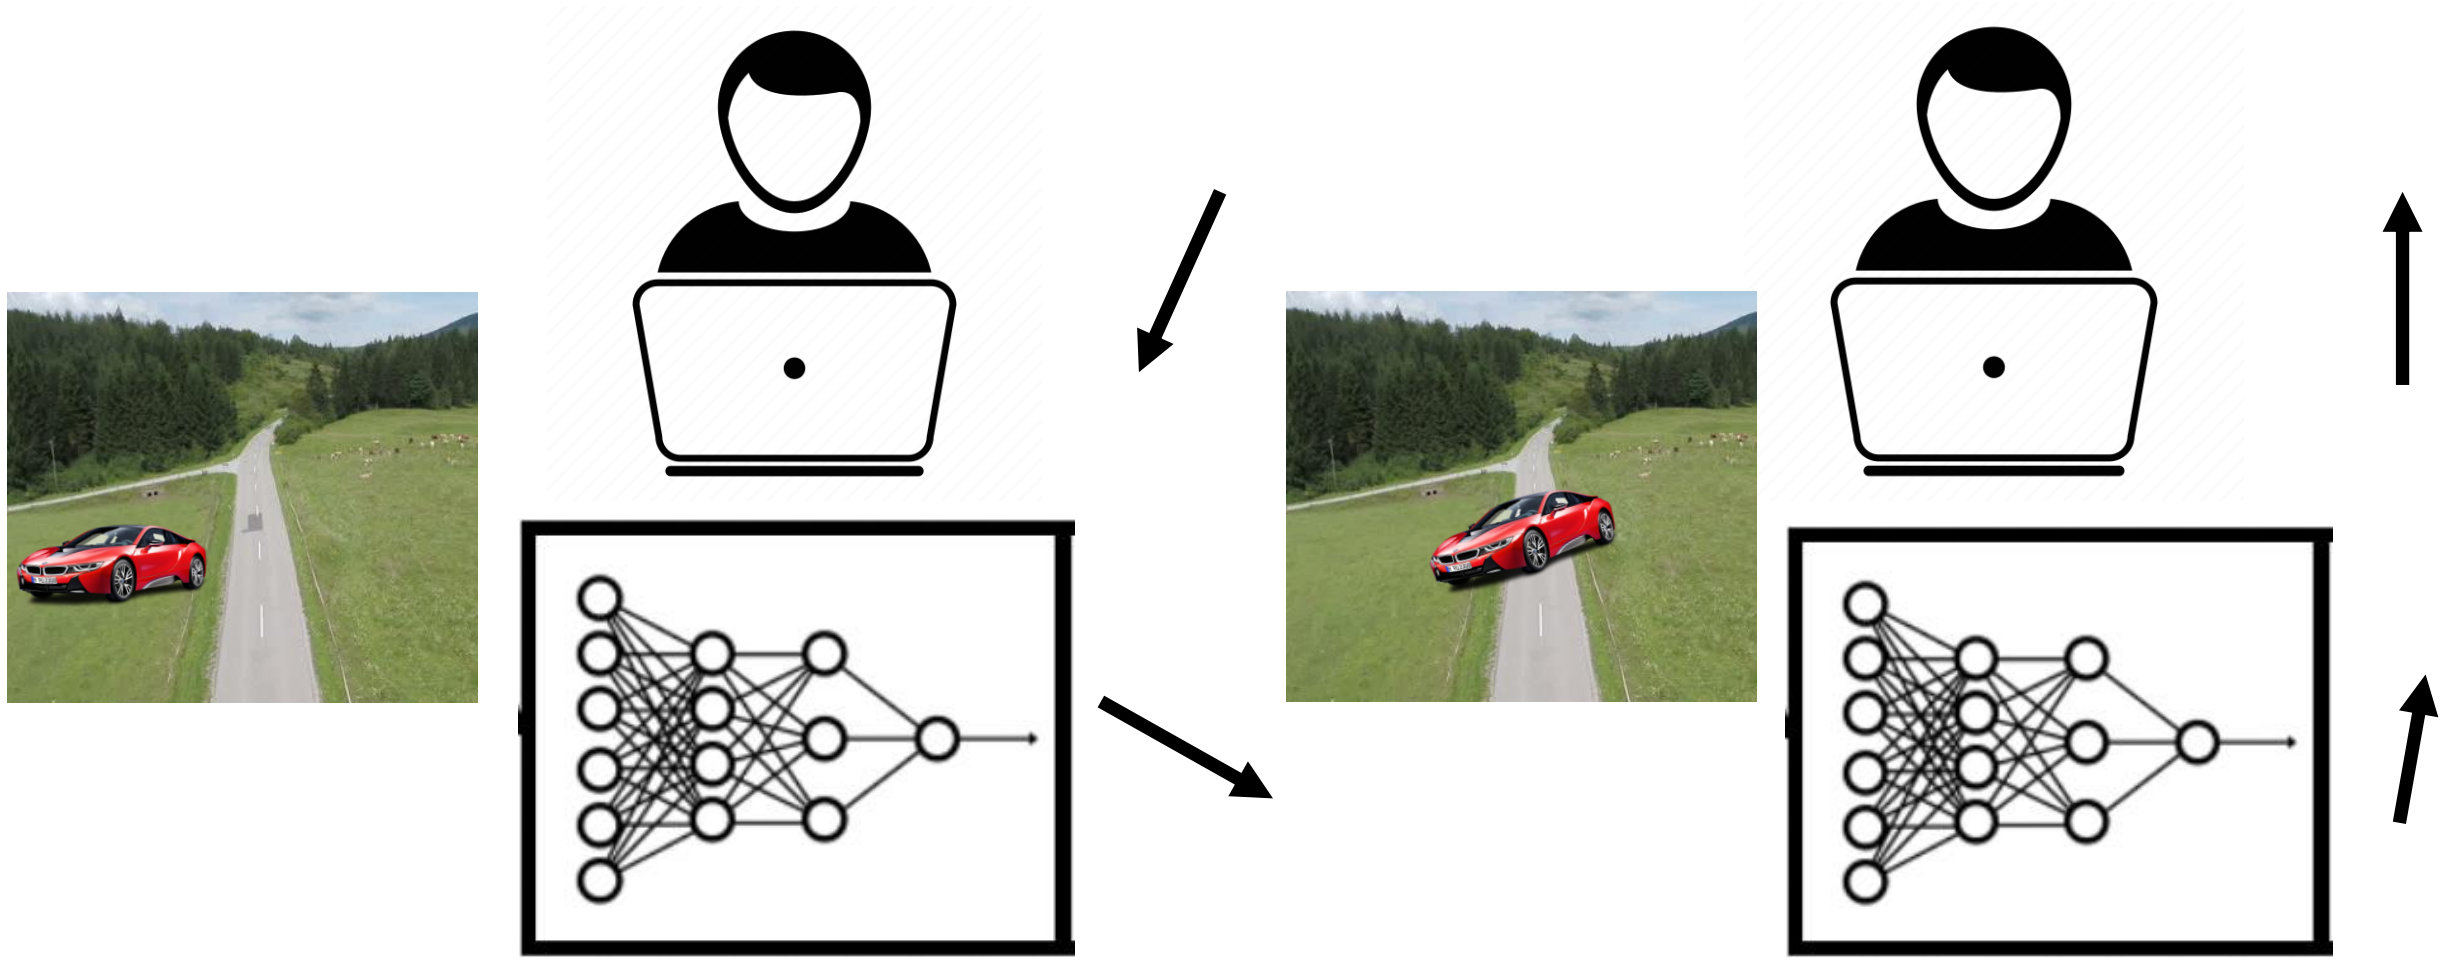

Risk Estimation

$$Q_{\mathcal{G}}^{\pi_r}(s_t, a_t) = \mathbb{E}_{\pi_r}\left[\sum_{t'=t}^{\infty} \gamma^{t'-t} \mathbb{1}_{\mathcal{G}}(s'_t) | s_t, a_t\right]$$

Risk Estimation

$$Q_{\mathcal{G}}^{\pi_r}(s_t, a_t) = \mathbb{E}_{\pi_r}\left[\sum_{t'=t}^{\infty} \gamma^{t'-t} \mathbb{1}_{\mathcal{G}}(s'_t) | s_t, a_t\right]$$

$$\mathbf{Risk}^{\pi_r}(s, a) = 1 - \hat{Q}_{\phi,\mathcal{G}}^{\pi_r}(s, a)$$

Risk Estimation

$$Q_{\mathcal{G}}^{\pi_r}(s_t, a_t) = \mathbb{E}_{\pi_r}\left[\sum_{t'=t}^{\infty} \gamma^{t'-t} \mathbb{1}_{\mathcal{G}}(s'_t) | s_t, a_t\right]$$

$$\mathbf{Risk}^{\pi_r}(s, a) = 1 - \hat{Q}_{\phi,\mathcal{G}}^{\pi_r}(s, a)$$

$$J_{\mathcal{G}}^{Q}(s_t, a_t, s_{t+1}; \phi) =$$

$$\frac{1}{2}\left(\hat{Q}_{\phi,\mathcal{G}}^{\pi_r}(s_t, a_t) - (\mathbb{1}_{\mathcal{G}}(s_t) + \ (1 - \mathbb{1}_{\mathcal{G}}(s_t))\gamma \hat{Q}_{\phi,\mathcal{G}}^{\pi_r}(s_{t+1}, \pi_r(s_{t+1})))\right)^2$$

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$
OR
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

Switch to SUPERVISOR MODE

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$
OR
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

→ Switch to SUPERVISOR MODE

SUPERVISOR MODE

$$||\pi_r(s_t) - \pi_h(s_t)||_2^2 < \delta_r$$
AND
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) < \beta_r$$

→ Switch to AUTONOMOUS MODE

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$
OR
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

→ Switch to SUPERVISOR MODE

SUPERVISOR MODE

$$||\pi_r(s_t) - \pi_h(s_t)||_2^2 < \delta_r$$
AND
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) < \beta_r$$

→ Switch to AUTONOMOUS MODE

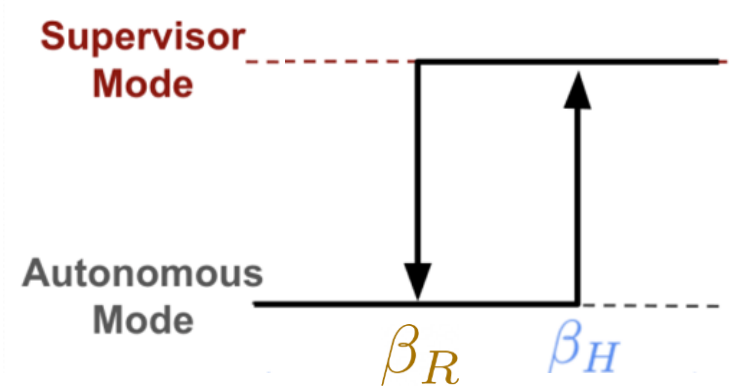

Wait, didn't we just double the number of hyperparameters?

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$

OR

$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

Set to 1-α quantiles of empirical data

→ Switch to SUPERVISOR MODE

SUPERVISOR MODE

$$||\pi_r(s_t) - \pi_h(s_t)||_2^2 < \delta_r$$

AND

$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) < \beta_r$$

→ Switch to AUTONOMOUS MODE

$$\alpha = \frac{\text{desired \# interventions}}{\text{\# robot actions}}$$

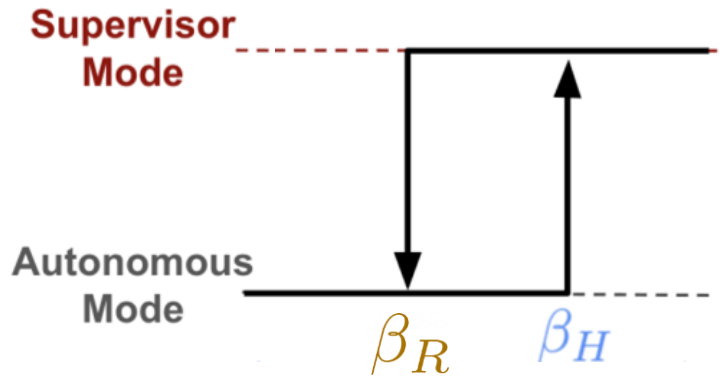

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$

OR

$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

Set to 1-α quantiles of empirical data

→ Switch to SUPERVISOR MODE

SUPERVISOR MODE

$$||\pi_r(s_t) - \pi_h(s_t)||_2^2 < \delta_r$$

AND

$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) < \beta_r$$

→ Switch to AUTONOMOUS MODE

Set to medians of empirical data

$$\alpha = \frac{\text{desired \# interventions}}{\text{\# robot actions}}$$

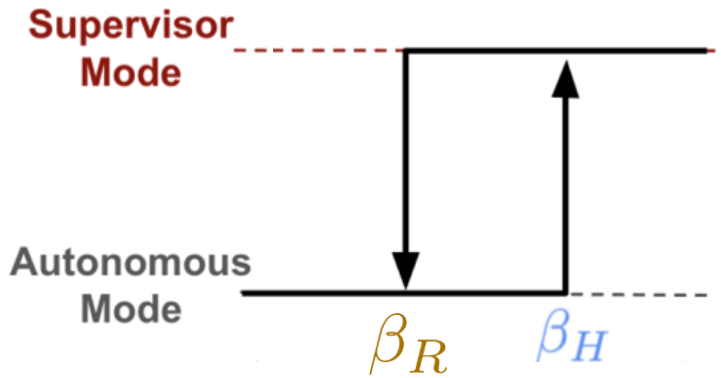

Putting it all together…

AUTONOMOUS MODE

$$\text{Novelty}(s_t) > \delta_h$$
$$\text{OR}$$
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) > \beta_h$$

→ Switch to SUPERVISOR MODE

SUPERVISOR MODE

$$||\pi_r(s_t) - \pi_h(s_t)||_2^2 < \delta_r$$
$$\text{AND}$$
$$\text{Risk}^{\pi_r}(s_t, \pi_r(s_t)) < \beta_r$$

→ Switch to AUTONOMOUS MODE

$$\alpha = \frac{\text{desired \# interventions}}{\text{\# robot actions}}$$

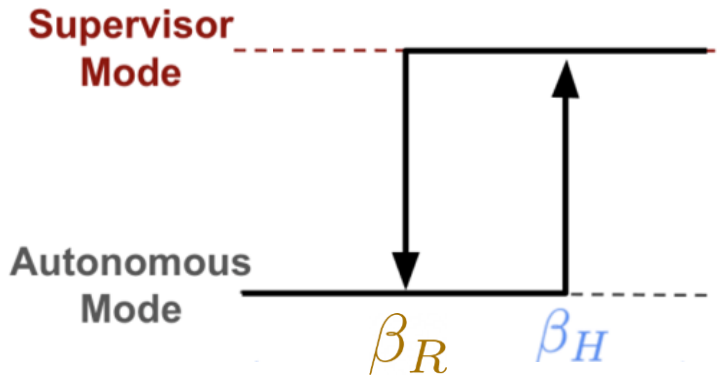

ThriftyDAgger

Target percent of time human wants to give interventions.

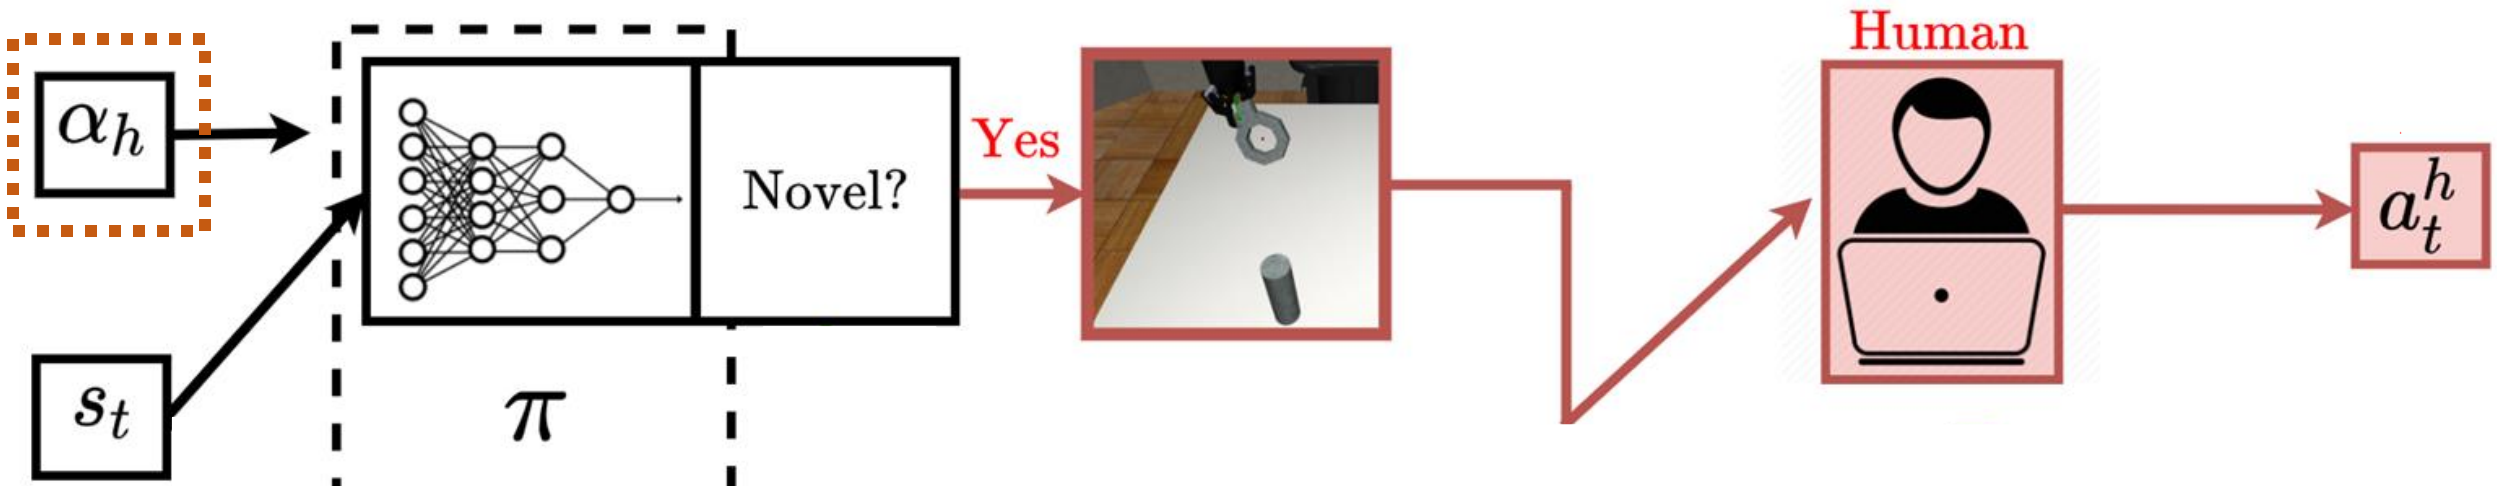

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

ThriftyDAgger

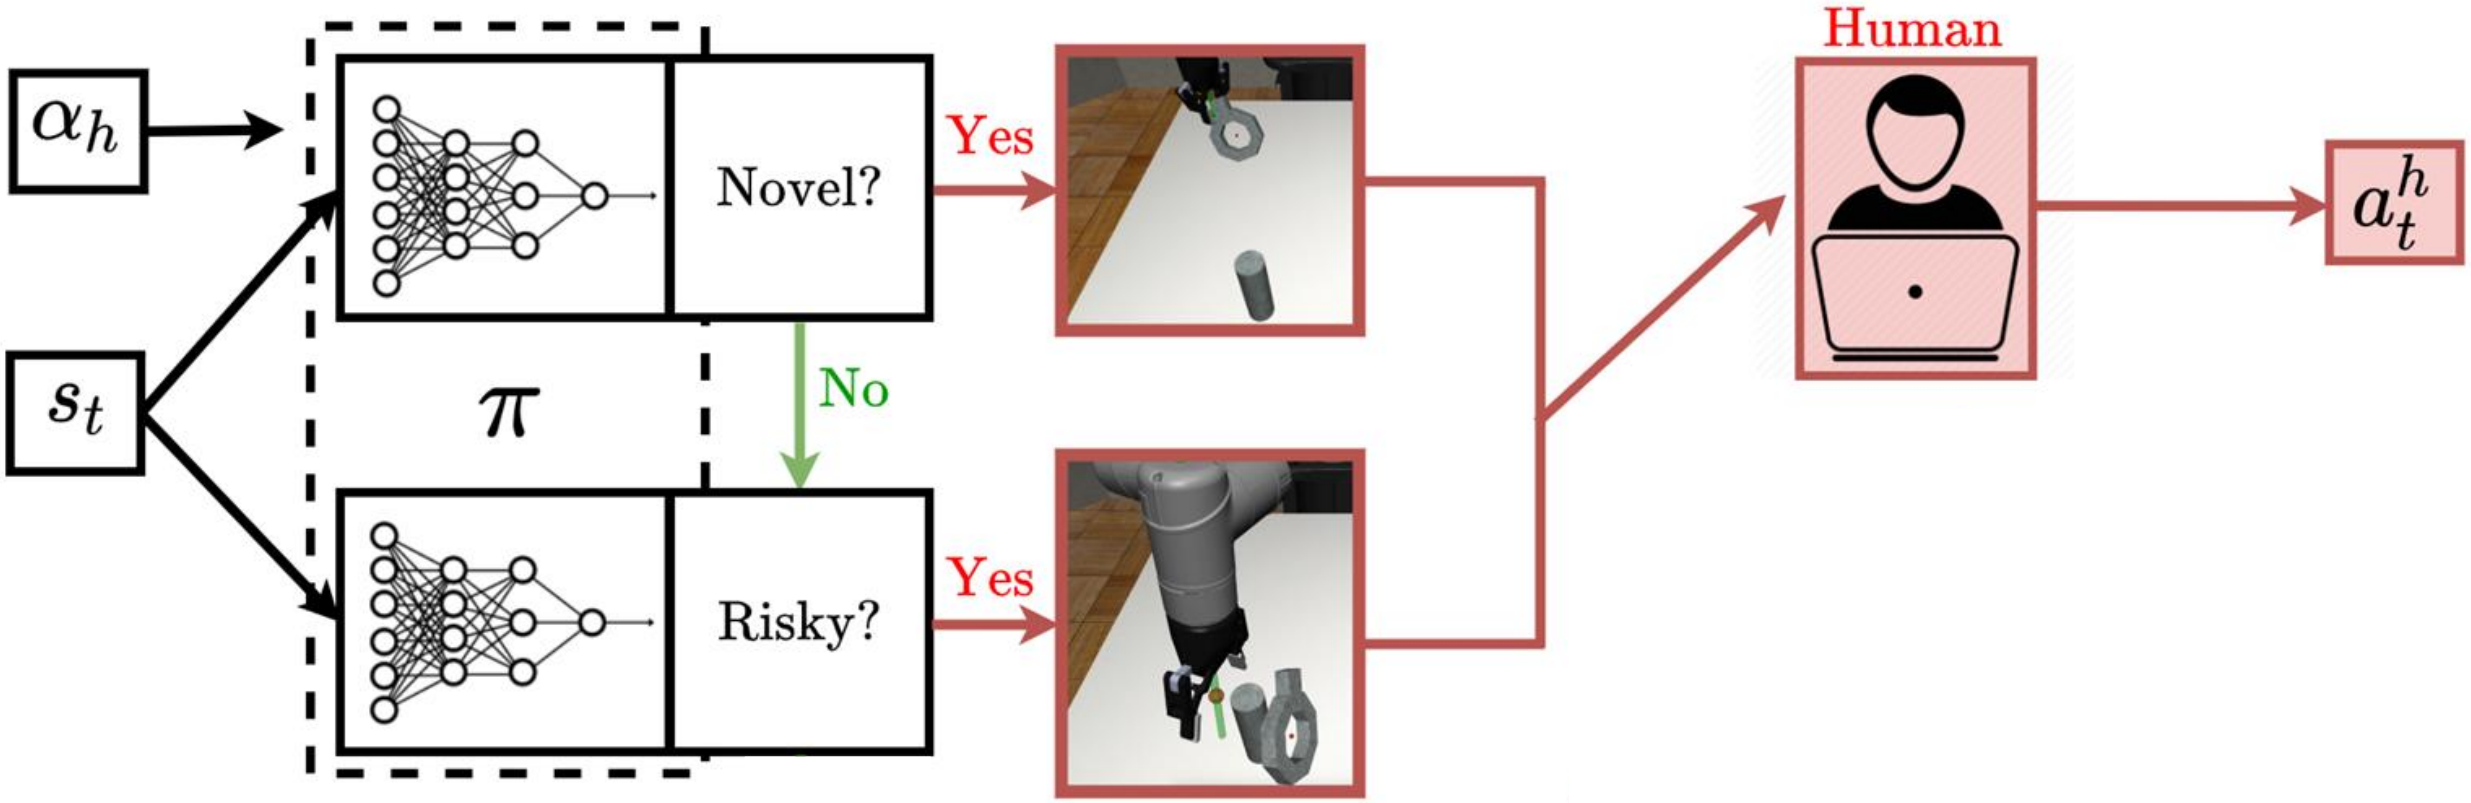

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

ThriftyDAgger

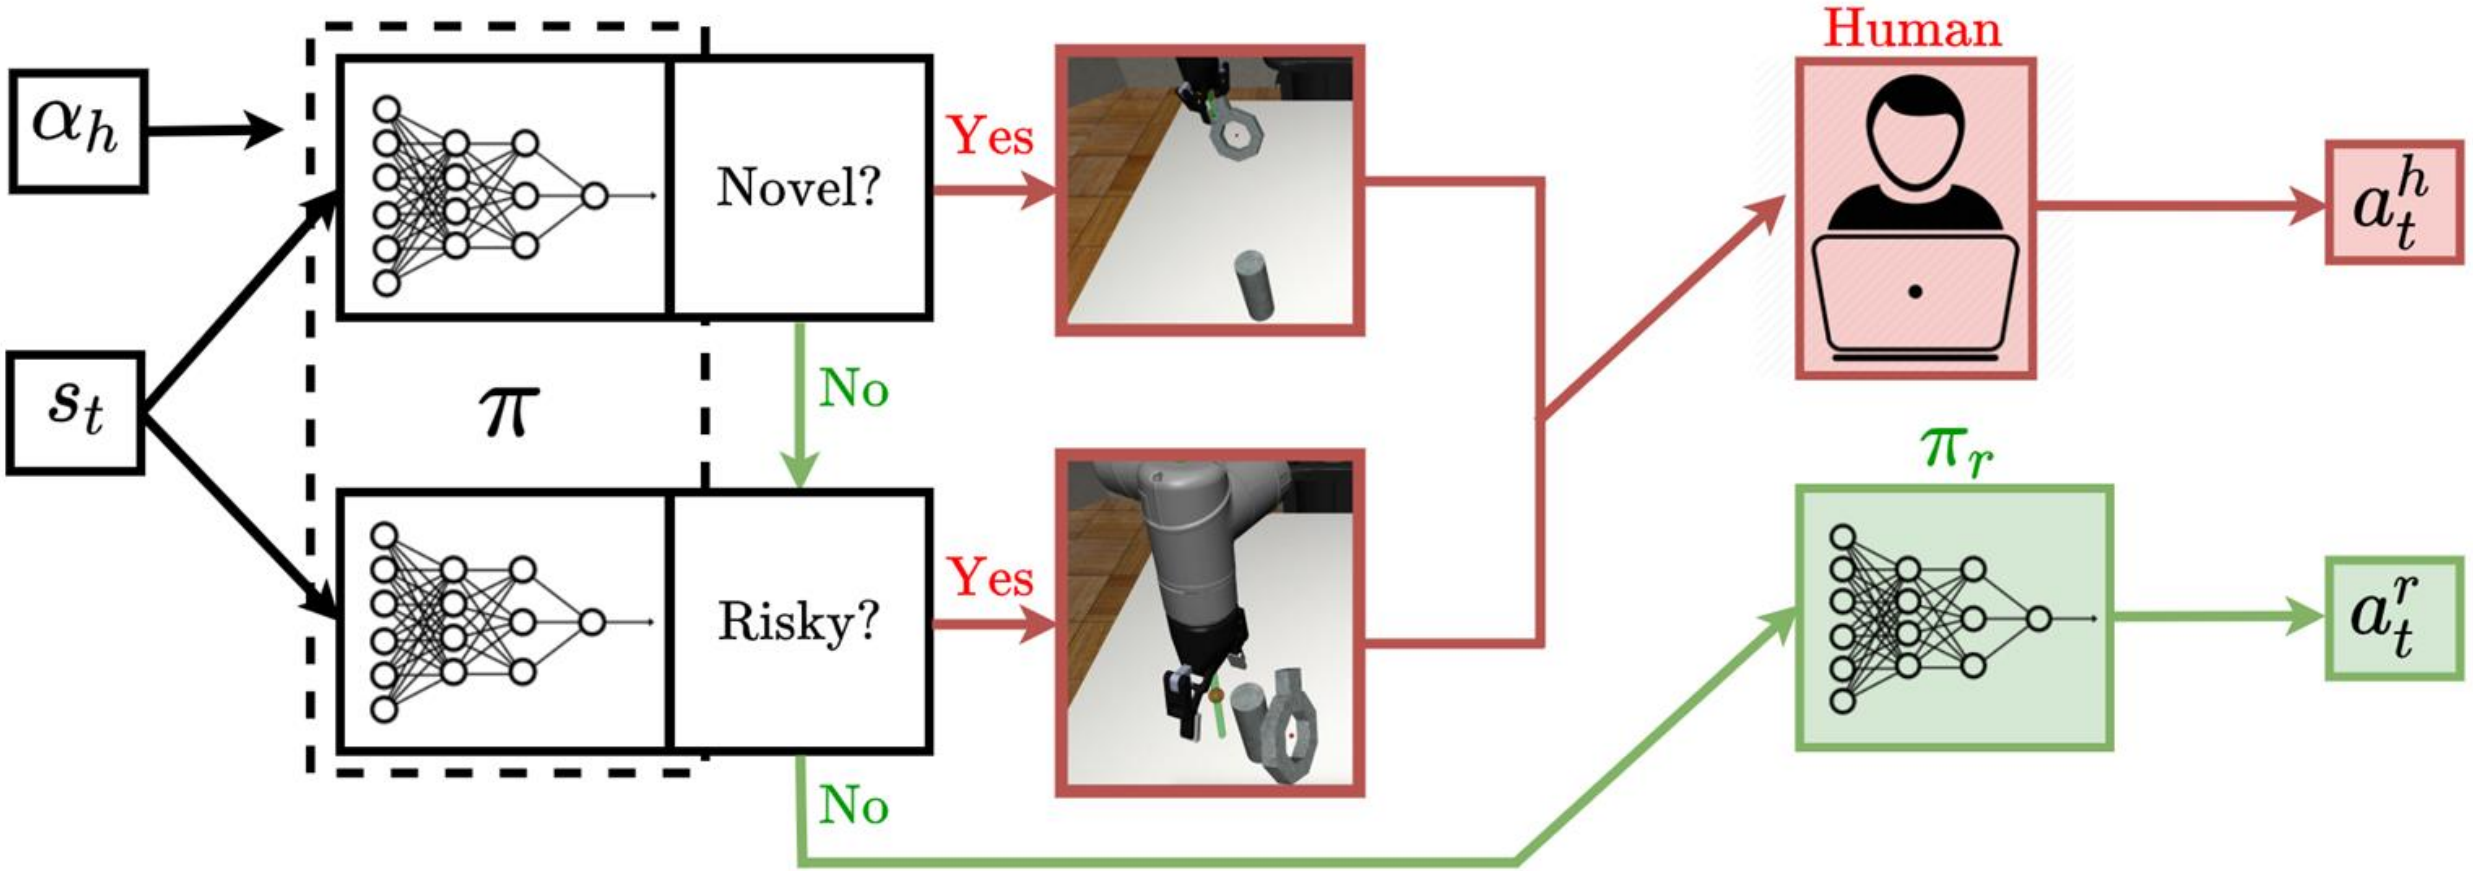

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

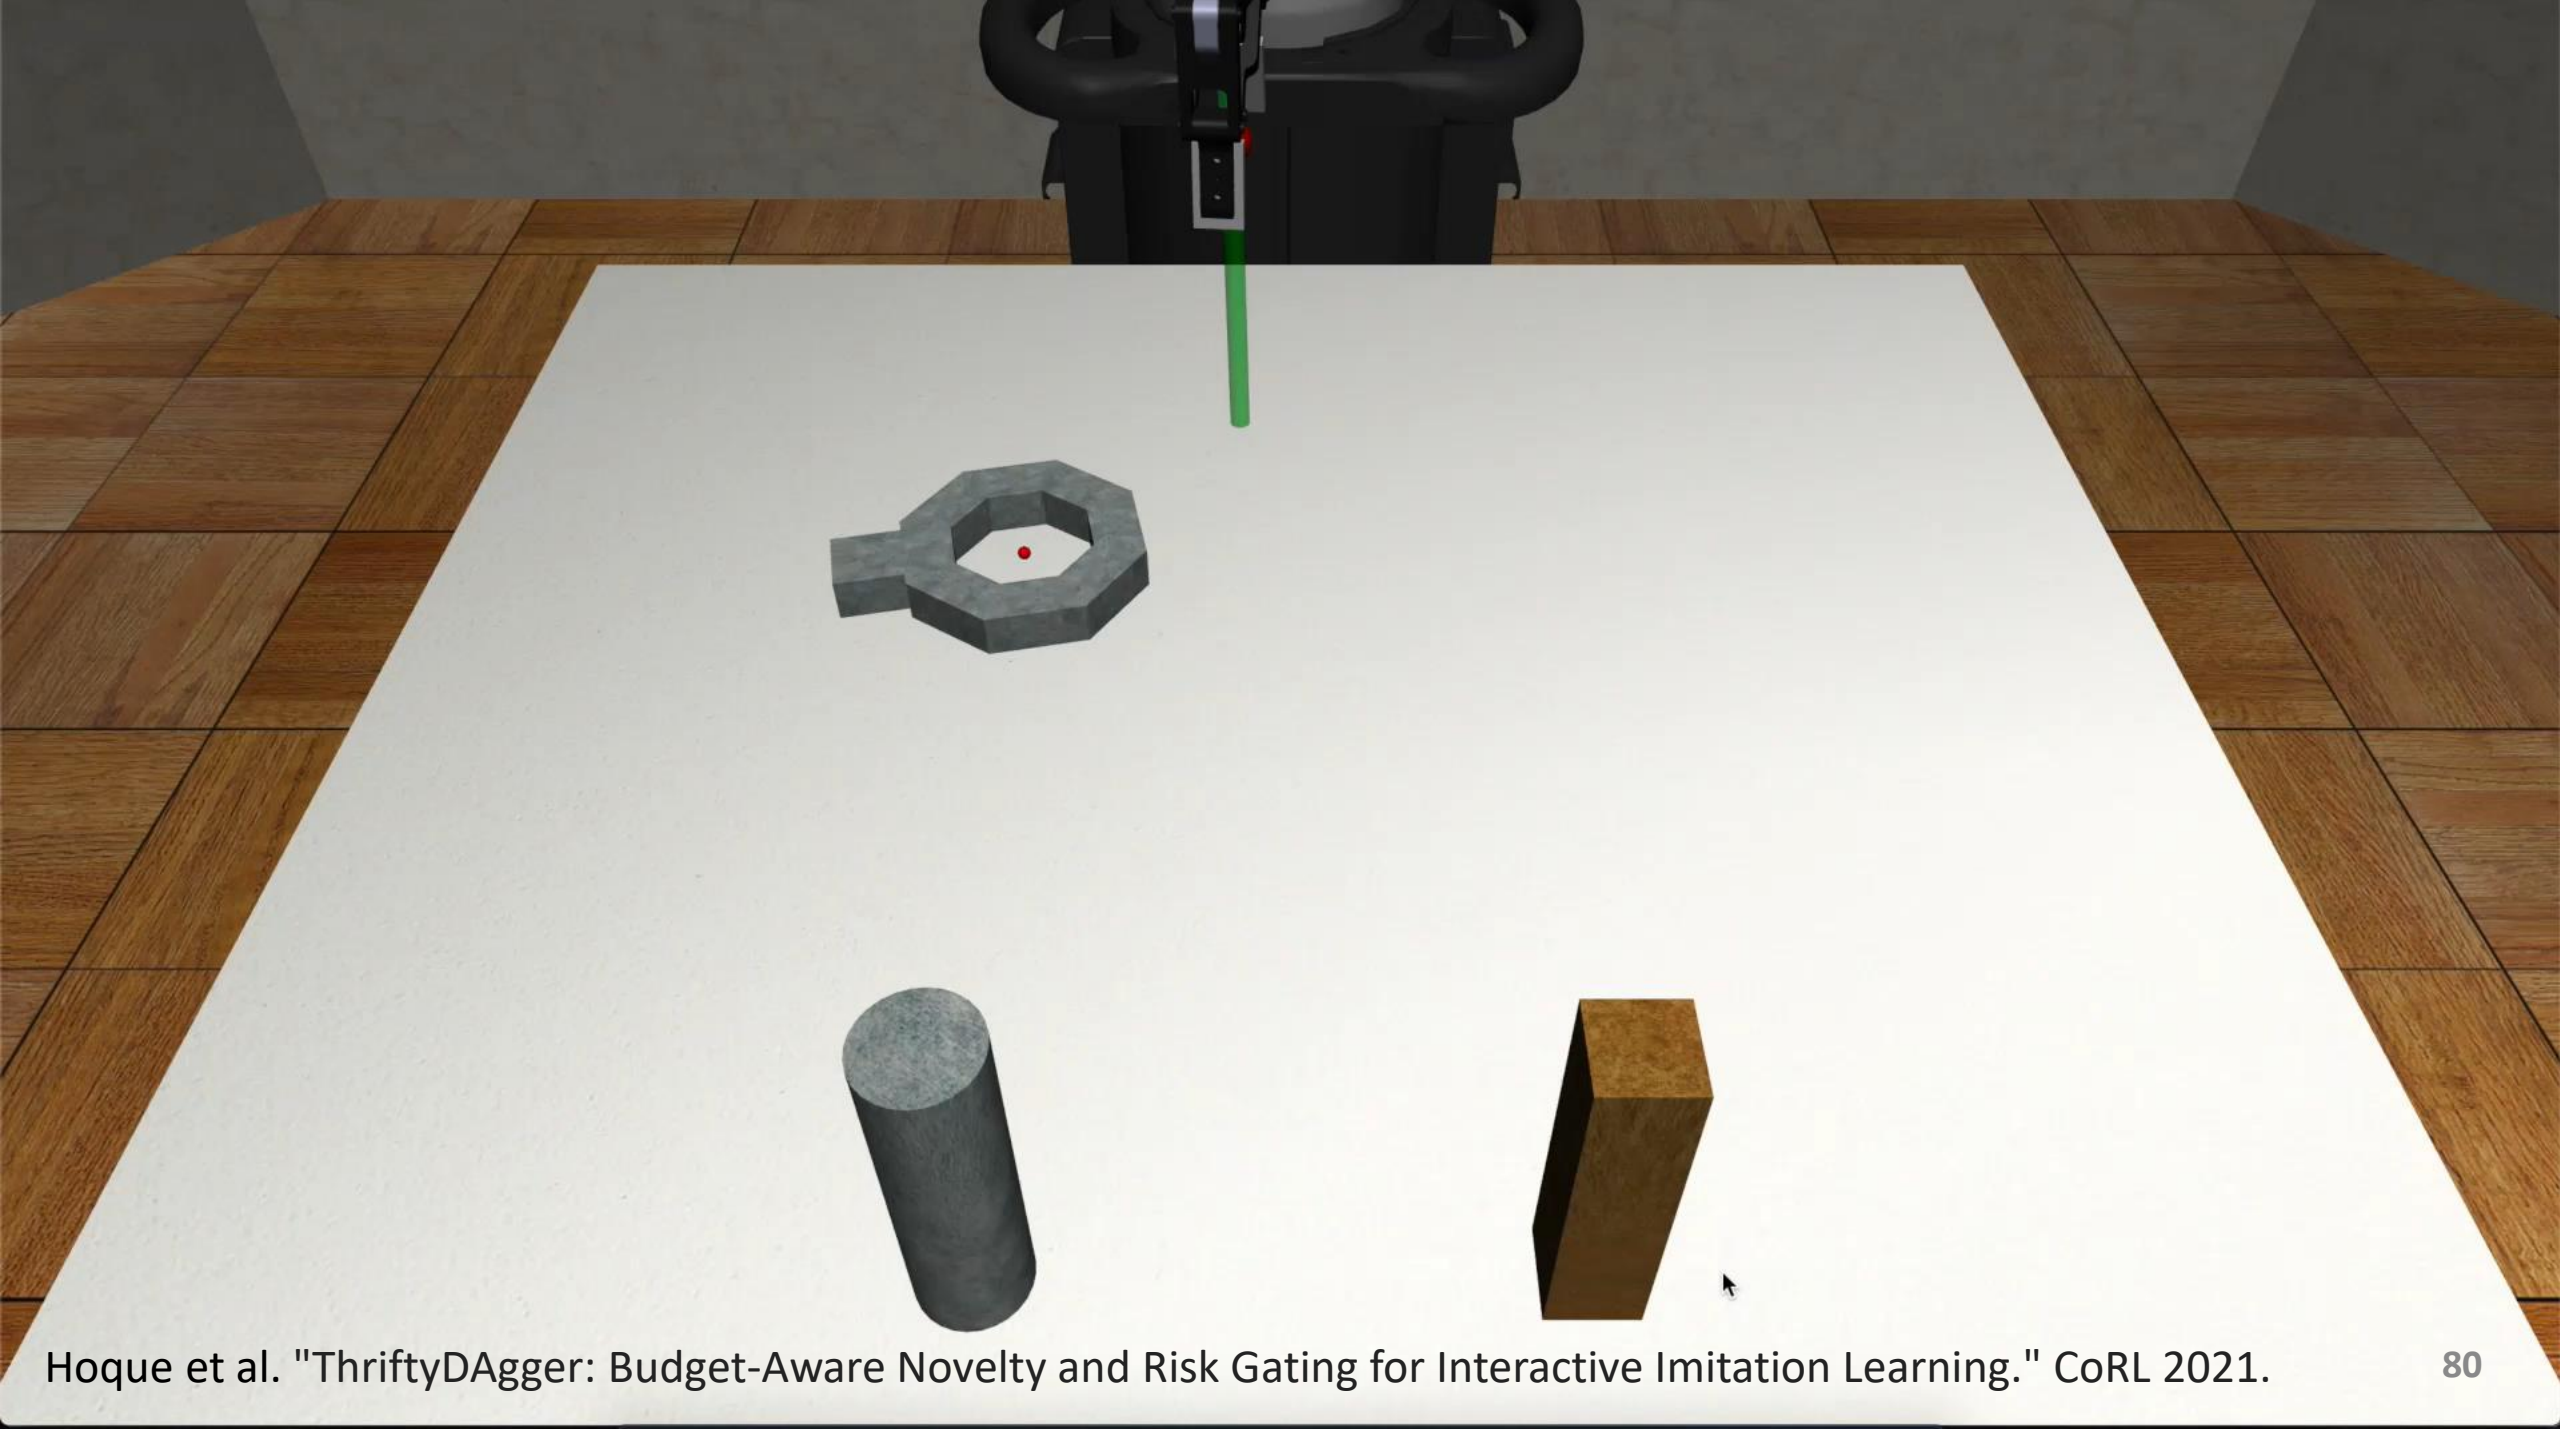

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

Autonomous Mode

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

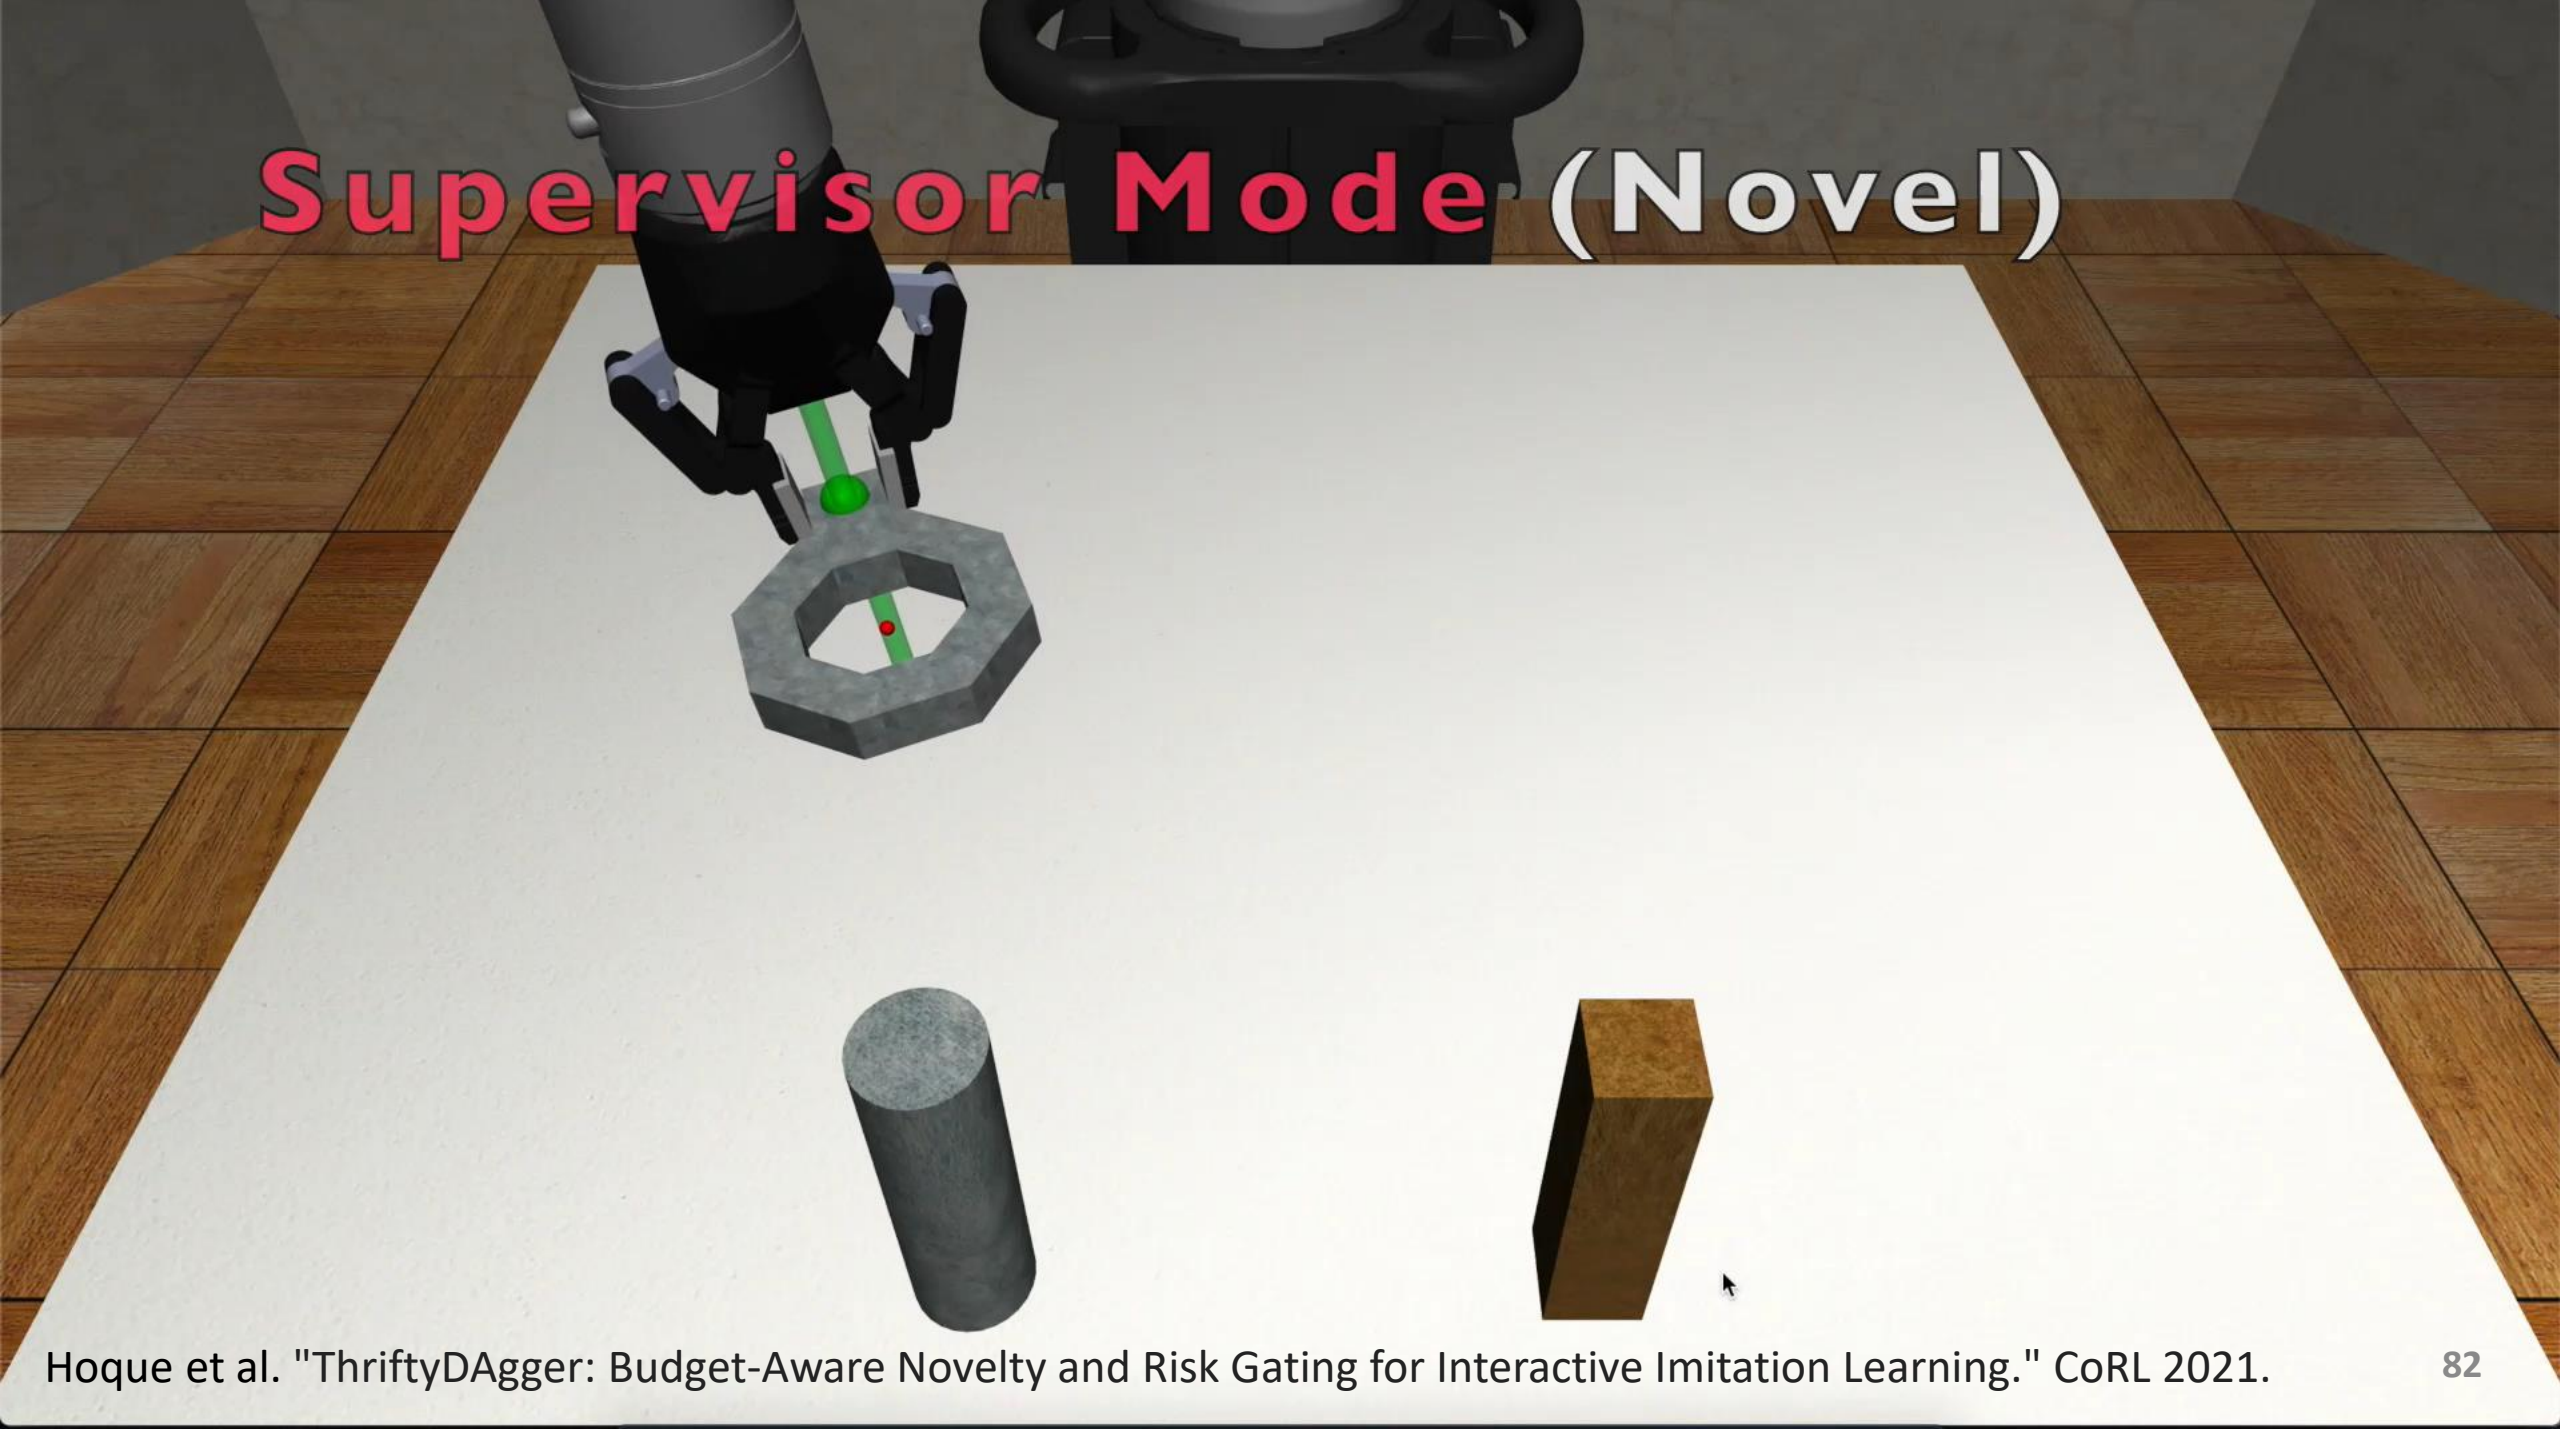

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

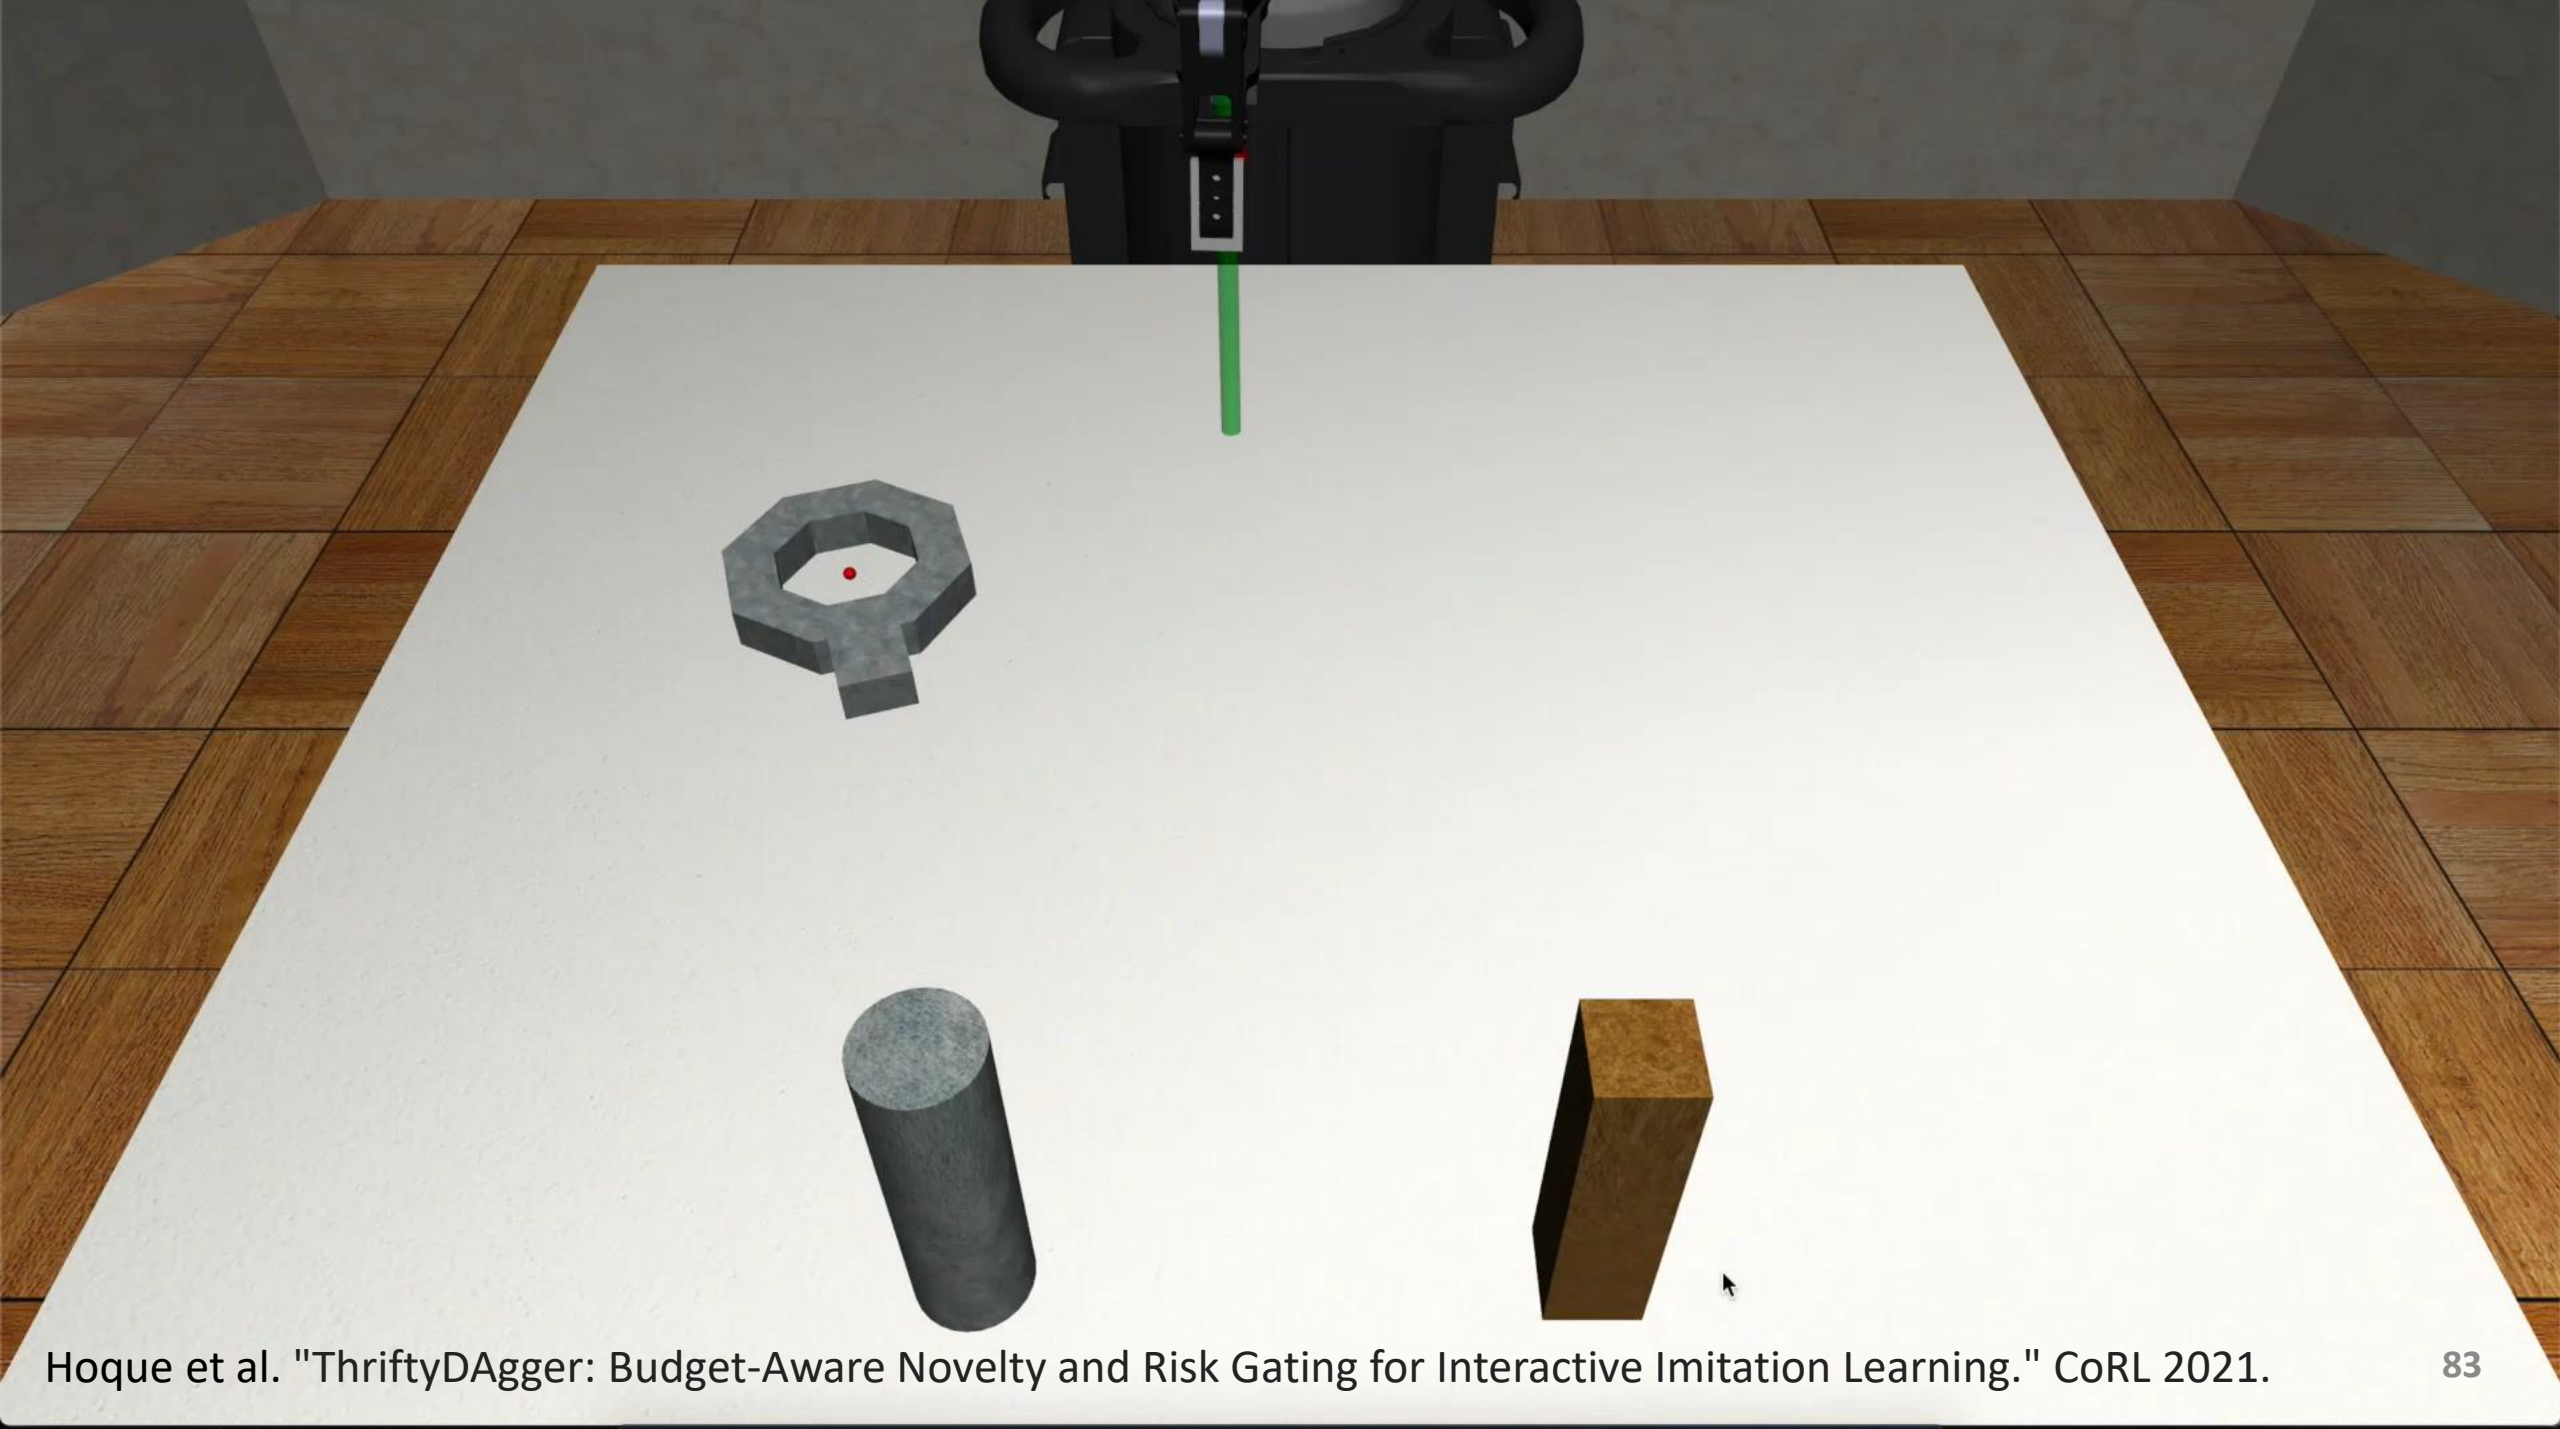

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

Supervisor Mode (Risk)

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

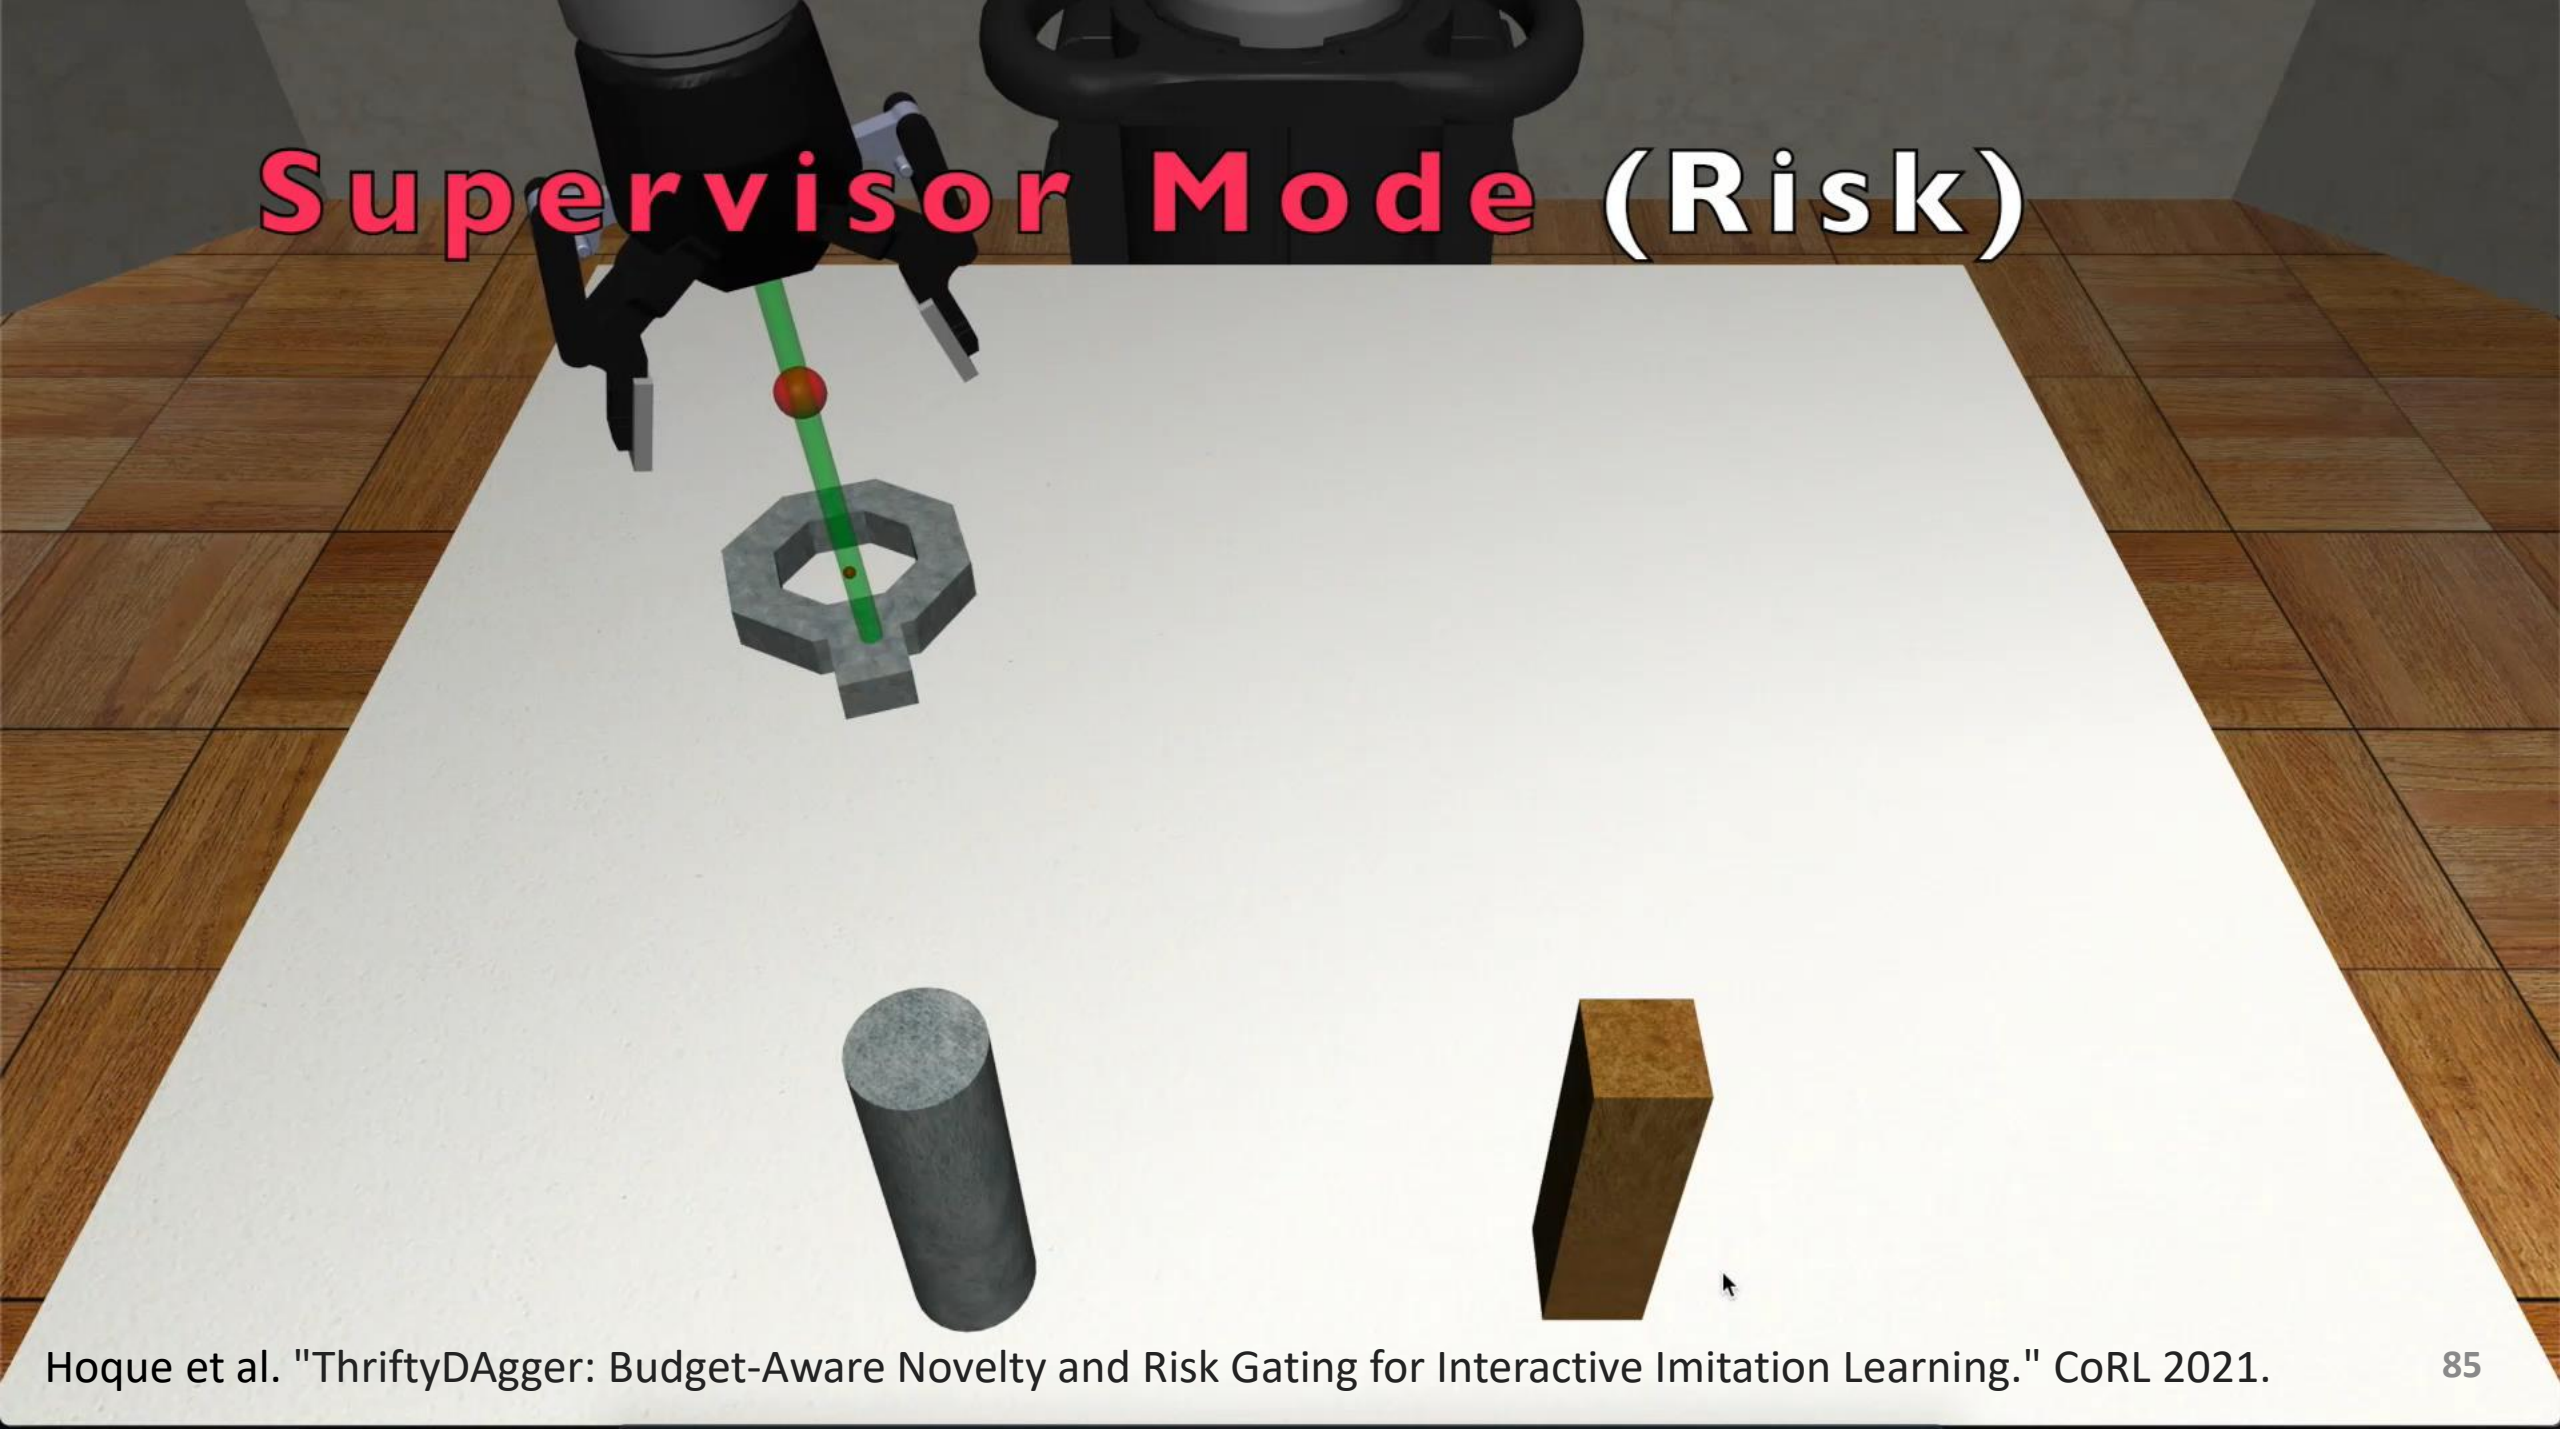

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

Human Demonstration

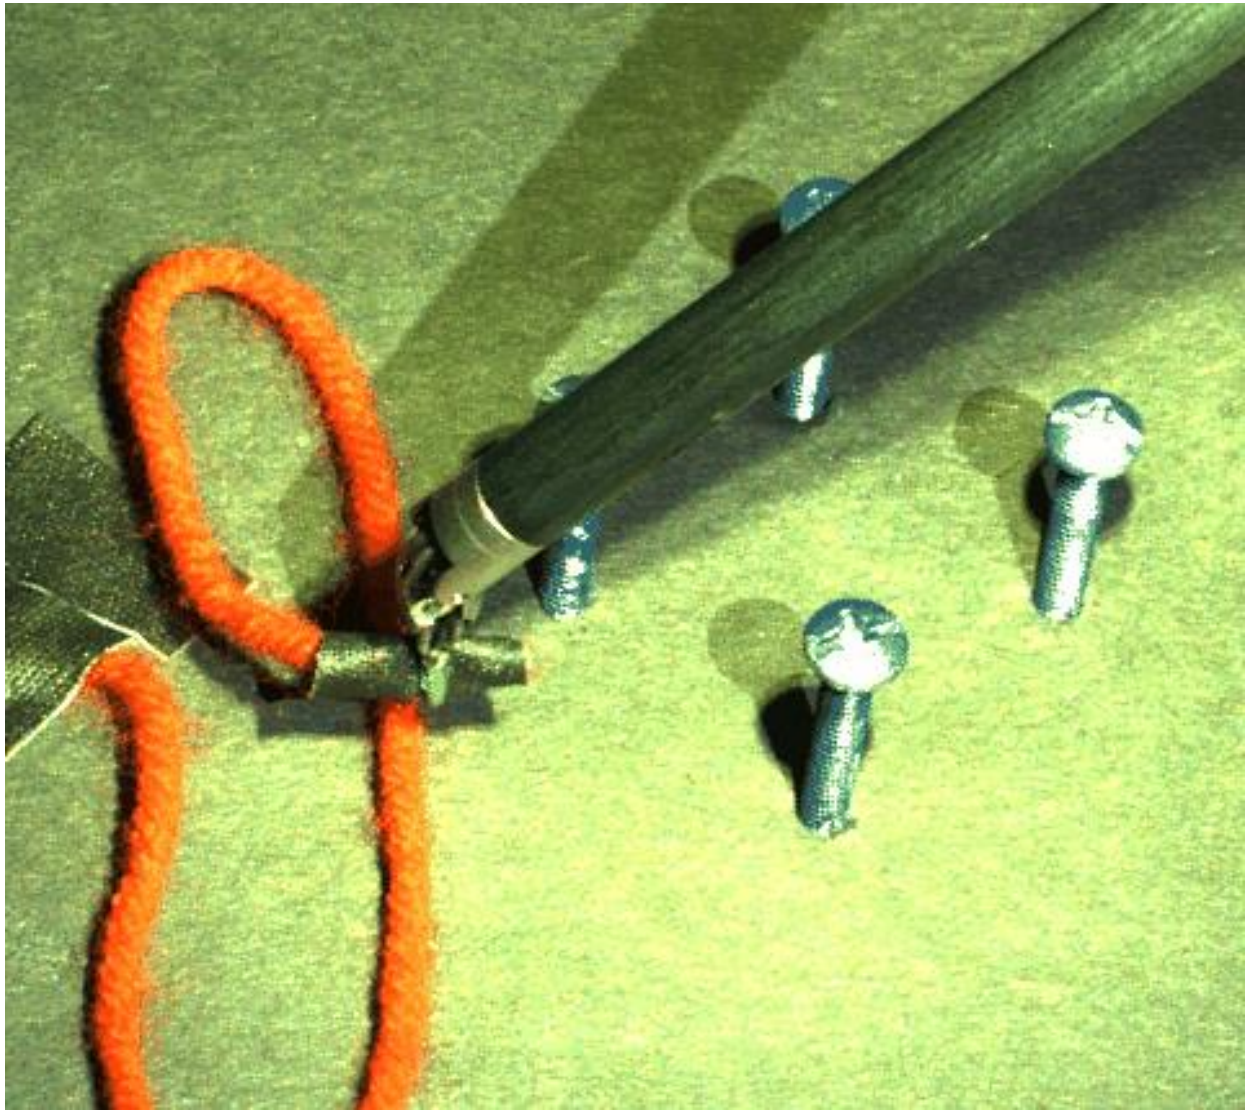

Behavior Cloning

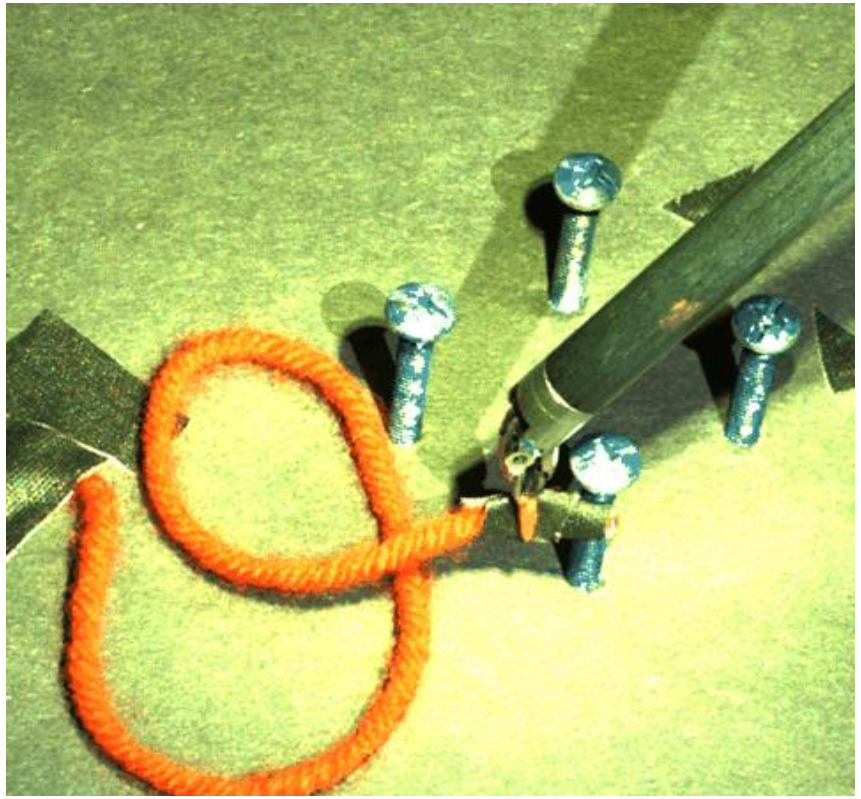

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

Behavior Cloning

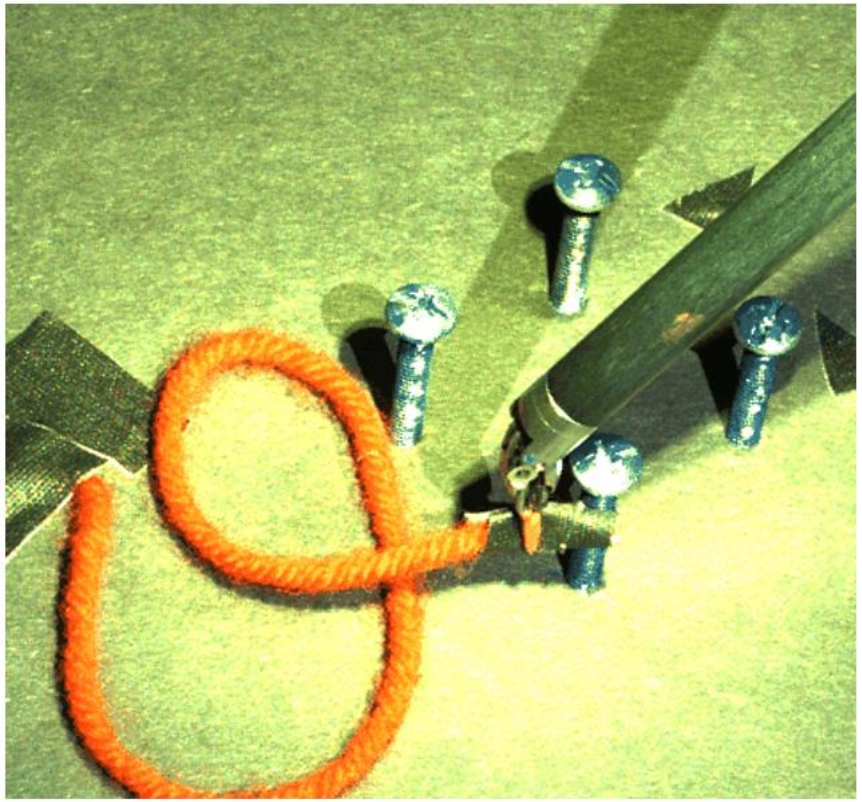

ThriftyDAgger (autonomous)

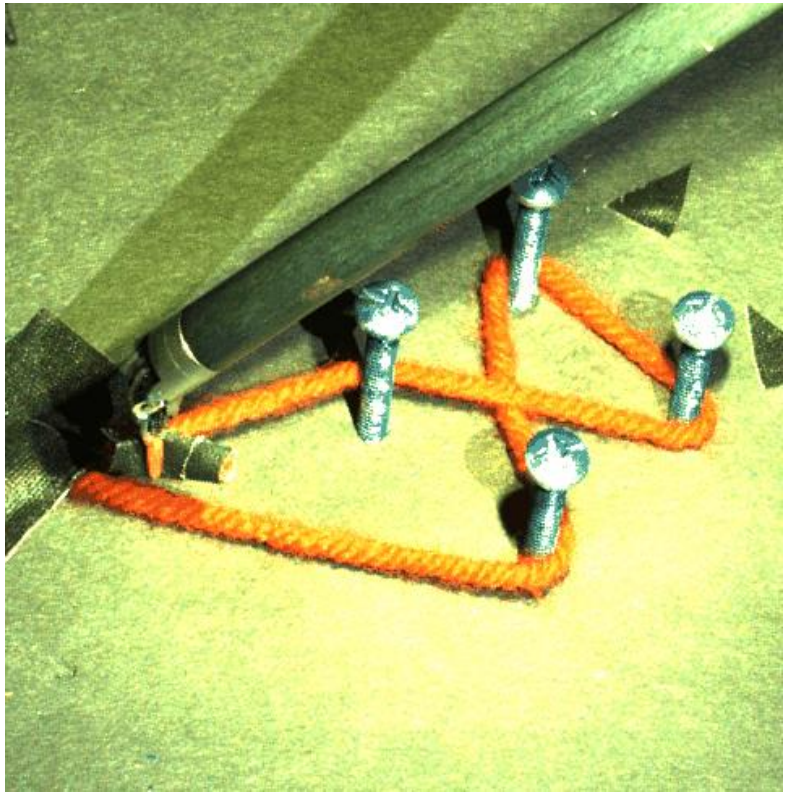

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

Behavior Cloning

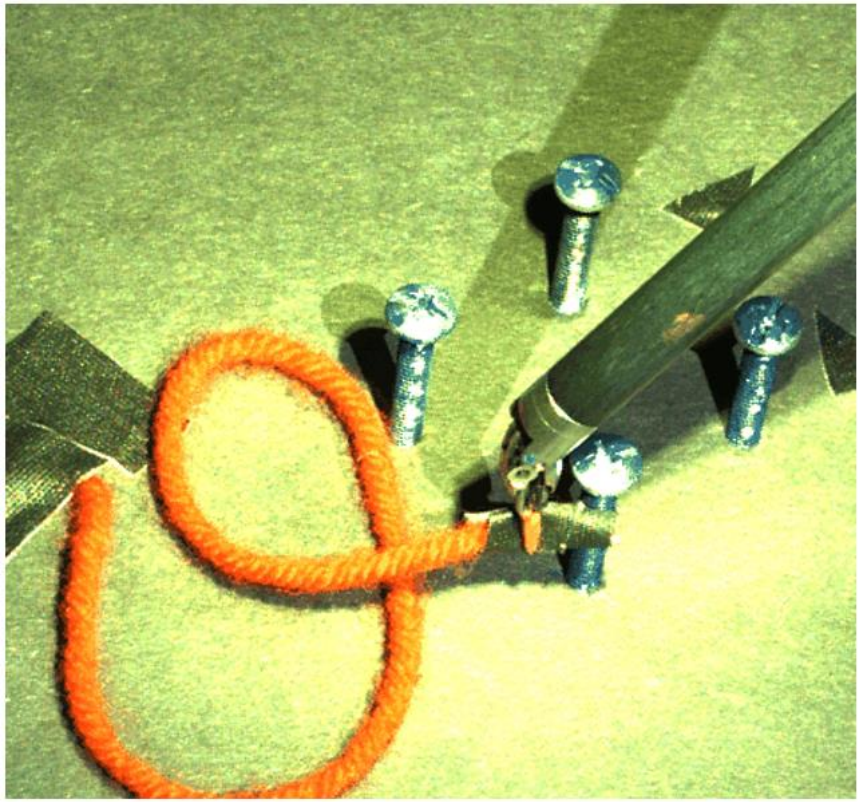

ThriftyDAgger (autonomous)

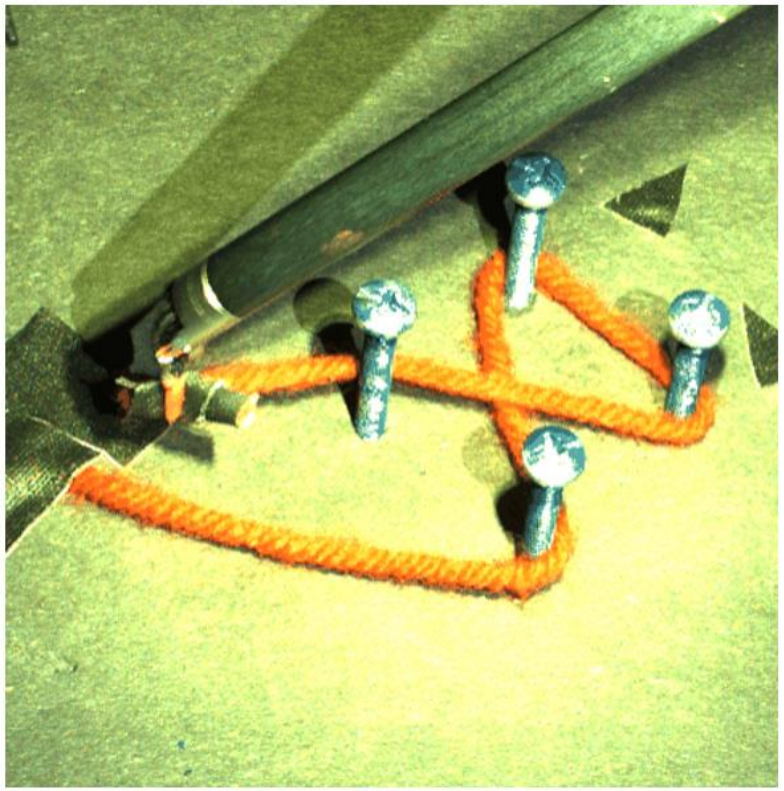

ThriftyDAgger (+human)

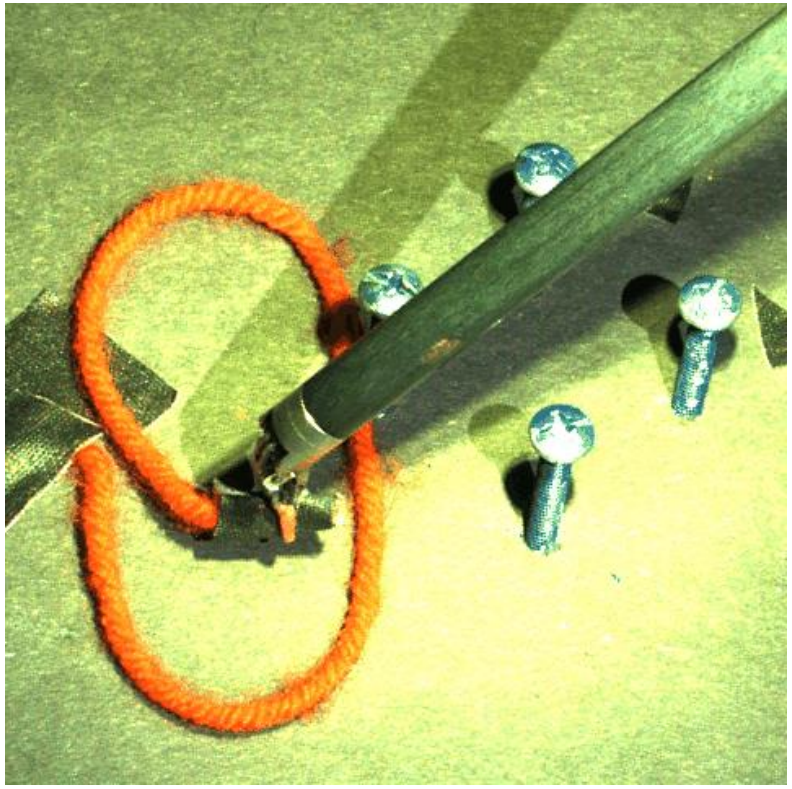

Hoque et al. "ThriftyDAgger: Budget-Aware Novelty and Risk Gating for Interactive Imitation Learning." CoRL 2021.

User Study

N=10 subjects each control 3 robots in simulation.

Robot-Gated

Memory: Non-Match

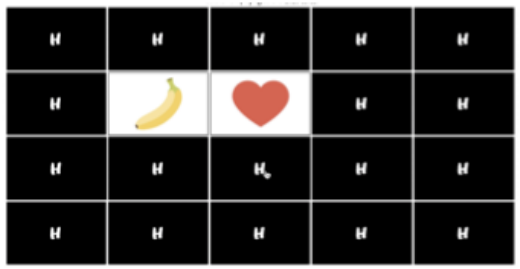

Memory: Match

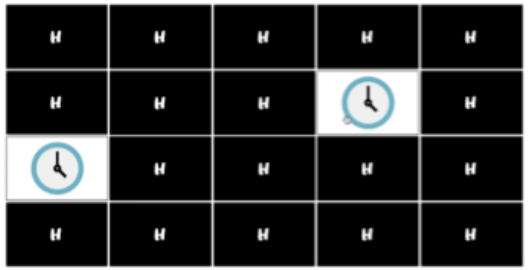

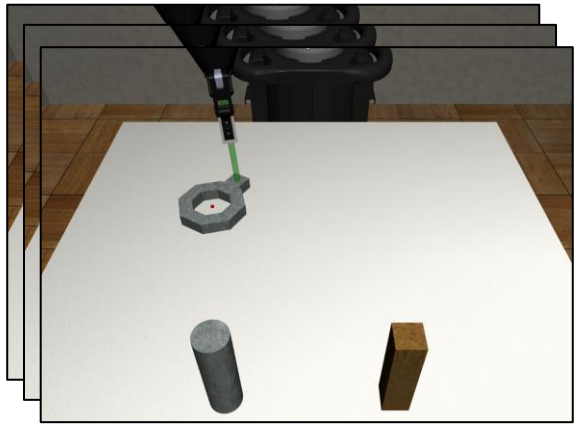

Human-Gated

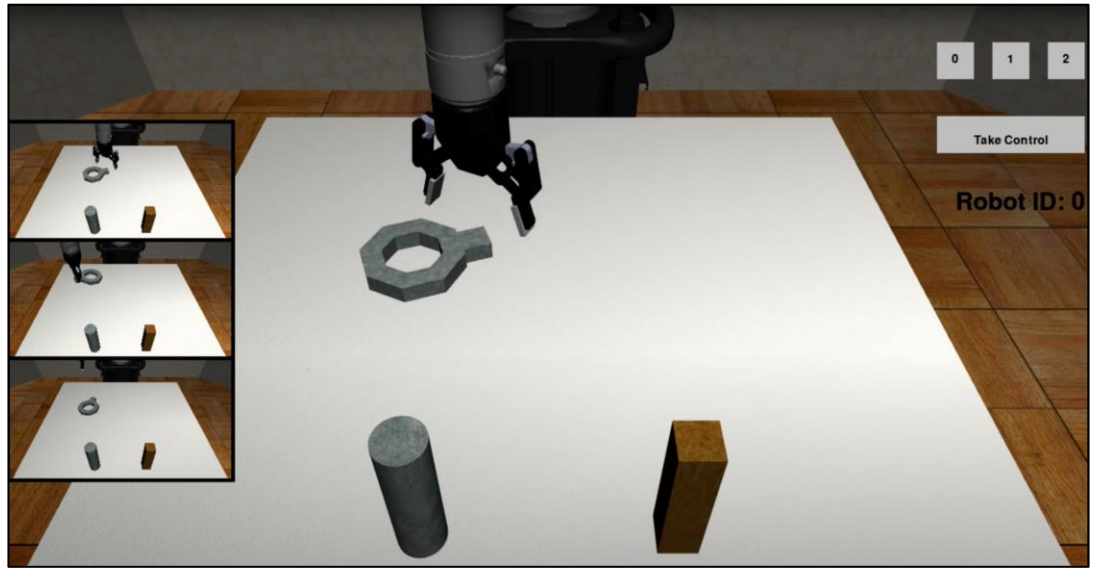

ThriftyDAgger Qualitative Results

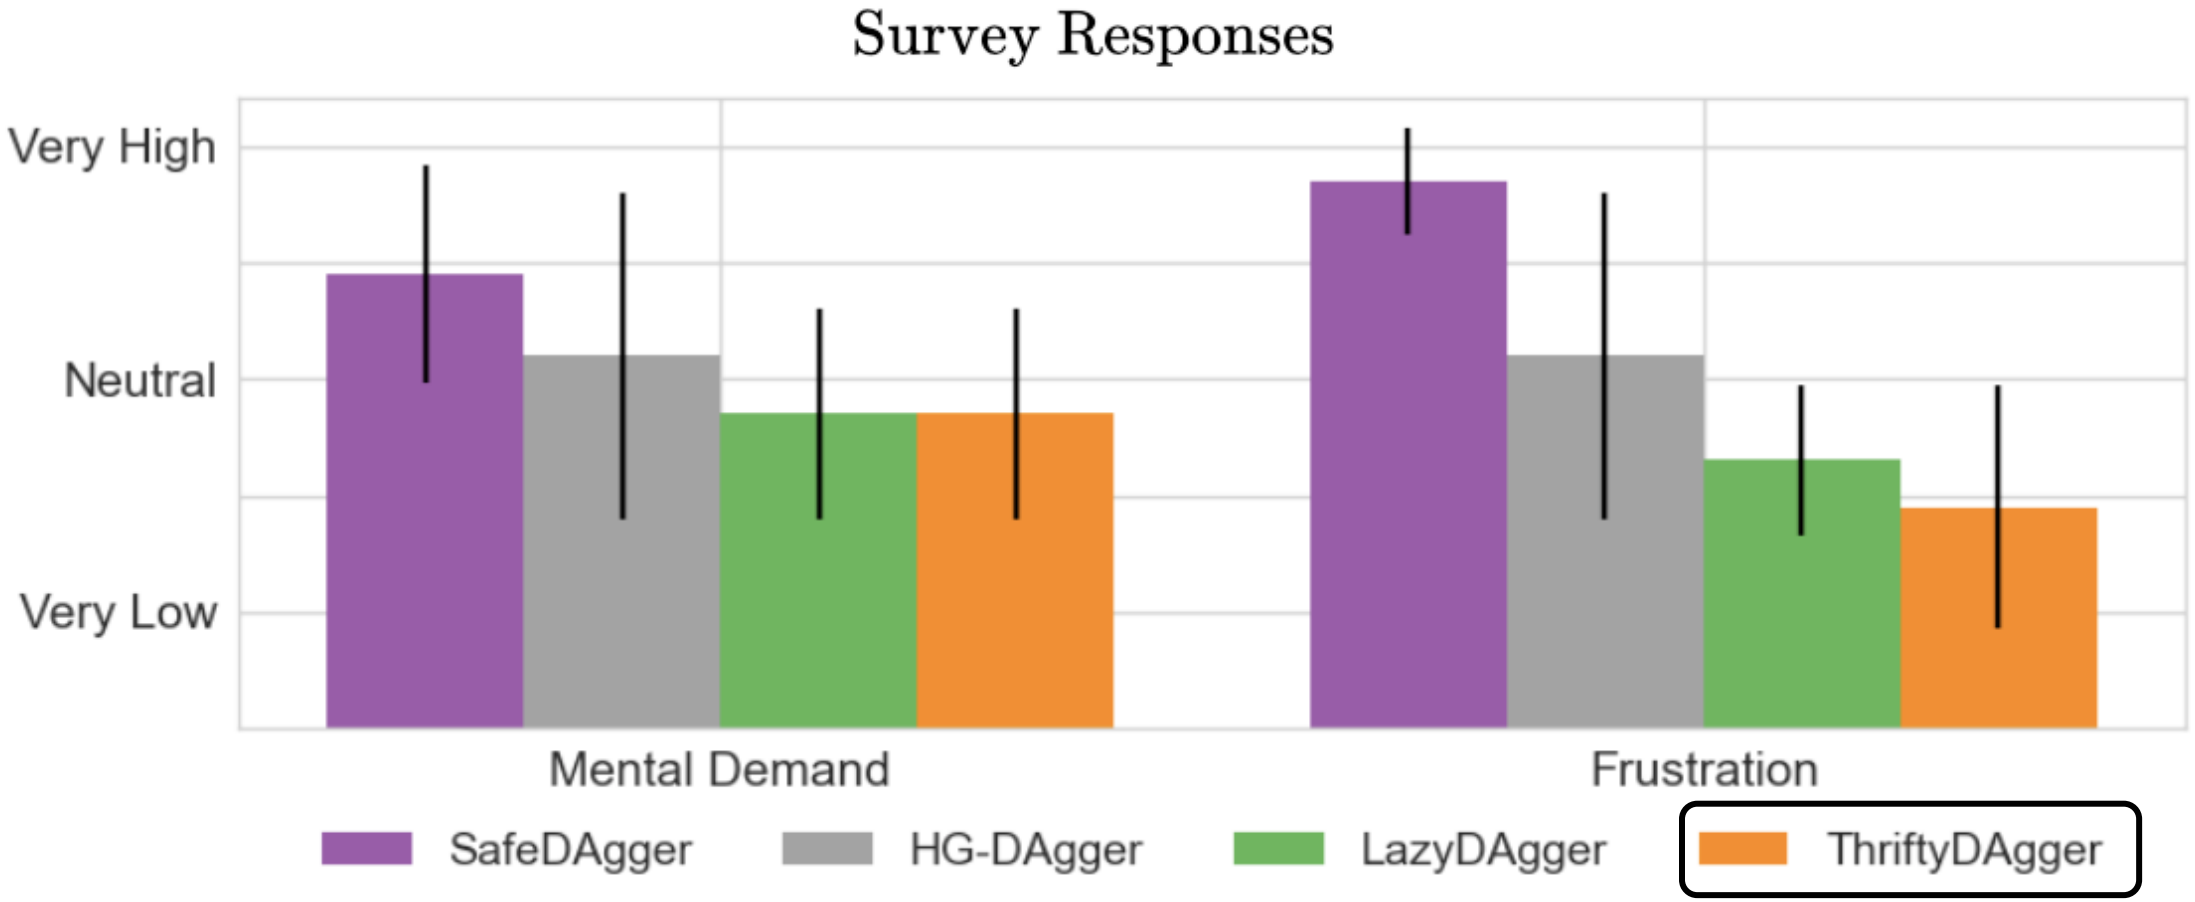

User Study Quantitative Results

ThriftyDAgger had

- 21% fewer human interventions
- 57% more concentration pairs found
- 80% more throughput

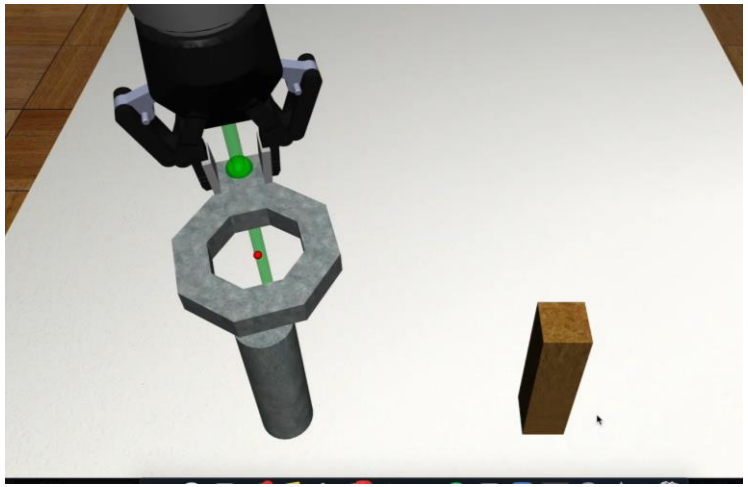

Scalable and safe robot fleets are possible when robots ask for help in ways that minimize human supervisor burden.