

Intro/Refresher on MDPs and Reinforcement Learning



Instructor: Daniel Brown

University of Utah

[Based on slides created by Dan Klein and Pieter Abbeel for CS188 Intro to AI at UC Berkeley. <http://ai.berkeley.edu>.]

Markov Decision Processes

(MDP)

- An MDP is defined by:


- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$, $R(s, a)$, or $R(s')$
- A start state
- Maybe a terminal state



- MDPs are non-deterministic search problems

- One way to solve them is with expectimax search
- We'll have a new tool soon

in general



Other examples of MDPs

- Checkers Boardgame



$S = \{ \text{all board states} \}$

$A = \{ \text{all legal actions} \}$

$T = P(s' | a, s)$

↑
include other players move

$R =$ ^{option A}
+1 capture
-1 lose piece

=
+1 win
-1 loss
0 draw

- Medication treatment

Other examples of MDPs

- Self-driving car
- Language Generation (ChatGPT)

What is Markov about MDPs?

- “Markov” generally means that given the present state, the future and the past are independent. **Conditional Independence!**
- For Markov decision processes, “Markov” means action outcomes depend only on the current state


$$\begin{aligned} &P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ &= \\ &P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$

- This is just like search, where the successor function could only depend on the current state (not the history)



Andrey Markov
(1856-1922)


Types of Markov Models



	System state is fully observable	System state is partially observable
System is autonomous	Markov chain <i>Markov model</i>	Hidden Markov model (HMM)
System is controlled	Markov decision process (MDP)	Partially observable Markov decision process (POMDP)


Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal **policy** $\pi^*: S \rightarrow A$
 - A policy π gives an action for each state
 - An optimal policy is one that maximizes expected utility if followed
 - An explicit policy defines a reflex agent




Optimal policy when $R(s, a, s') = -0.03$
for all non-terminals s


Optimal Policies




$$R(s) = -0.01$$



$$R(s) = -0.03$$



$$R(s) = -0.4$$



$$R(s) = -2.0$$

Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially

$$\gamma \in (0, 1)$$



1

Worth Now



γ

Worth Next Step



γ^2

Worth In Two Steps

Discounting

- How to discount?


- Each time we descend a level, we multiply in the discount once

- Why discount?

- Sooner rewards probably do have higher utility than later rewards
- Also helps our algorithms converge

- Example: discount of 0.5


- $U([1,2,3]) = 1*1 + 0.5*2 + 0.25*3$
- $U([1,2,3]) < U([3,2,1])$



Stationary Preferences

- Theorem: if we assume **stationary preferences**:

$$\begin{aligned} [a_1, a_2, \dots] &\succ [b_1, b_2, \dots] \\ &\Updownarrow \\ [r, a_1, a_2, \dots] &\succ [r, b_1, b_2, \dots] \end{aligned}$$




- Then: there are only two ways to define utilities

- Additive utility: $U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$
- Discounted utility: $U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$

Quiz: Discounting

- Given: reward



- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

- Quiz 1: For $\gamma = 1$, what is the optimal policy?

- Quiz 2: For $\gamma = 0.1$, what is the optimal policy?

- Quiz 3: For which γ are West and East equally good when in state d?



$$\gamma = 10\gamma^3$$

$$1 = 10\gamma^2$$

$$\gamma = \sqrt{1/10} \approx 0.316$$

$$\begin{aligned}
 &\rightarrow 1 \cdot \gamma^0 \\
 &\quad \gamma^1 \\
 &\quad \gamma^2 \\
 &\quad \gamma^3 \\
 &\quad \gamma^4 \\
 &\quad \gamma^5 \\
 &\quad \gamma^6 \\
 &\quad \gamma^7 \\
 &\quad \gamma^8 \\
 &\quad \gamma^9 \\
 &\quad \gamma^{10} \\
 &\quad \gamma^{11} \\
 &\quad \gamma^{12} \\
 &\quad \gamma^{13} \\
 &\quad \gamma^{14} \\
 &\quad \gamma^{15} \\
 &\quad \gamma^{16} \\
 &\quad \gamma^{17} \\
 &\quad \gamma^{18} \\
 &\quad \gamma^{19} \\
 &\quad \gamma^{20} \\
 &\quad \gamma^{21} \\
 &\quad \gamma^{22} \\
 &\quad \gamma^{23} \\
 &\quad \gamma^{24} \\
 &\quad \gamma^{25} \\
 &\quad \gamma^{26} \\
 &\quad \gamma^{27} \\
 &\quad \gamma^{28} \\
 &\quad \gamma^{29} \\
 &\quad \gamma^{30} \\
 &\quad \gamma^{31} \\
 &\quad \gamma^{32} \\
 &\quad \gamma^{33} \\
 &\quad \gamma^{34} \\
 &\quad \gamma^{35} \\
 &\quad \gamma^{36} \\
 &\quad \gamma^{37} \\
 &\quad \gamma^{38} \\
 &\quad \gamma^{39} \\
 &\quad \gamma^{40} \\
 &\quad \gamma^{41} \\
 &\quad \gamma^{42} \\
 &\quad \gamma^{43} \\
 &\quad \gamma^{44} \\
 &\quad \gamma^{45} \\
 &\quad \gamma^{46} \\
 &\quad \gamma^{47} \\
 &\quad \gamma^{48} \\
 &\quad \gamma^{49} \\
 &\quad \gamma^{50} \\
 &\quad \gamma^{51} \\
 &\quad \gamma^{52} \\
 &\quad \gamma^{53} \\
 &\quad \gamma^{54} \\
 &\quad \gamma^{55} \\
 &\quad \gamma^{56} \\
 &\quad \gamma^{57} \\
 &\quad \gamma^{58} \\
 &\quad \gamma^{59} \\
 &\quad \gamma^{60} \\
 &\quad \gamma^{61} \\
 &\quad \gamma^{62} \\
 &\quad \gamma^{63} \\
 &\quad \gamma^{64} \\
 &\quad \gamma^{65} \\
 &\quad \gamma^{66} \\
 &\quad \gamma^{67} \\
 &\quad \gamma^{68} \\
 &\quad \gamma^{69} \\
 &\quad \gamma^{70} \\
 &\quad \gamma^{71} \\
 &\quad \gamma^{72} \\
 &\quad \gamma^{73} \\
 &\quad \gamma^{74} \\
 &\quad \gamma^{75} \\
 &\quad \gamma^{76} \\
 &\quad \gamma^{77} \\
 &\quad \gamma^{78} \\
 &\quad \gamma^{79} \\
 &\quad \gamma^{80} \\
 &\quad \gamma^{81} \\
 &\quad \gamma^{82} \\
 &\quad \gamma^{83} \\
 &\quad \gamma^{84} \\
 &\quad \gamma^{85} \\
 &\quad \gamma^{86} \\
 &\quad \gamma^{87} \\
 &\quad \gamma^{88} \\
 &\quad \gamma^{89} \\
 &\quad \gamma^{90} \\
 &\quad \gamma^{91} \\
 &\quad \gamma^{92} \\
 &\quad \gamma^{93} \\
 &\quad \gamma^{94} \\
 &\quad \gamma^{95} \\
 &\quad \gamma^{96} \\
 &\quad \gamma^{97} \\
 &\quad \gamma^{98} \\
 &\quad \gamma^{99}
 \end{aligned}$$

Infinite Utilities?!

- Problem: What if the game lasts forever? Do we get infinite rewards?

- Solutions:


- Finite horizon: (similar to depth-limited search)
 - Terminate episodes after a fixed T steps (e.g. life)
 - Gives nonstationary policies (π depends on time left)

- Discounting: use $0 < \gamma < 1$

Why?

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$

- Smaller γ means smaller “horizon” – shorter term focus
- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “overheated” for racing)



$$A = \sum_{t=0}^{\infty} R_{\max} \gamma^t = R_{\max} + \cancel{\gamma R_{\max}} + \cancel{\gamma^2 R_{\max}} + \dots$$

$$B = \gamma \sum_{t=0}^{\infty} R_{\max} \gamma^t = \cancel{\gamma R_{\max}} + \cancel{\gamma^2 R_{\max}} + \cancel{\gamma^3 R_{\max}} \dots$$

$$A - B = R_{\max}$$

$$\underbrace{\left(\sum_{t=0}^{\infty} R_{\max} \gamma^t \right)}_{\parallel R_{\max}} (1 - \gamma) = R_{\max}$$

$$\frac{R_{\max}}{1 - \gamma}$$


MDP Notation

- Markov decision processes:

- Set of states S
- Start state s_0
- Set of actions A
- Transitions $P(s' | s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)


- Important MDP quantities:

- Policy = Choice of action for each state
- Utility = expected sum of (discounted) rewards = “expected return”




Fixed Policies

Choosing actions



Do what π says to do



- If we fixed some policy $\pi(s)$, then the computation is simpler – only one action per state
 - ... though the performance now depend on which policy we fixed


Performance of a Fixed Policy

- Goal: compute the utility of a state s under a fixed (generally non-optimal) policy
- Define the utility of a state s , under a fixed policy π :
 $V^\pi(s)$ = expected total discounted rewards starting in s and following π
- Recursive relation (one-step look-ahead):

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


$P(s' | s, \pi(s))$

$$= \mathbb{E}_{s'} [R(s, \pi(s), s') + \gamma V^\pi(s')]$$




Example: Policy Evaluation

Always Go Right



Always Go Forward



Example: Policy Evaluation

Always Go Right



$$\gamma = 0.9$$

Always Go Forward



Policy Evaluation


- How do we calculate the V 's for a fixed policy π ?
- Idea 1: Turn recursive Bellman equations into updates

$$V_0^\pi(s) = 0$$


$$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

$V_1^\pi(s)$

- Efficiency: $O(S^2)$ per iteration
- Idea 2: Just a linear system
 - Solve with Numpy or Matlab (or your favorite linear system solver)



Policy Evaluation



- Idea 2: The Policy Evaluation Bellman equations are just a linear system

- Solve with Numpy or Matlab (or your favorite linear system solver)

$$V = \begin{bmatrix} V(s_1) \\ \vdots \\ V(s_n) \end{bmatrix}$$

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')] \quad \text{distribute}$$

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') R(s, \pi(s), s') + \gamma \sum_{s'} T(s, \pi(s), s') V^\pi(s')$$

$$V^\pi(s) = \bar{R}(s) + \gamma \sum_{s'} T(s, \pi(s), s') V^\pi(s')$$

$$T^\pi(i, j) = P(j|i, \pi)$$

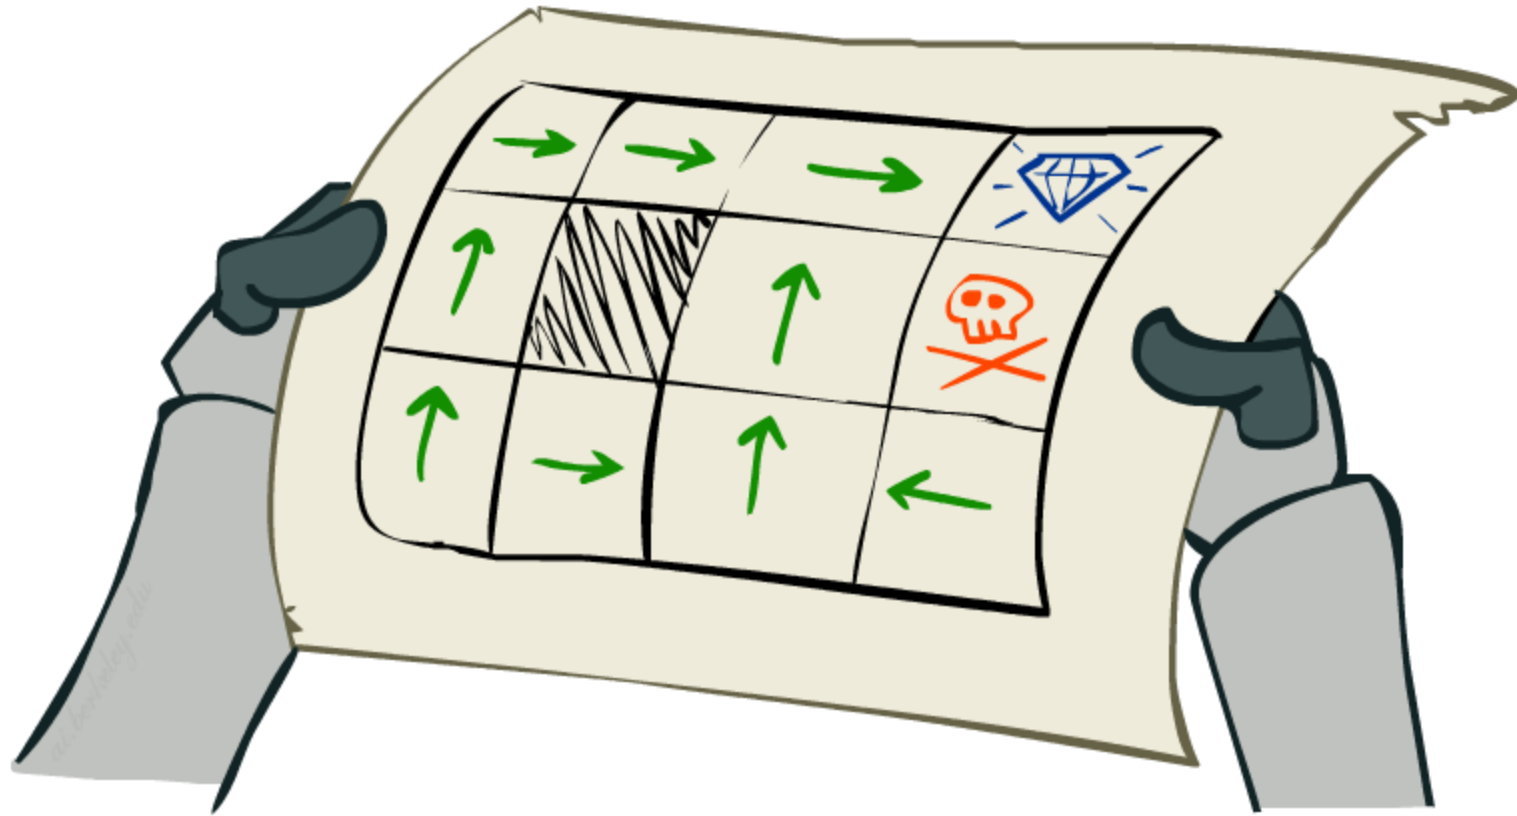
$$V^\pi = \bar{R} + \gamma T^\pi V^\pi \Rightarrow V^\pi - \gamma T^\pi V^\pi = \bar{R} \Rightarrow (I - \gamma T^\pi) V^\pi = \bar{R}$$

151×1 151×1 151×151

$$(I - \gamma T^\pi) V^\pi = \bar{R} \Rightarrow V^\pi = (I - \gamma T^\pi)^{-1} \bar{R}$$


$Ax = b$

Solving MDPs




Optimal Quantities

- The value (utility) of a state s :
 $V^*(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a) :
 $Q^*(s,a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally
- The optimal policy:
 $\pi^*(s)$ = optimal action from state s
$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



Can we write the optimal policy in terms of Q^* ?

The Bellman Equations



How to be optimal:

Step 1: Take correct first action


Step 2: Keep being optimal

Bellman Equations

- Fundamental operation: compute the (expectimax) value of a state
 - Expected utility under optimal action
 - Average sum of (discounted) rewards
 - This is just what expectimax computed!

det $\pi: S \rightarrow A$

stoch $\pi: S \rightarrow \mathcal{A}(A)$



- Recursive definition of value:

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Aside: Different ways to write Bellman Eqns

- What if R only depends on state and action? e.g. $R(s,a,s') = R(s,a)$

$$V^*(s) = \max_a Q^*(s, a)$$

$$\begin{aligned} Q^*(s, a) &= \sum_{s'} T(s, a, s') [R(s, a) + \gamma V^*(s')] \\ &= \sum_{s'} T(s, a, s') R(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s') \\ &= R(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s') \end{aligned}$$

Aside: Different ways to write Bellman Eqns

- What if R only depends on state? e.g. $R(s,a,s') = R(s)$ ~~R~~ HW 3

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = R(s) + \gamma \sum_{s'} T(s, a, s') V^*(s')$$

$\gamma \uparrow$
 $\tau \rightarrow \infty$

Value Iteration

- Start with $V_0(s) = 0$: no time steps left means an expected reward sum of zero
- Given vector of $V_k(s)$ values, do one ply of expectimax from each state:


$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Bellman Update Equation


- Repeat until convergence

$$\max_s |V_{k+1}(s) - V_k(s)| < \epsilon$$

- Complexity of each iteration: $O(S^2A)$
- Theorem: will converge to unique optimal values
 - Basic idea: approximations get refined towards optimal values
 - Policy may converge long before values do



$k=0$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=1




Noise = 0.2
Discount = 0.9
Living reward = 0

k=2



Noise = 0.2
Discount = 0.9
Living reward = 0

k=3



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=4$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=5



Noise = 0.2
Discount = 0.9
Living reward = 0

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=8



Noise = 0.2
Discount = 0.9
Living reward = 0

k=9



Noise = 0.2
Discount = 0.9
Living reward = 0

k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

k=12



Noise = 0.2
Discount = 0.9
Living reward = 0

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration


- Bellman equations **characterize** the optimal values:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


- Value iteration **computes** them:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Value iteration is just a fixed point solution method
 - ... though the V_k vectors are also interpretable as time-limited values



Policy Extraction



Computing Actions from Values

- Let's imagine we have the optimal values $V^*(s)$
- How should we act?
 - It's not obvious!
- We need to do a mini-expectimax (one step)



$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$


Handwritten red text: $Q^(s, a)$ with an arrow pointing to the summation term.*

- This is called **policy extraction**, since it gets the policy implied by the values

Computing Actions from Q-Values


- Let's imagine we have the optimal q-values:
- How should we act?
 - Completely trivial to decide!

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



- Important lesson: actions are easier to select from q-values than values!

Policy Iteration




Problems with Value Iteration

- Value iteration repeats the Bellman updates:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Problem 1: It's slow – $O(S^2A)$ per iteration
- Problem 2: The “max” at each state rarely changes
- Problem 3: The policy often converges long before the values



k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

Step 0: start w/ random π

- Alternative approach for optimal values:
 - **Step 1: Policy evaluation:** calculate utilities for some fixed policy (not optimal utilities!) until convergence
 - **Step 2: Policy improvement:** update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
 - Repeat steps until policy converges
- This is **policy iteration**
 - It's still optimal!
 - Can converge (much) faster under some conditions

Policy Iteration

- Evaluation: For fixed current policy π_i find values with policy evaluation:
 - Iterate until values converge:
- alternative: run for K iterations*

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

- Improvement: For fixed values, get a better policy using policy extraction
 - One-step look-ahead:

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

$Q^{\pi_i}(s, a)$

stop when $\pi_{i+1} = \pi_i$

Comparison

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In value iteration:
 - Every iteration updates both the values and (implicitly) the policy
 - We don't track the policy, but taking the max over actions implicitly recomputes it
- In policy iteration:
 - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - The new policy will be better (or we're done)
- Both are dynamic programs for solving MDPs

Aside: Linear Programming

x_1, x_2 : Decision variables

max $350x_1 + 300x_2$
subject to

$$\begin{aligned}x_1 + x_2 &\leq 200 \\9x_1 + 6x_2 &\leq 1566 \\12x_1 + 16x_2 &\leq 2880 \\x_1, x_2 &\geq 0\end{aligned}$$

Objective function

Constraints

General form

$$\begin{aligned}\max \quad & c^T x \\ \text{s.t.} \quad & Ax \leq b \\ & x \geq 0\end{aligned}$$

Primal Linear Programming Solutions

Basic idea: we can capture the constraint

$$V(s) \geq R(s) + \gamma \max_{a \in \mathcal{A}} \sum_{s' \in \mathcal{S}} P(s, |s, a) V(s')$$

via the set of $|\mathcal{A}|$ linear constraints

$Ax \geq b$ {

$$V(s) \geq R(s) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall a \in \mathcal{A}$$

Primal Linear Programming Solutions

Now consider the linear program

We want $V^*(s) \forall s$

$$\begin{aligned} & \underset{V}{\text{minimize}} && \sum_s V(s) \\ & \text{subject to} && V(s) \geq R(s) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall a \in \mathcal{A}, s \in \mathcal{S} \end{aligned}$$

Dual Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\pi : S \times A \mapsto [0, 1]$$



$$u_{sa} = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t \mathbf{1}_{(s_t=s, a_t=a)} \right]$$

Expected discounted # times
we take a in s

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\pi(a|s) = \frac{u_{sa}}{\sum_a u_{sa}}$$



$$u_{sa} = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t \mathbf{1}_{(s_t=s, a_t=a)} \right]$$

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \sum_{s,a} r_{sa} u_{sa}$$

Reward for taking
action a in state s

$R(s,a)$

such that

$$\sum_a u_{sa} = p_0(s) + \gamma \sum_{s',a} u_{s'a} P(s', a, s), \forall s$$

$$u_{sa} \geq 0, \forall s, a$$

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \sum_{s,a} r_{sa} u_{sa}$$

$$u_{sa} = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t \mathbf{1}_{(s_t=s, a_t=a)} \right]$$

State Occupancies

such that

*Bellman
Flow
constraint \leftrightarrow*

$$\sum_a u_{sa} = p_0(s) + \gamma \sum_{s',a} u_{s'a} P(s', a, s), \forall s$$

$$u_{sa} \geq 0, \forall s, a$$

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \quad \sum_{s,a} r_{sa} u_{sa}$$

such that

$$\sum_a u_{sa} = \overset{\text{Initial state distribution}}{p_0(s)} + \gamma \sum_{s',a} u_{s'a} P(s', a, s), \forall s$$

$$u_{sa} \geq 0, \forall s, a$$

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \sum_{s,a} r_{sa} u_{sa}$$

such that

$$\sum_a u_{sa} = p_0(s) + \gamma \sum_{s',a} u_{s'a} \overset{\text{Transition Probability}}{P(s', a, s)}, \forall s$$

$$u_{sa} \geq 0, \forall s, a$$

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \sum_{s,a} r_{sa} u_{sa}$$

Reward for taking action a in state s

such that

$$\sum_a u_{sa} = p_0(s) + \gamma \sum_{s',a} u_{s'a} P(s', a, s), \forall s$$

Initial state distribution

Transition Probability

State Occupancies

$$u_{sa} \geq 0, \forall s, a$$

Discount factor

Linear Programming for MDPs

- One-to-one correspondence between stochastic policies and state-action occupancy frequencies:

$$\max_u \sum_{s,a} r_{sa} u_{sa}$$



such that

How often do I
start in s ?

How often do I visit
other states s' and then
transition to state s ?

How often do
I visit state s ?

$$\sum_a u_{sa} = p_0(s) + \gamma \sum_{s',a} u_{s'a} P(s', a, s), \forall s$$

$$u_{sa} \geq 0, \forall s, a$$

Summary: MDP Algorithms

- So you want to....
 - Compute optimal values: use value iteration or policy iteration
 - Compute values for a particular policy: use policy evaluation
 - Turn your values into a policy: use policy extraction (one-step lookahead)
- These all look the same!
 - They basically are – they are all variations of Bellman updates
 - They all use one-step lookahead computations